

The Continuing Saga of the Lock-free Queue

Part 3 of N

Tony Van Eerd C++Now May 2018

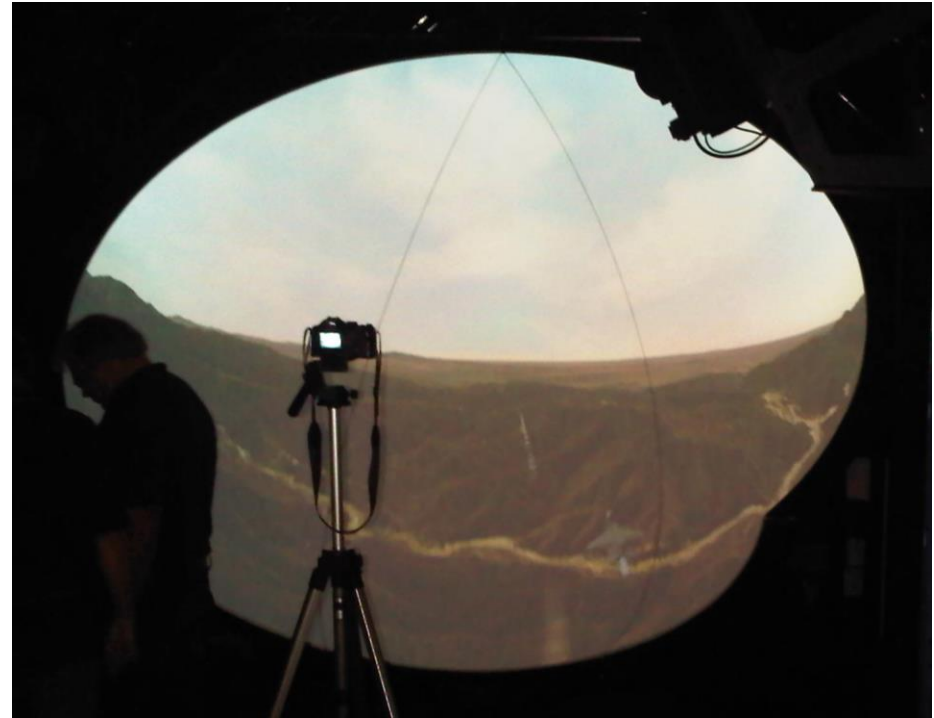
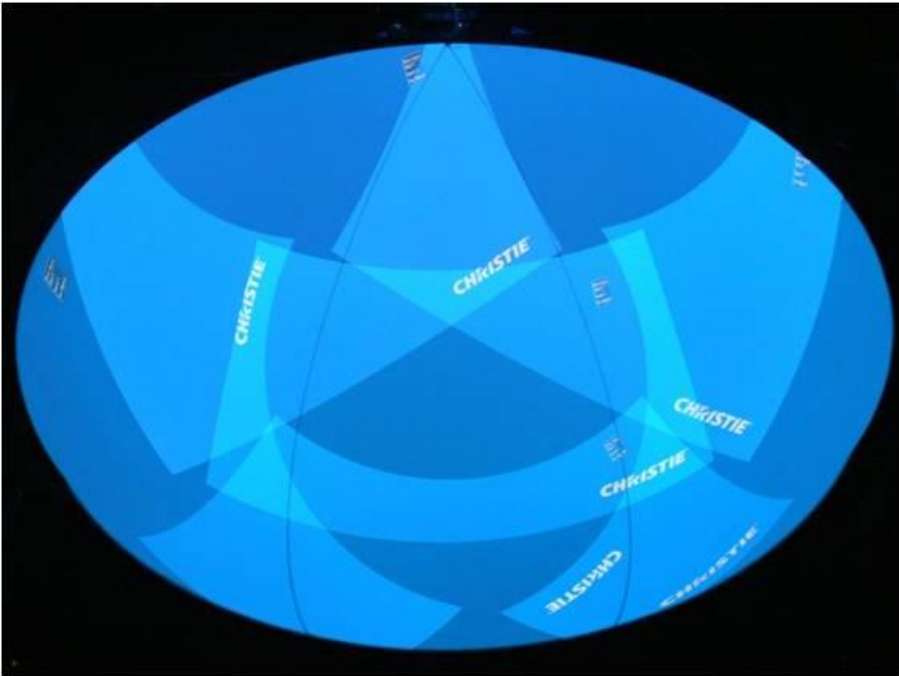


CHRISTIE®

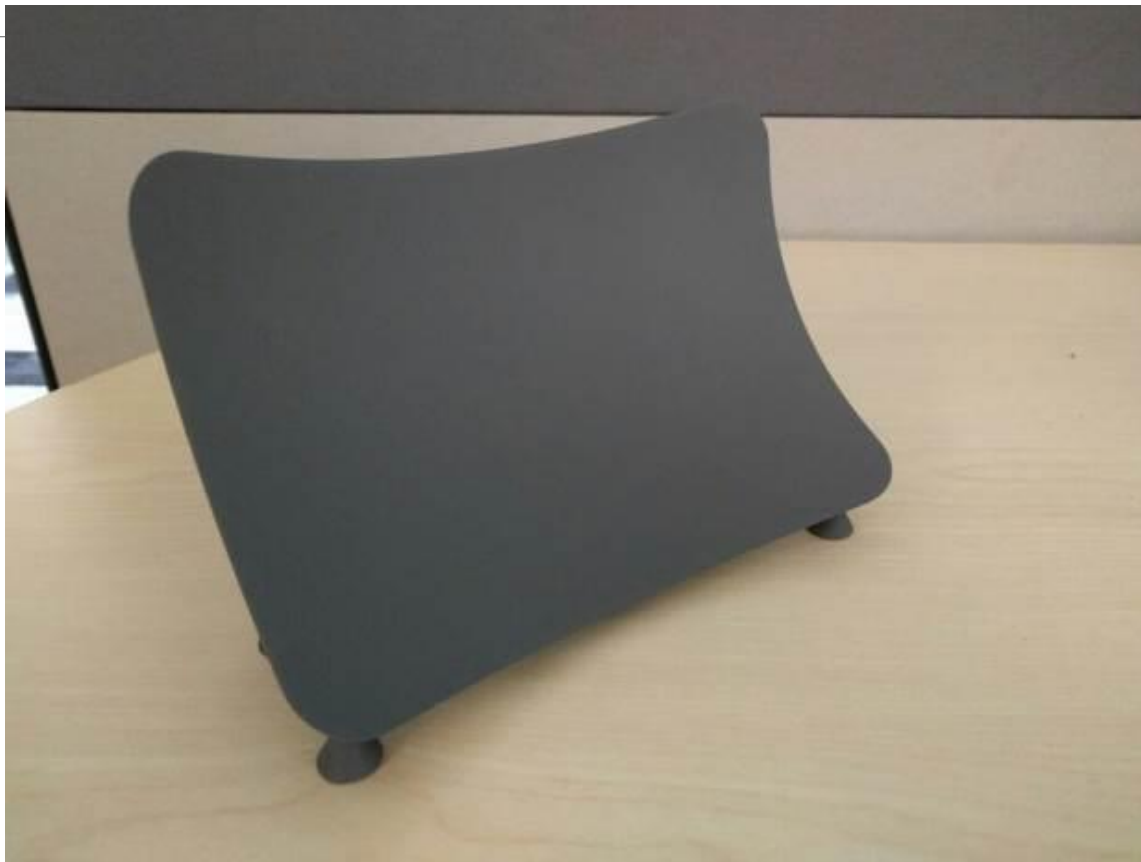


















CHRISTIE

CHRISTIE

Christie Mystique
Spectacular experiences. Simple solutions.

Design

Install

Operate

ALASKA

BOEHME

Christie Soft
Superior color
warping and







youtu.be/sql1kD6_Uok





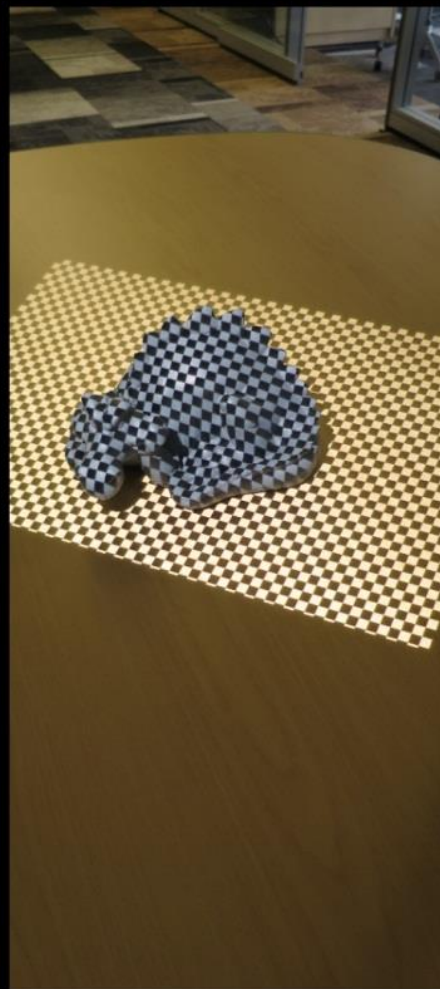
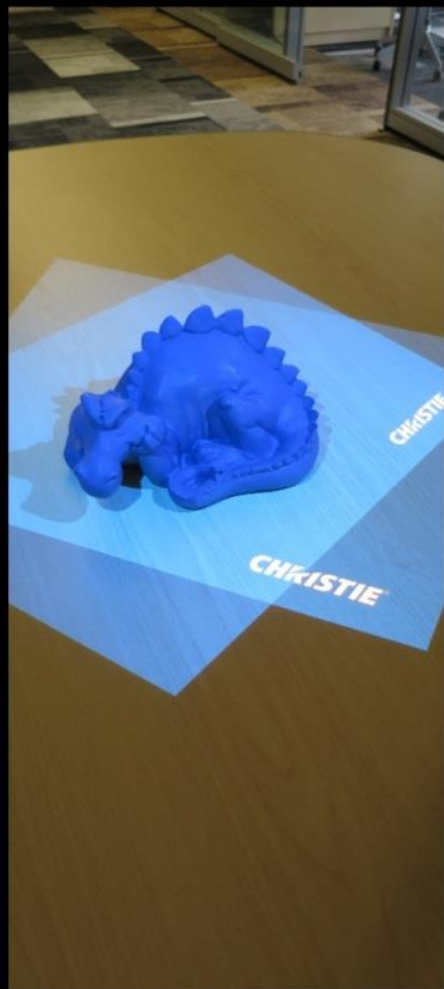












CHRISTIE®



Guide to Threaded Coding

Guide to Threaded Coding

1. Stop Sharing (forget what you learned in kindergarten)
2. OK, well Use Locks then
(don't call unknown code while holding a lock)
3. Measure
4. Measure
5. Change your Algorithm
6. GOTO 1

∞. **Lock-free**

Lock-free coding is the last thing you want to do.

Guide to Threaded Coding

1. Stop Sharing (forget what you learned in kindergarten)
2. OK, well Use Locks then
(don't call unknown code while holding a lock)
3. Measure
4. Measure
5. Change your Algorithm
6. GOTO 1

∞. **Lock-free**

∞+1. **Measure. Measure.**

Lock-free coding is the last thing you want to do.

Guide to Threaded Coding

Don't Share
Use Locks

Rules of Lock-free Coding

- 1.

Rules of Lock-free Coding

1. Don't talk about lock-free coding

Guide to Coding

Guide to Coding

MACROS_ARE_EVIL

Notes:

`acquire = std::memory_order_acquire`

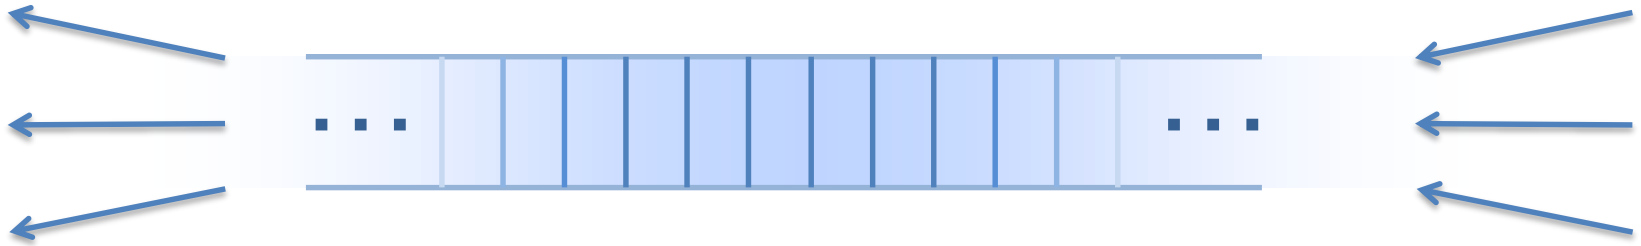
`release = std::memory_order_release`

(etc)

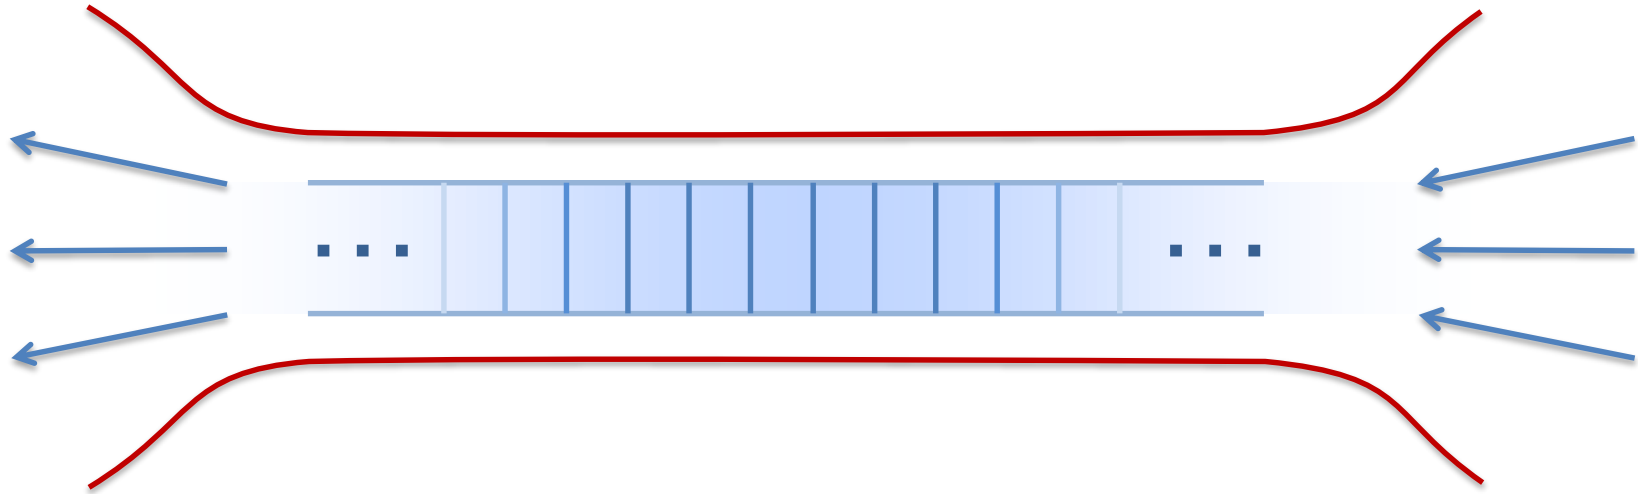
CAS = “Compare and Set/Swap” = `std::compare_exchange()` and friends

Not my coding style/structure

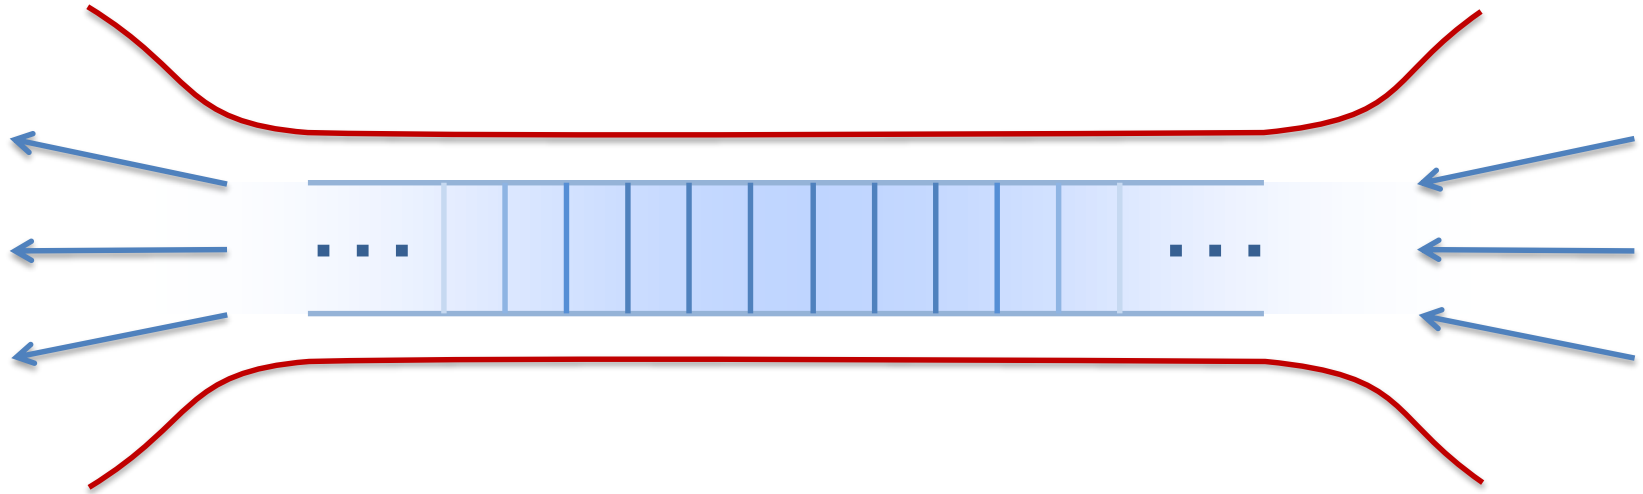
MPMC Queue

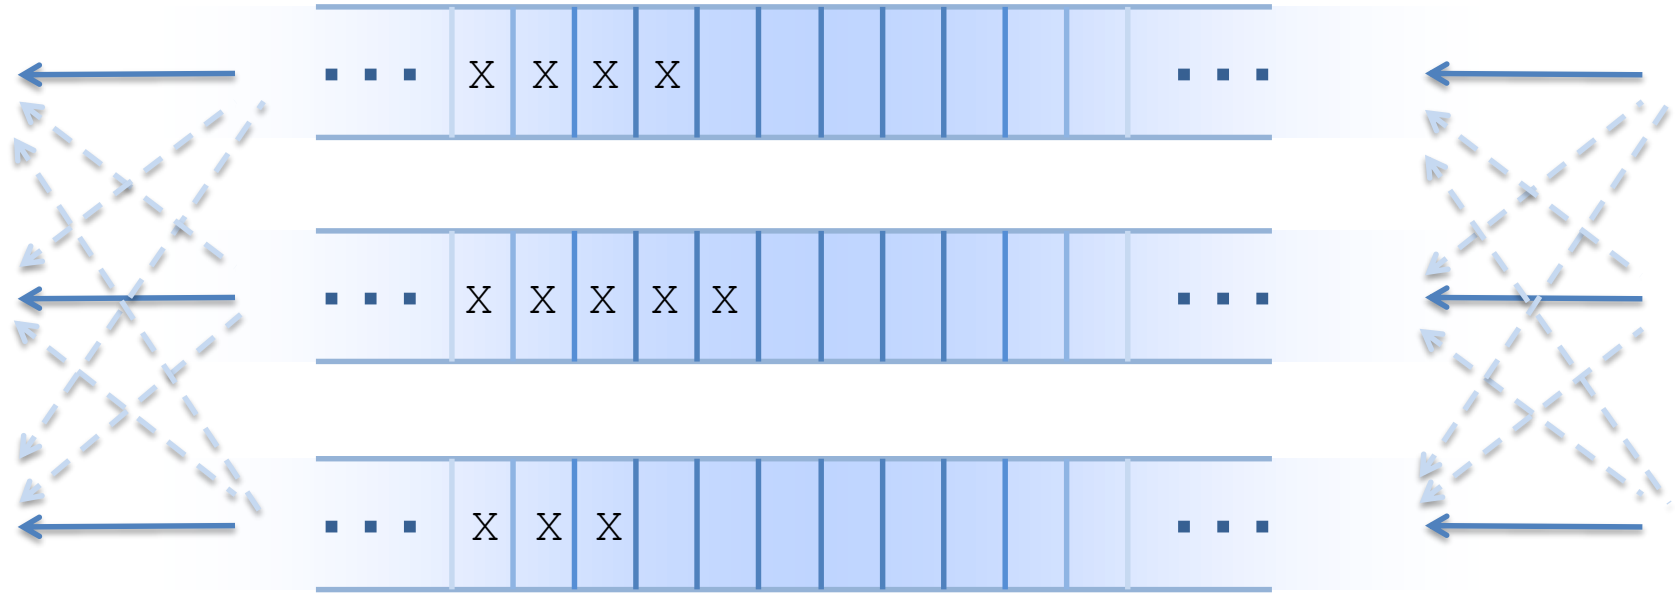


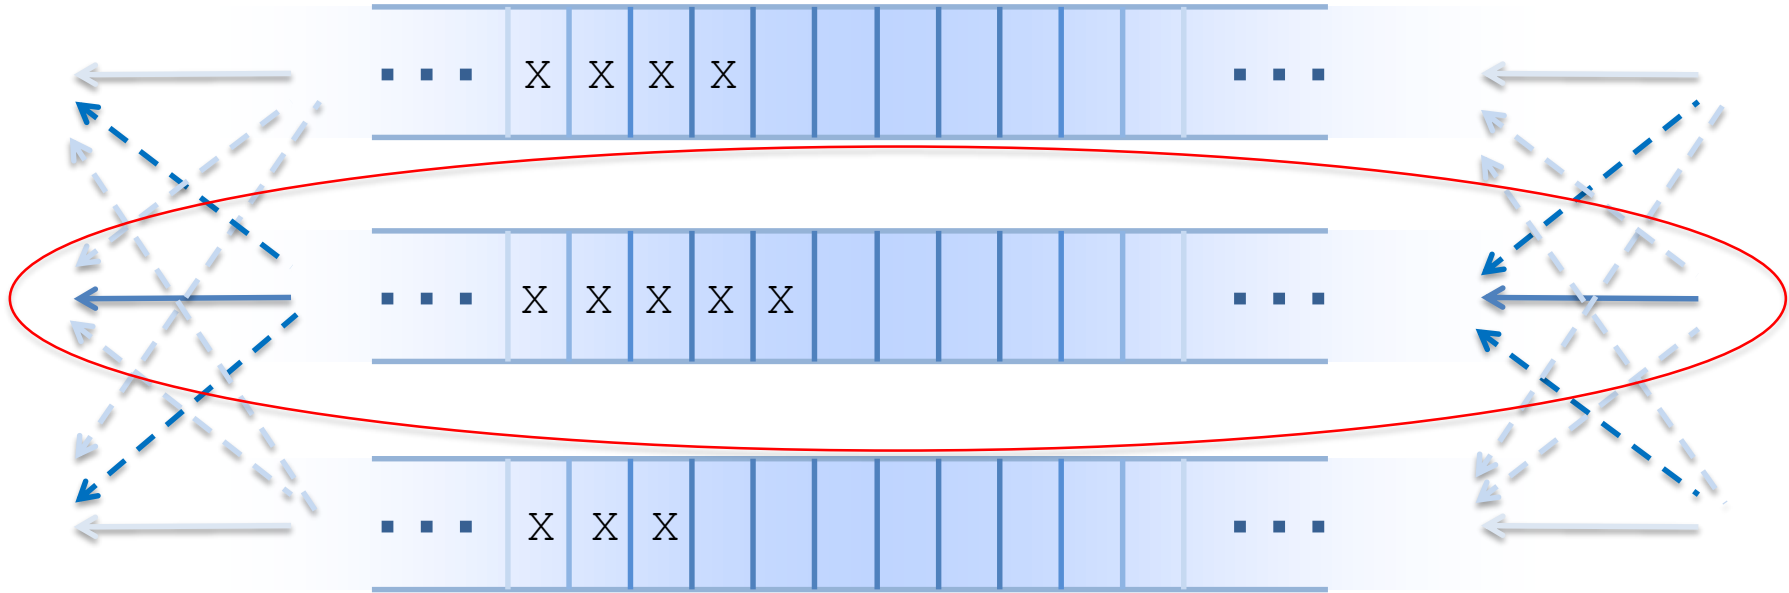
MPMC Queue

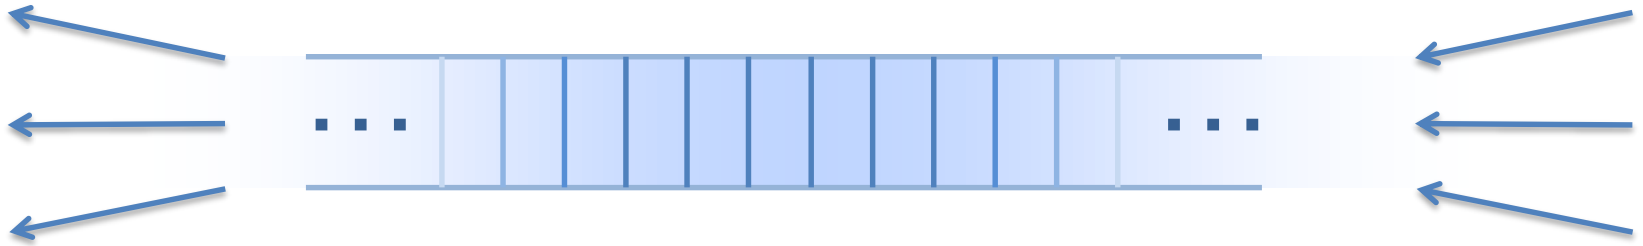


Bottleneck

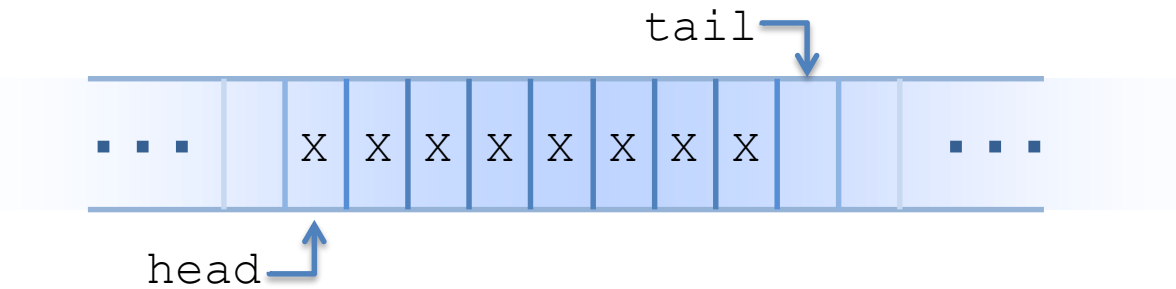






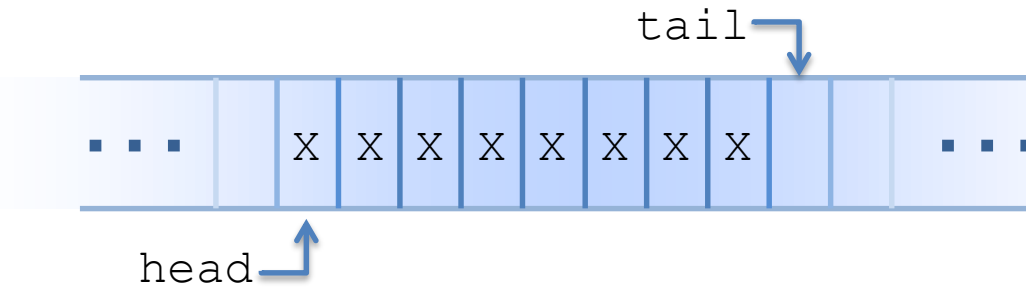


Review



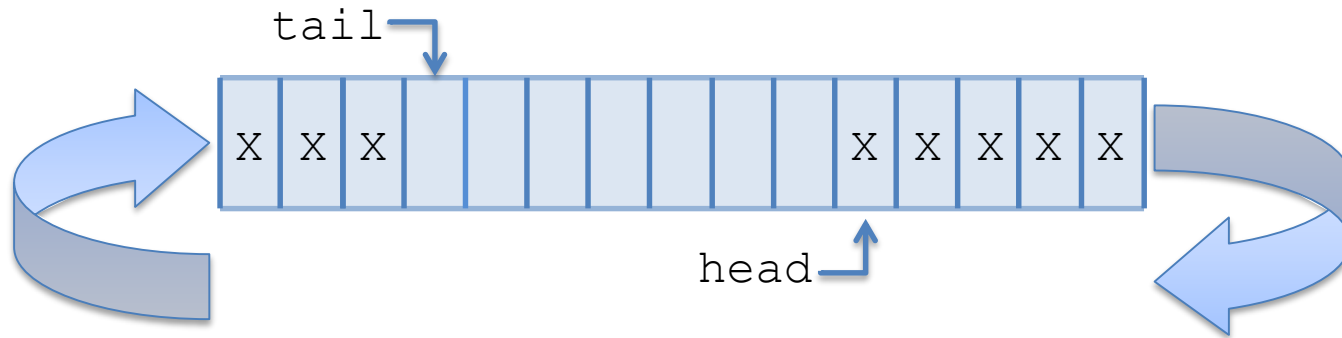
Review

```
class Queue {  
    ?  
};
```



Review

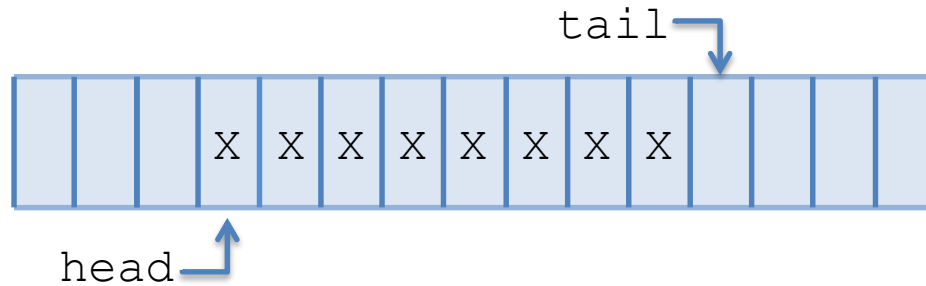
```
class Queue {  
    T buffer[SIZE];  
    int head;  
    int tail;  
};
```



Review

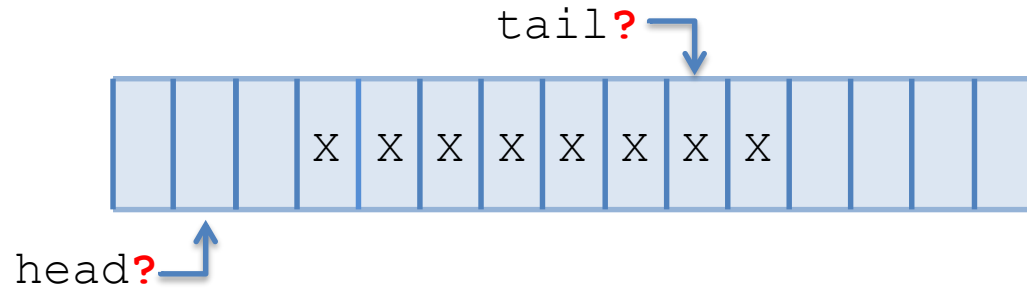
Compromise

```
class Queue {  
    int buffer[SIZE];  
    int head;  
    int tail;  
};
```



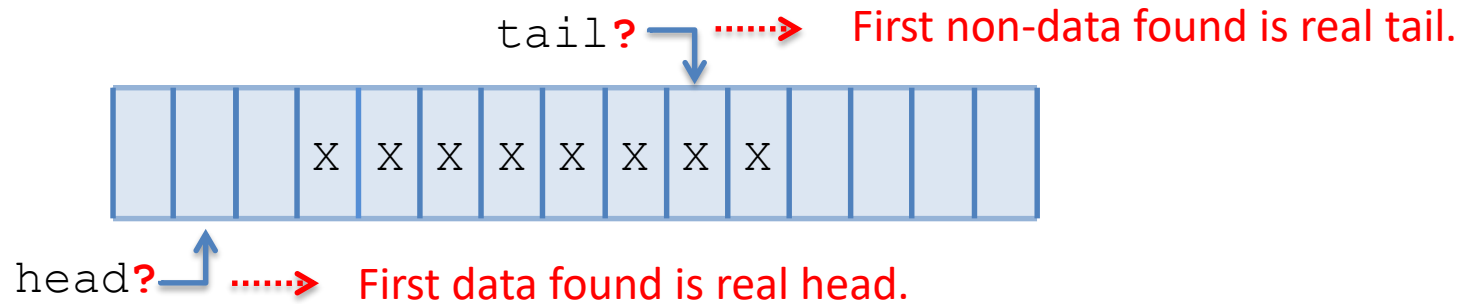
Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



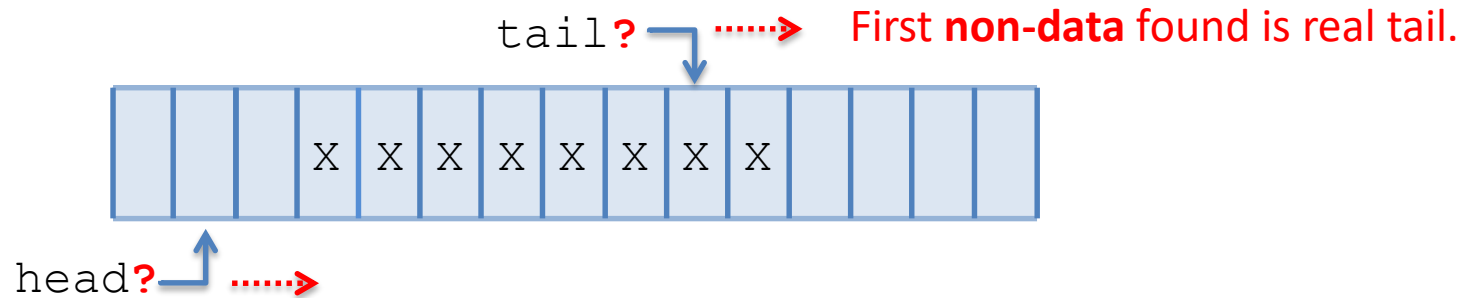
Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



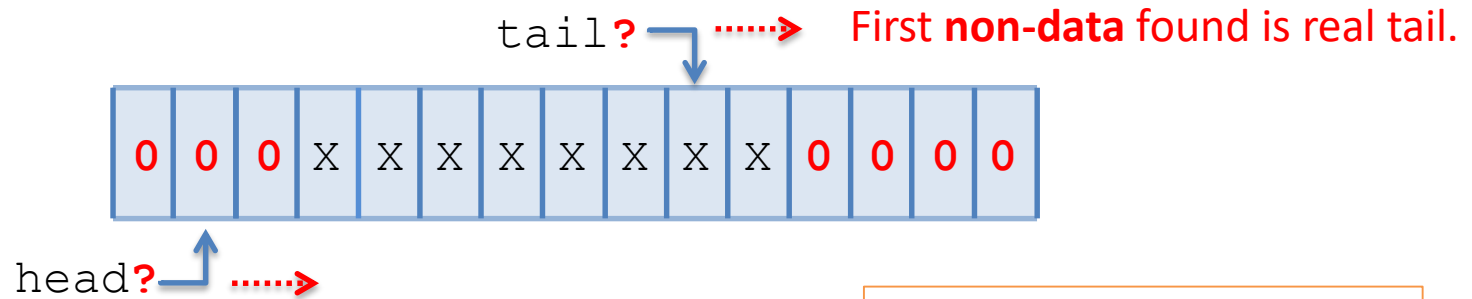
Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Review

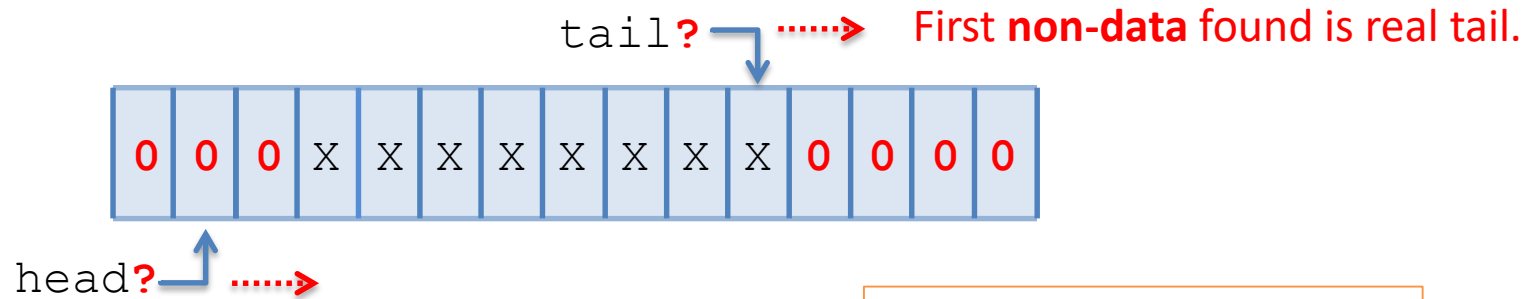
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Data == 0

Review

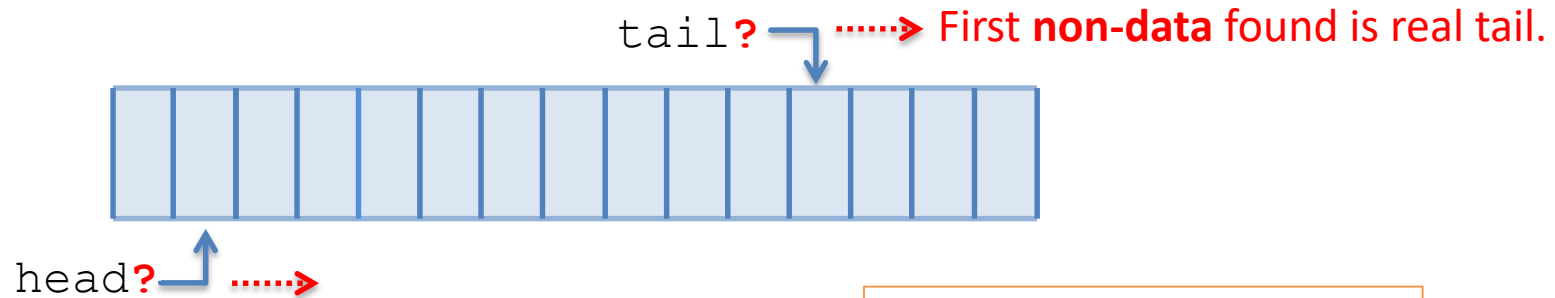
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Data == 0

Review

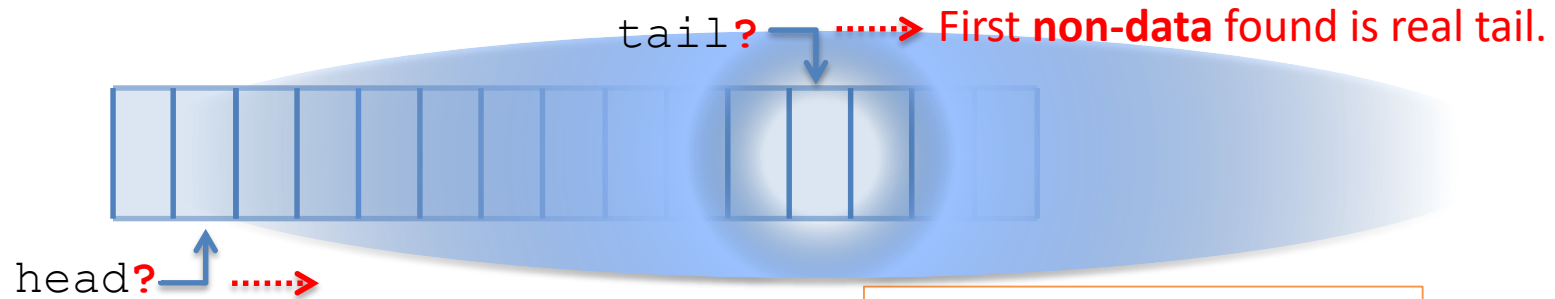
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Data == 0

Review

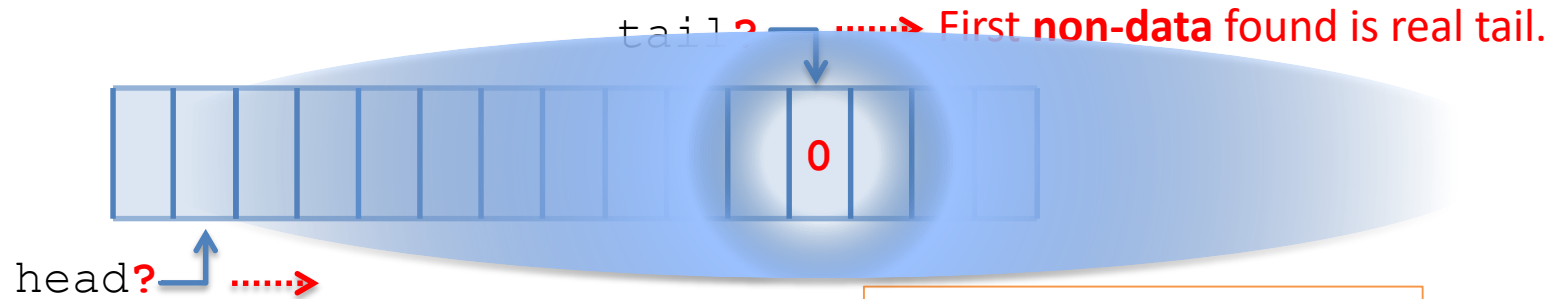
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Data == 0

Review

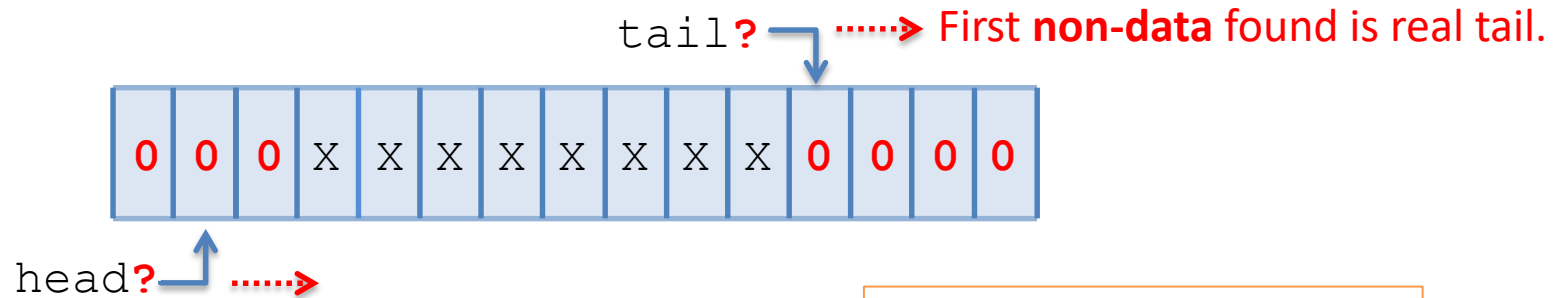
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Data == 0

Review

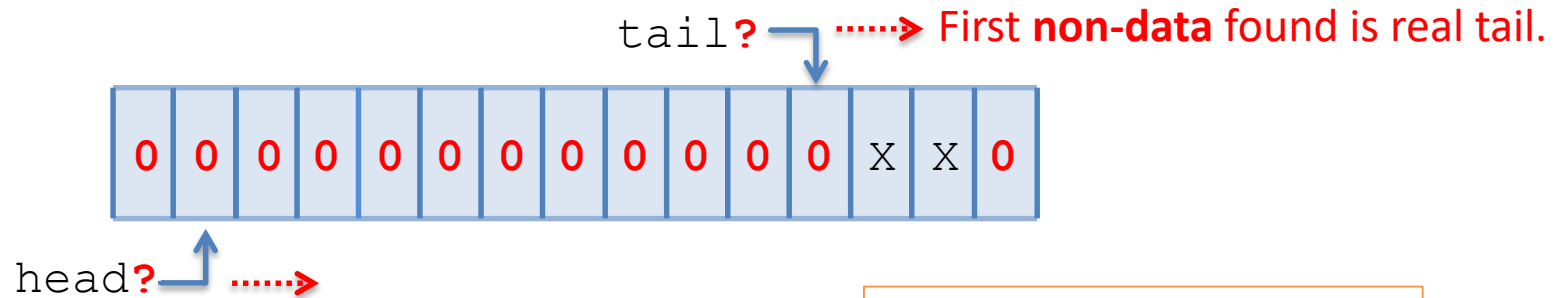
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Data == 0

Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```

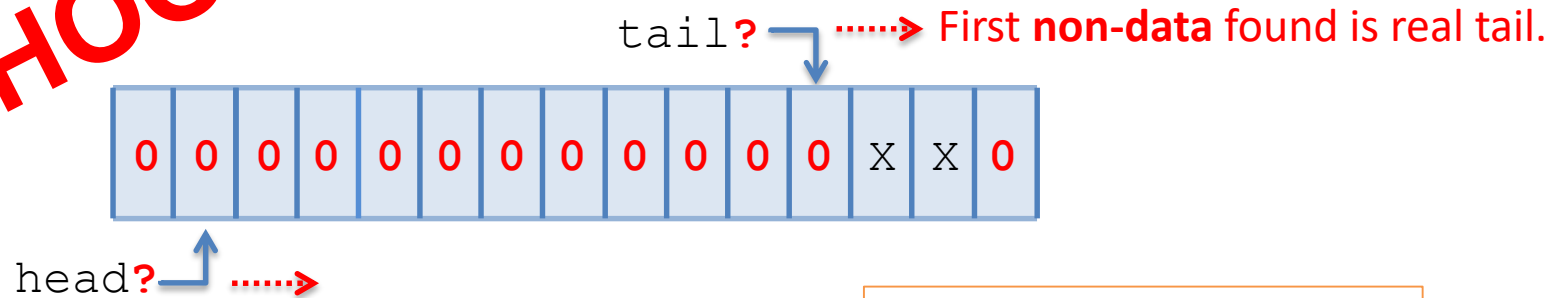


Data == 0

Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```

WHOOOPS!

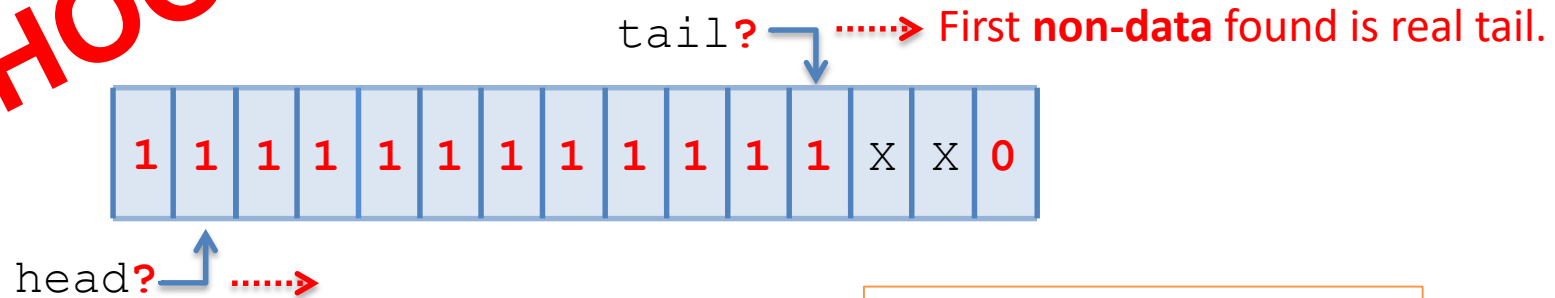


Data == 0

Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```

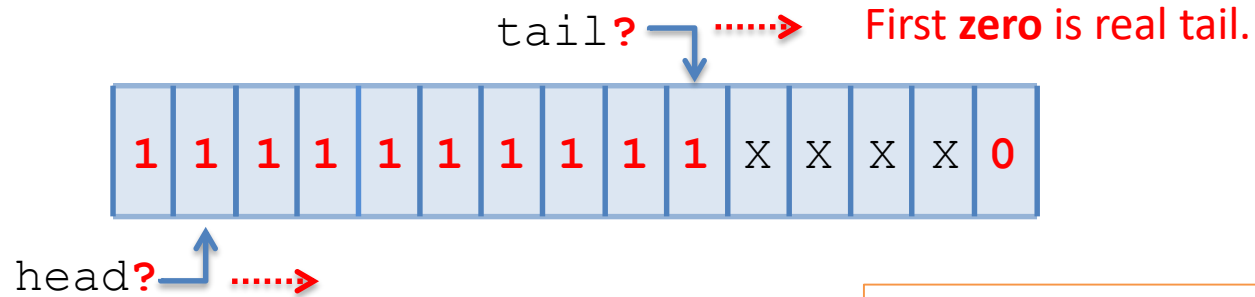
WHOOOPS!



Compromise

Review

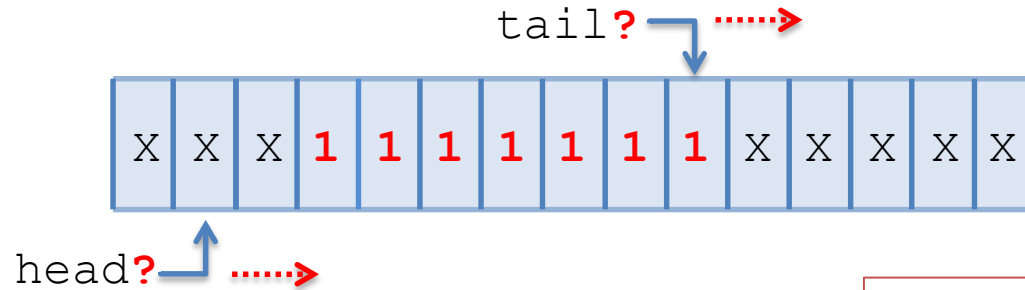
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Compromise

Review

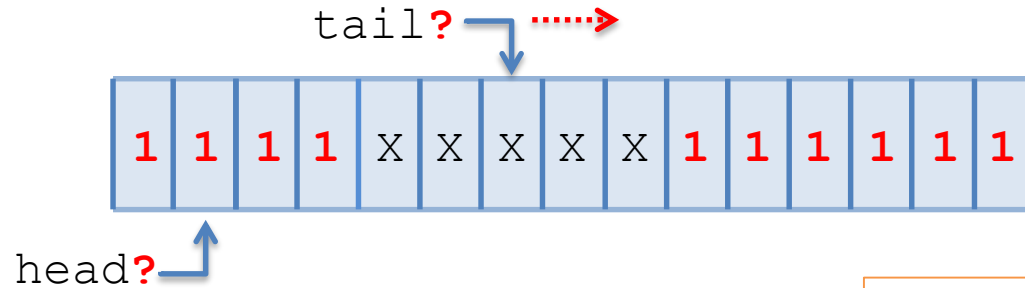
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Circular Buffer

Review

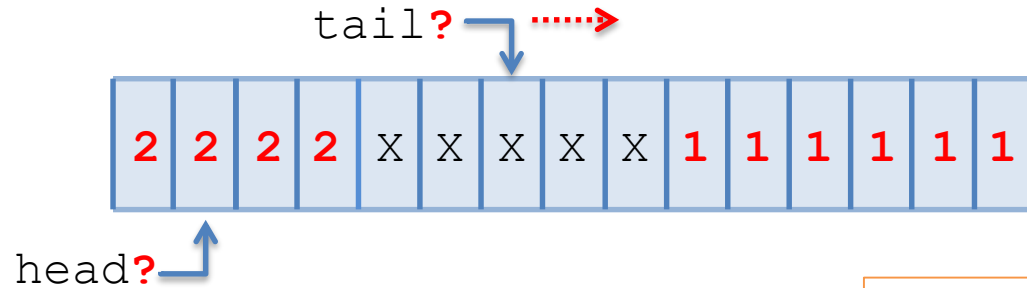
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Compromise

Review

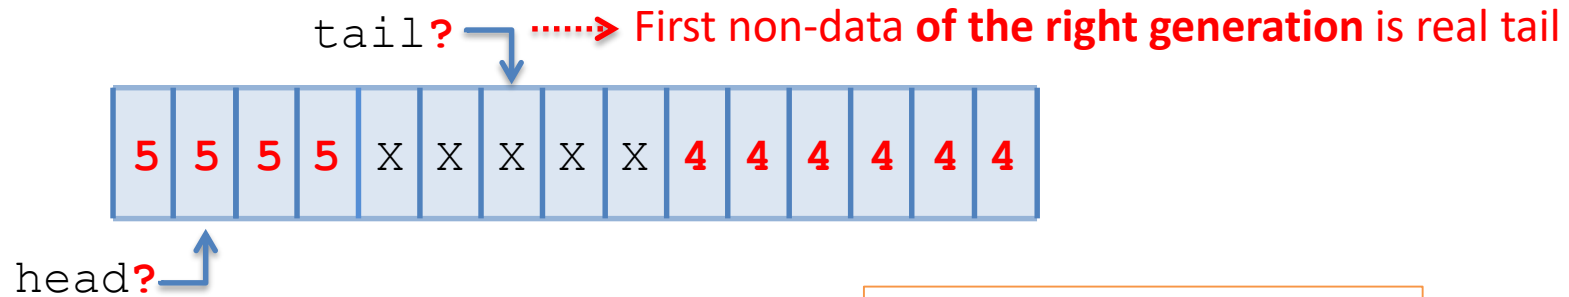
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
};
```



Compromise

Review

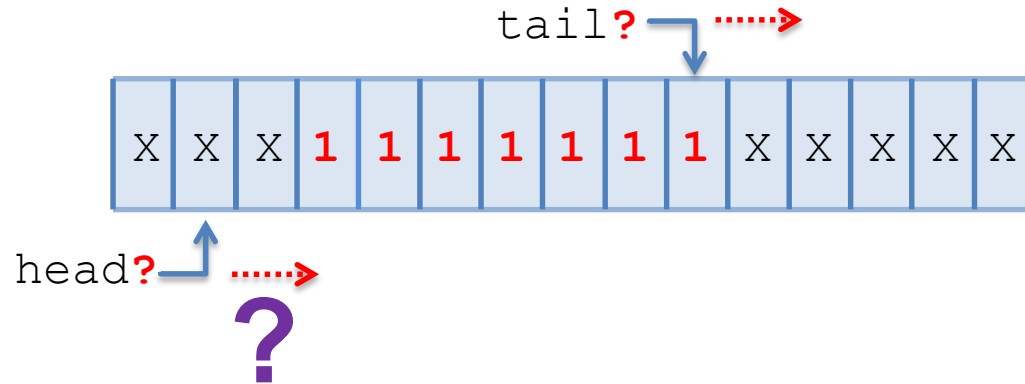
```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



Compromise

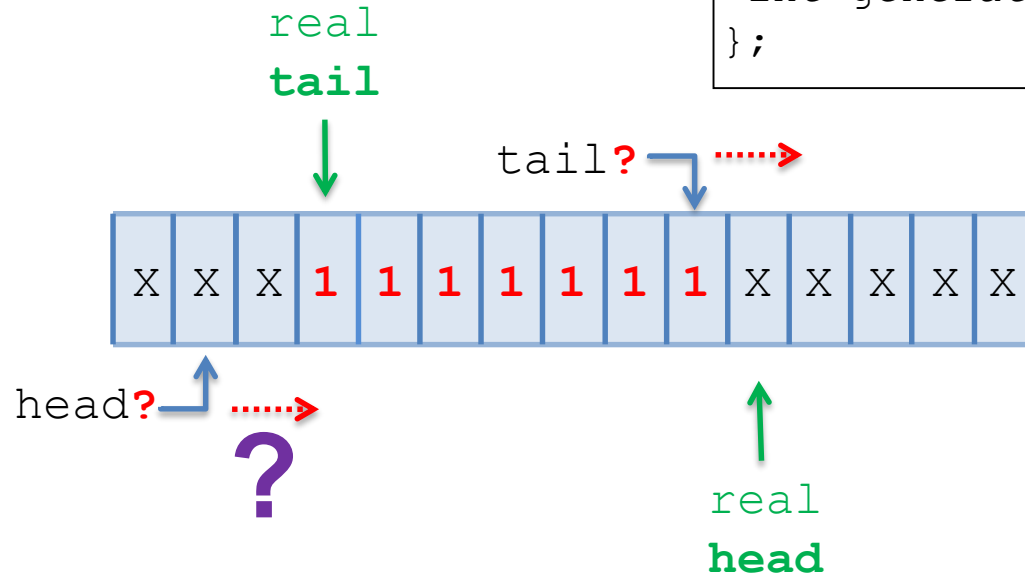
Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



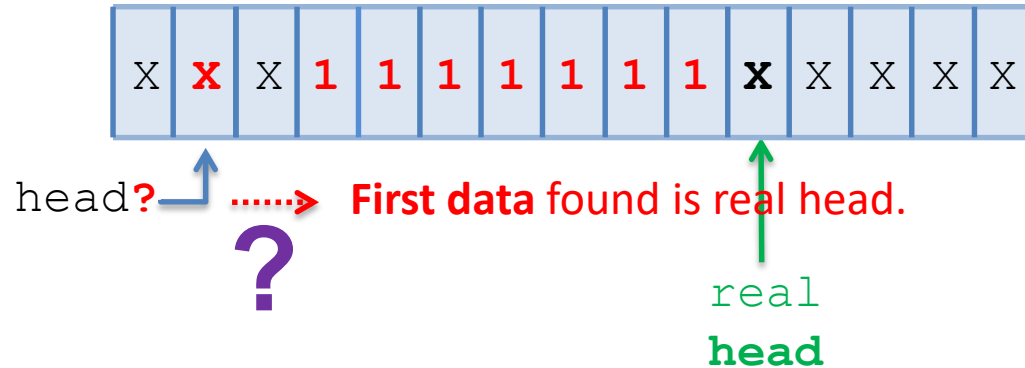
Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



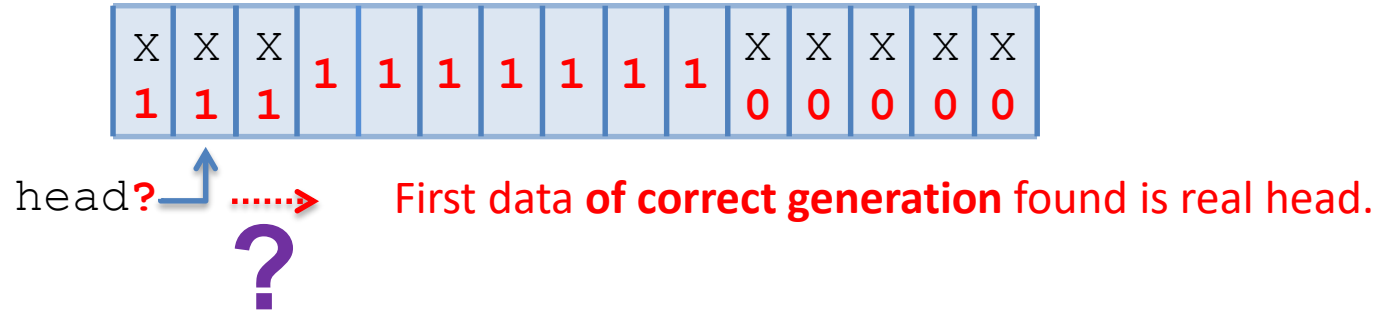
Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



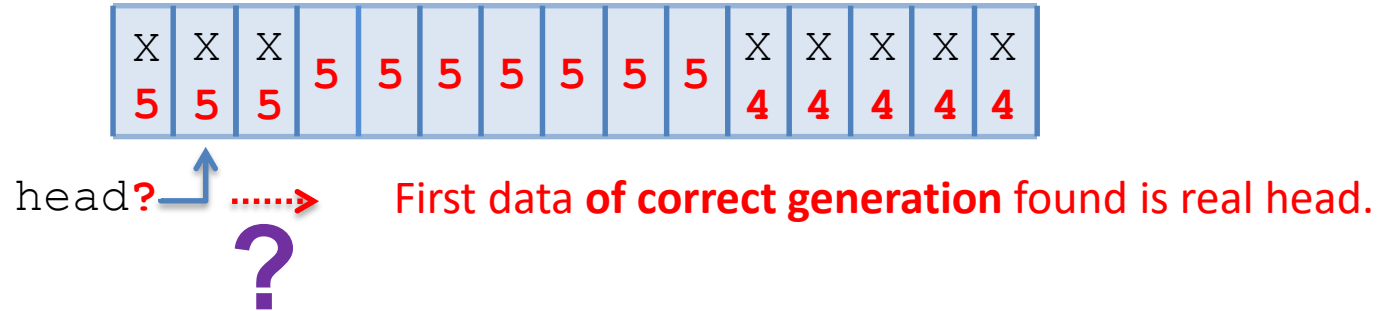
Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



Review

```
class Queue {  
    int buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```

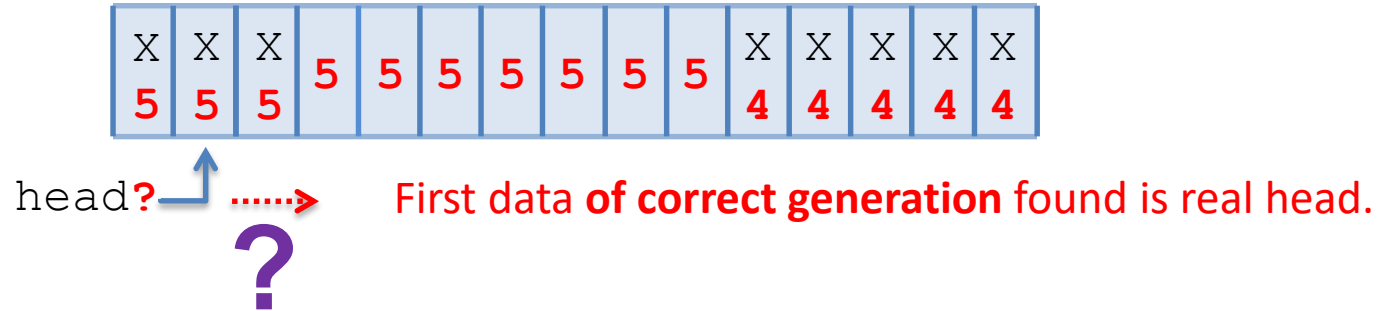


Review

```
struct entry {  
    int data;  
    int generation;  
};
```

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```

Compromise





Review

```
struct entry {  
    int data;  
    int generation;  
};
```

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```

Also: data != 0

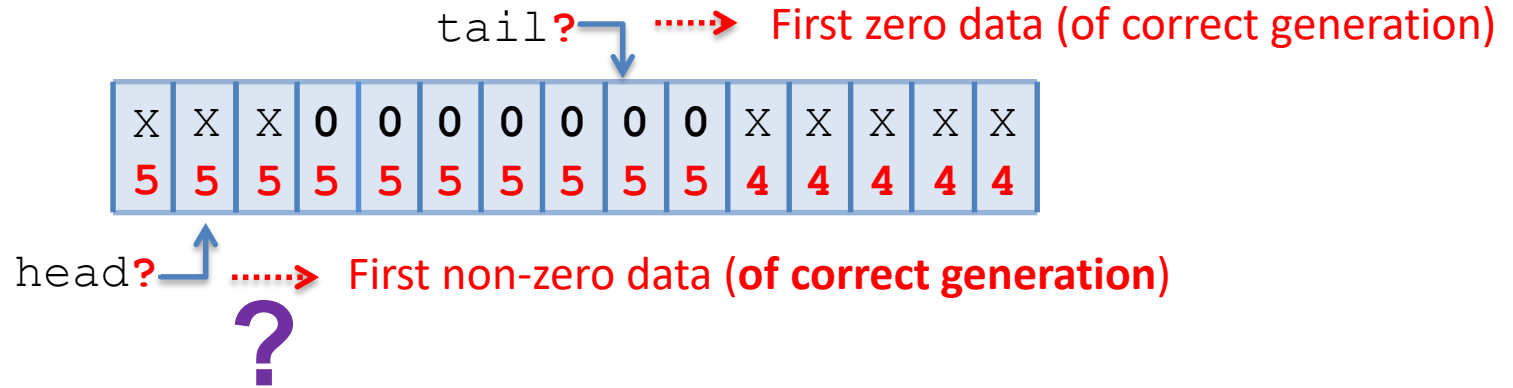
X	X	X	0	0	0	0	0	0	0	X	X	X	X	X
5	5	5	5	5	5	5	5	5	5	4	4	4	4	4

head? 


First data of correct generation found is real head.

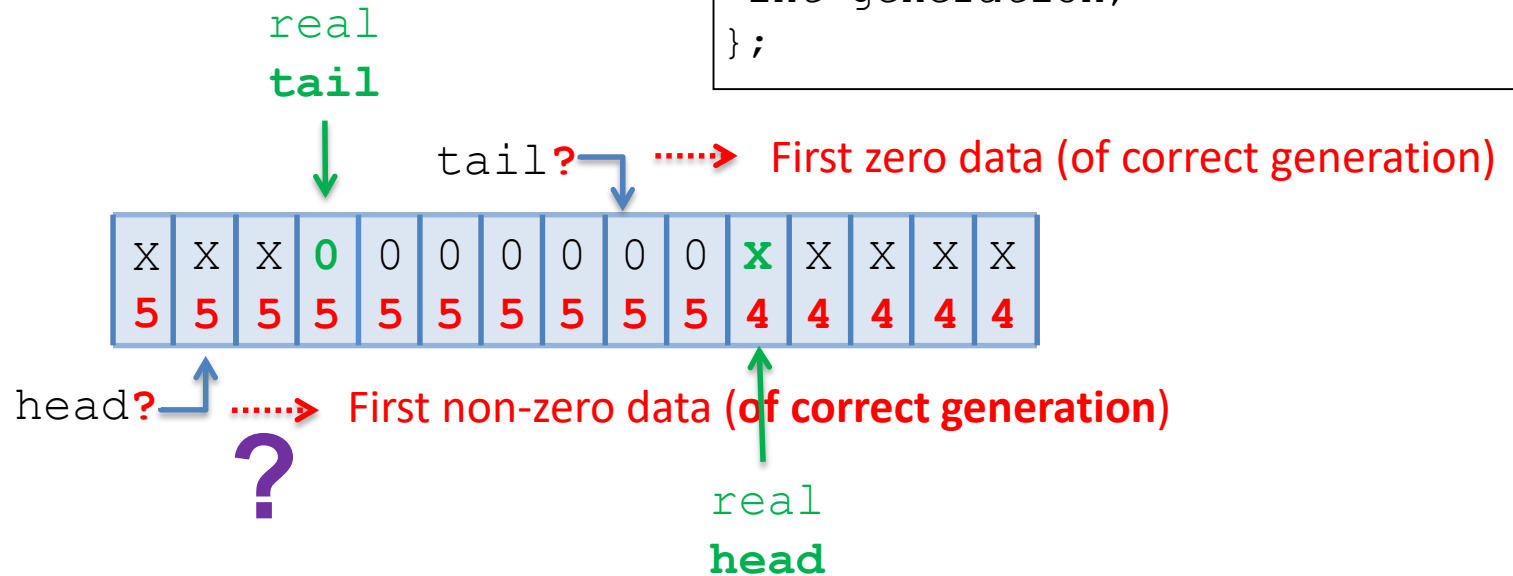
Review

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



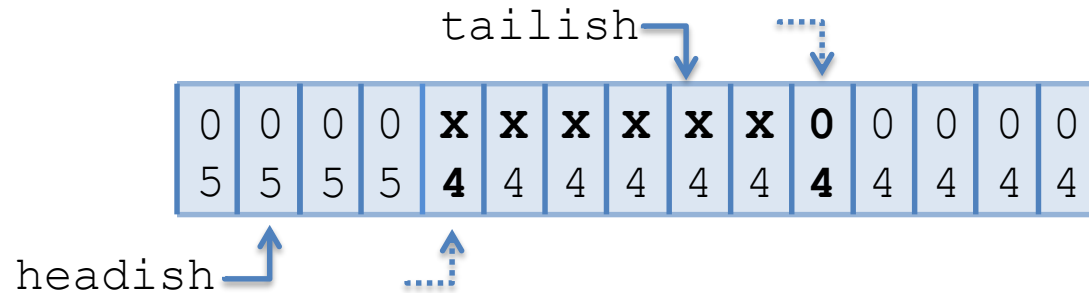
Review

```
class Queue {
    entry buffer[SIZE];
    int headish;
    int tailish;
    int generation;
};
```



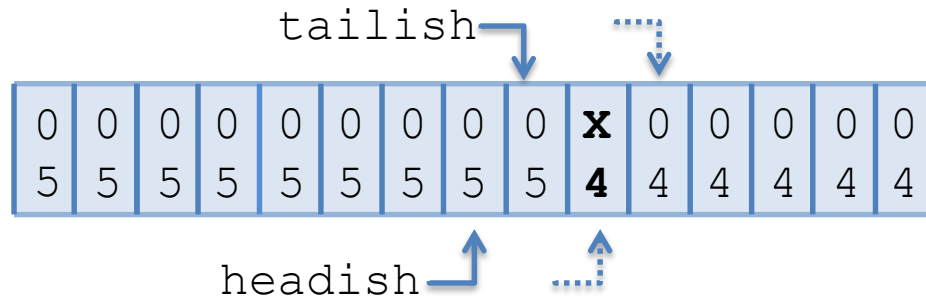
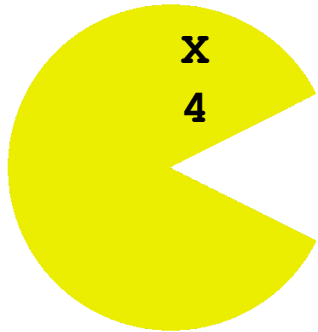
Examples

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



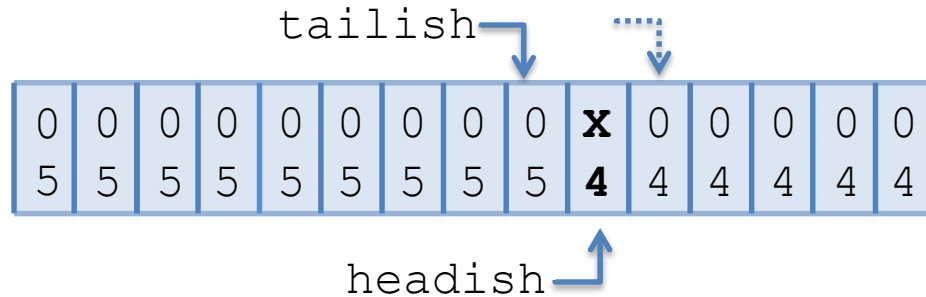
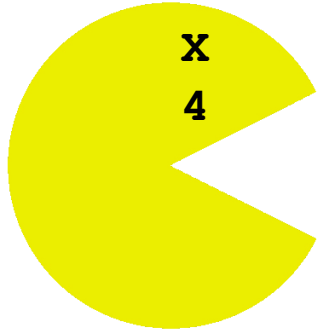
Examples

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



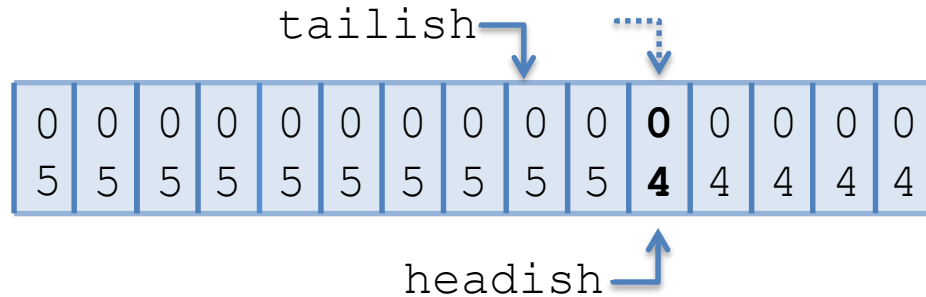
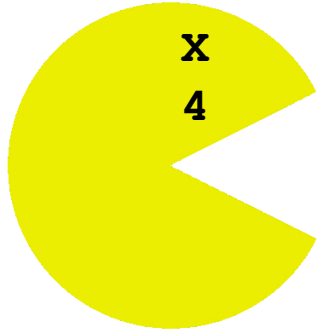
Examples

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



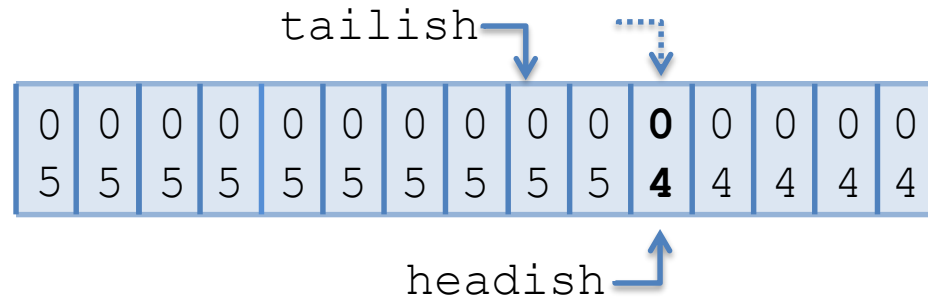
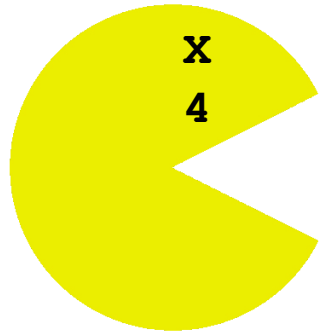
Examples

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



Examples

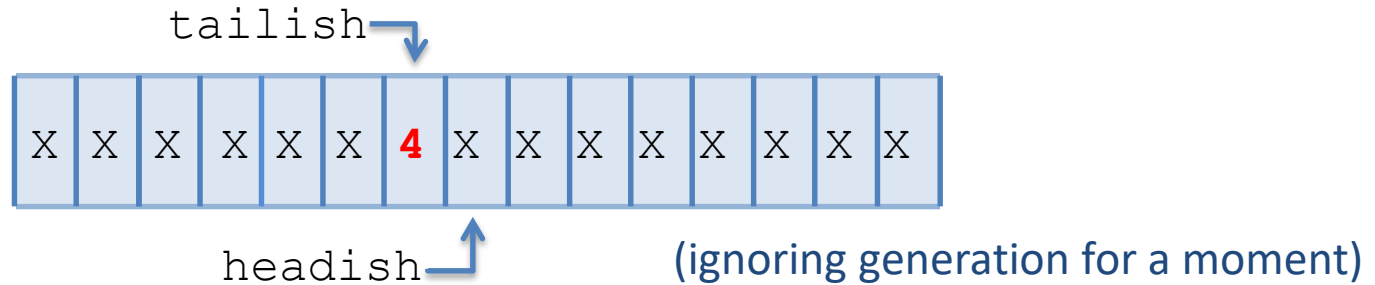
```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



Queue is Empty!

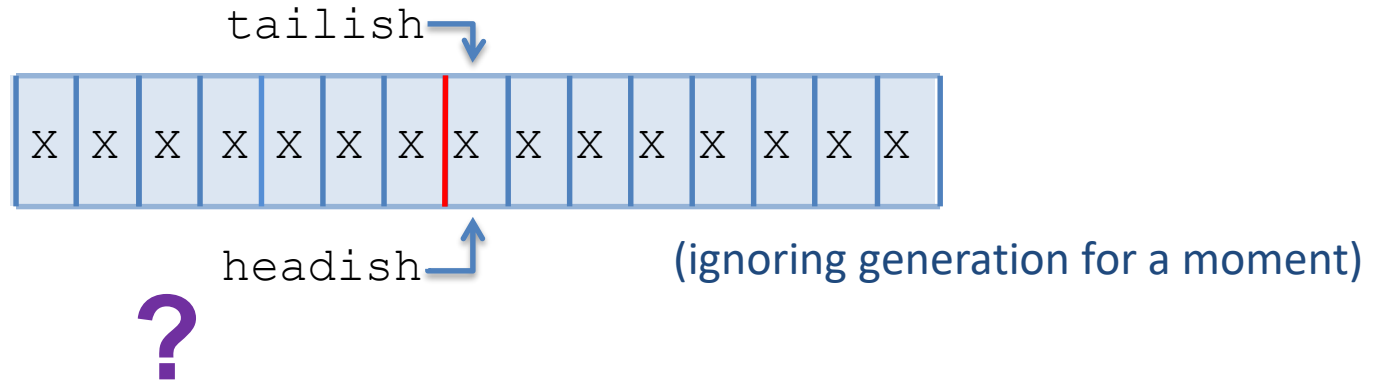
Examples

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



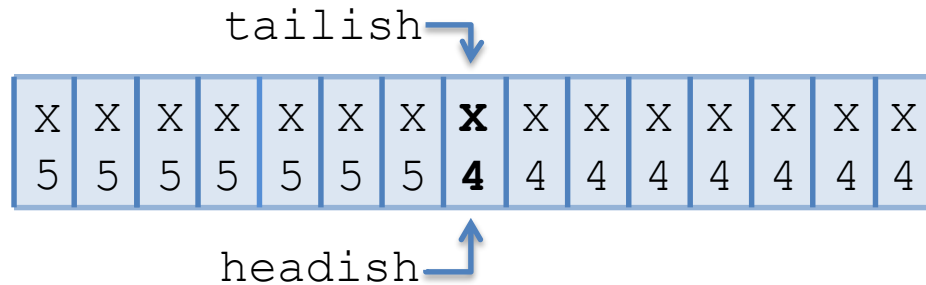
Examples

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



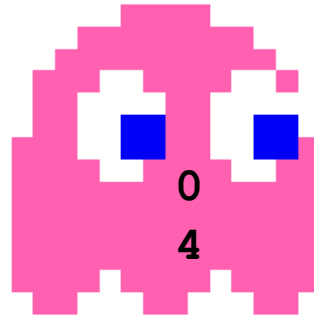
Examples

```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



Queue is Full!

Examples



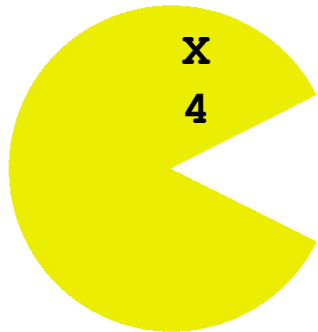
```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```

tailish

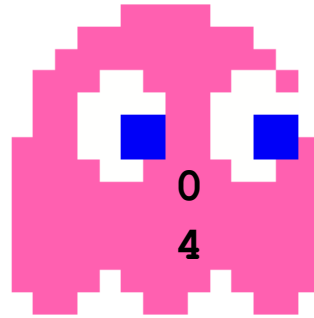


X	X	X	X	X	X	X	X	X	X	X	X	X	X	X
5	5	5	5	5	5	5	4	4	4	4	4	4	4	4

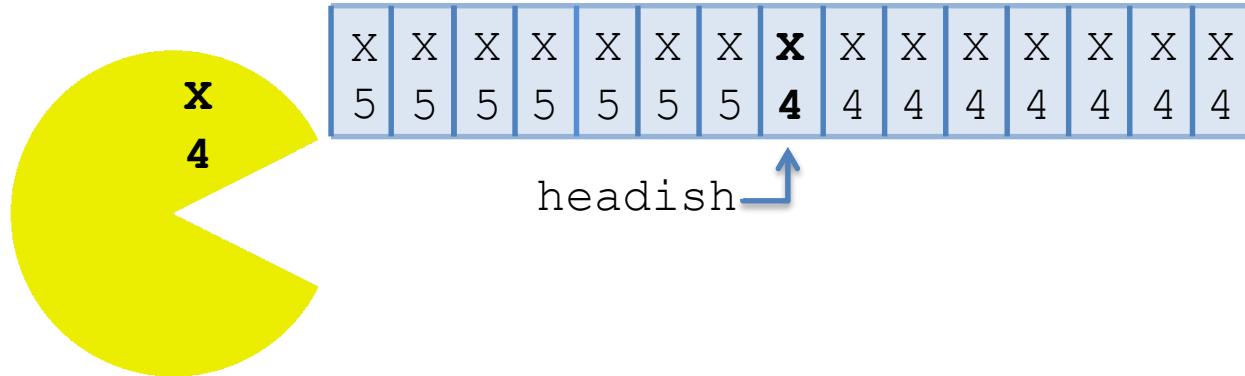
headish



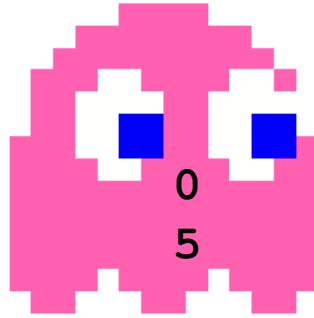
Examples



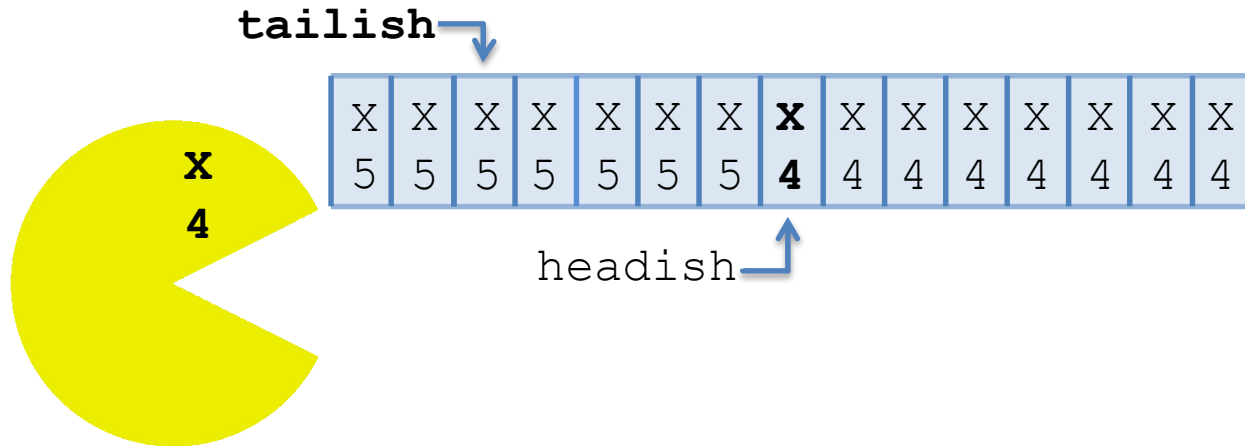
```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```



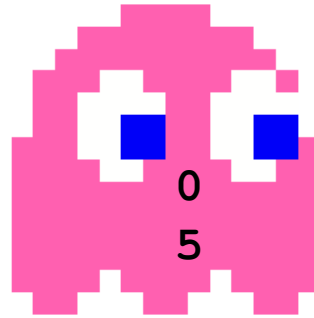
Examples



```
class Queue {
    entry buffer[SIZE];
    int headish;
    int tailish;
    int generation;
};
```



Examples



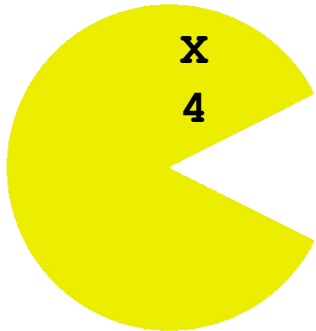
```
class Queue {  
    entry buffer[SIZE];  
    int headish;  
    int tailish;  
    int generation;  
};
```

tailish → 4 < 5

X	X	X	X	X	X	X	X	X	X	X	X	X	X	X
5	5	5	5	5	5	5	4	4	4	4	4	4	4	4

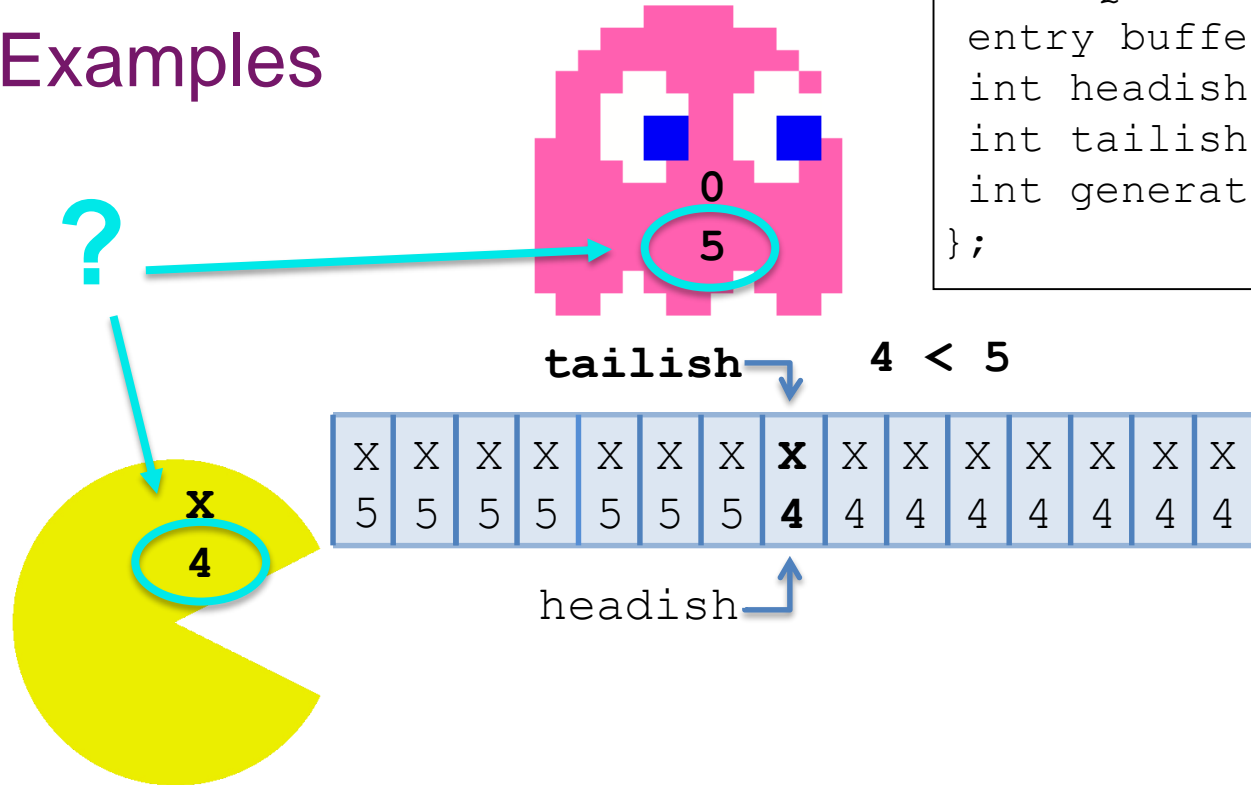
headish ↗

Queue is Full!



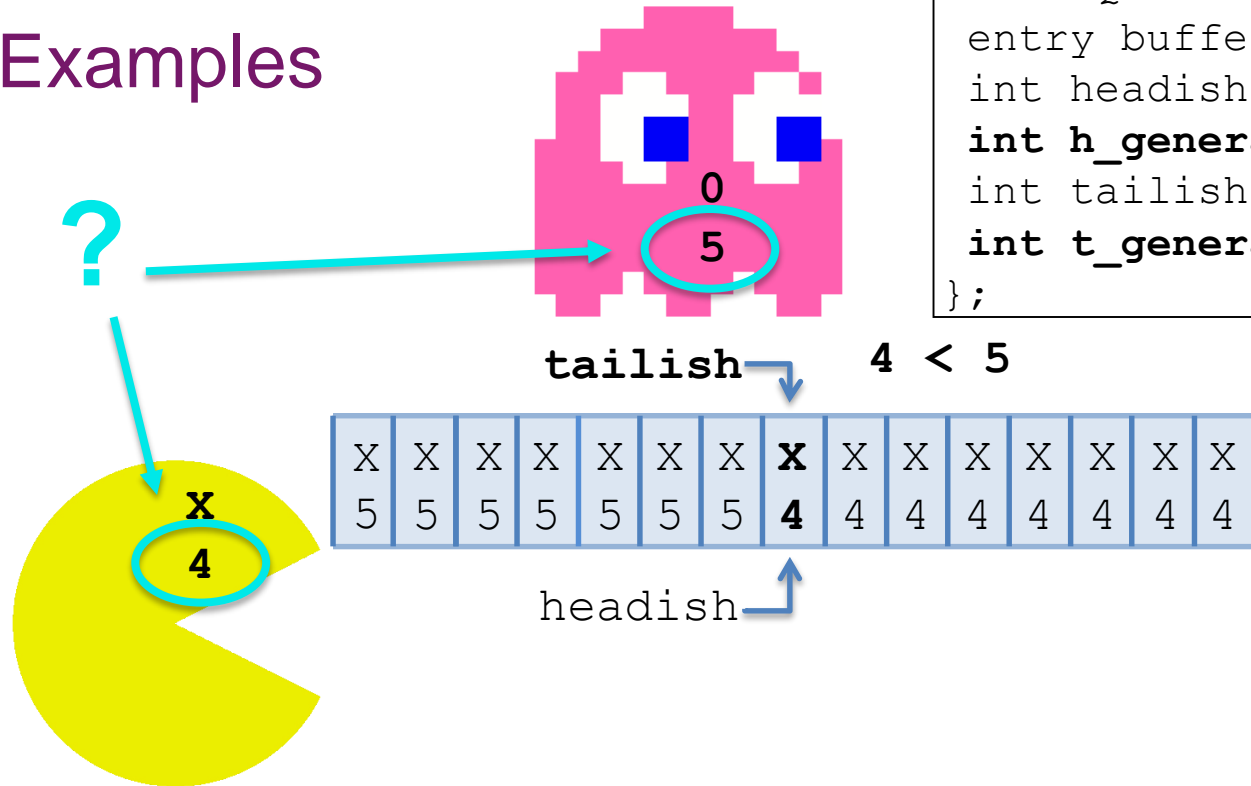
Examples

```
class Queue {
    entry buffer[SIZE];
    int headish;
    int tailish;
    int generation;
};
```



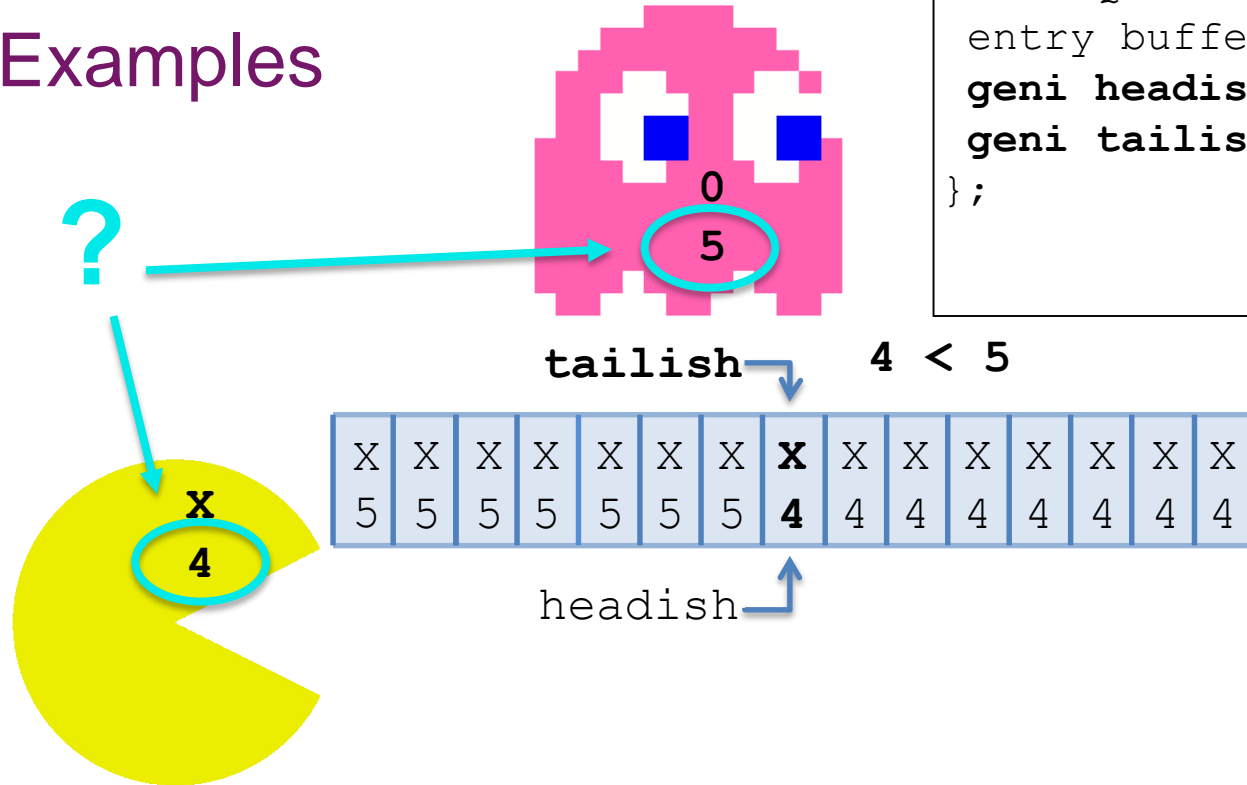
Queue is Full!

Examples



```
class Queue {
    entry buffer[SIZE];
    int headish;
    int h_generation;
    int tailish;
    int t_generation;
};
```

Examples

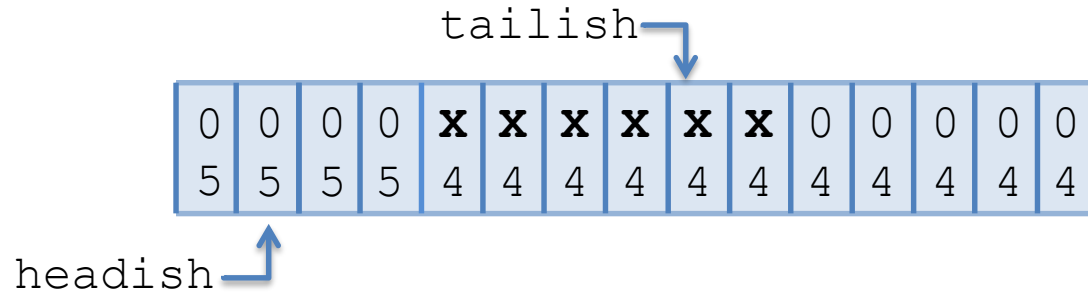


```
class Queue {
    entry buffer[SIZE];
    geni headish;
    geni tailish;
};
```

```
struct geni {
    int val, gen;
    void incr()
    { if (++val%SIZE == 0)
      {   val = 0;
          gen++;
        }
    }
    operator int() {
        return val; }
    operator<() = default;
    bool is_data(int g)
    { return val != 0
      && gen == g; }
    bool is_zero(int g)
    { return val == 0
      && gen == g; }
};
```

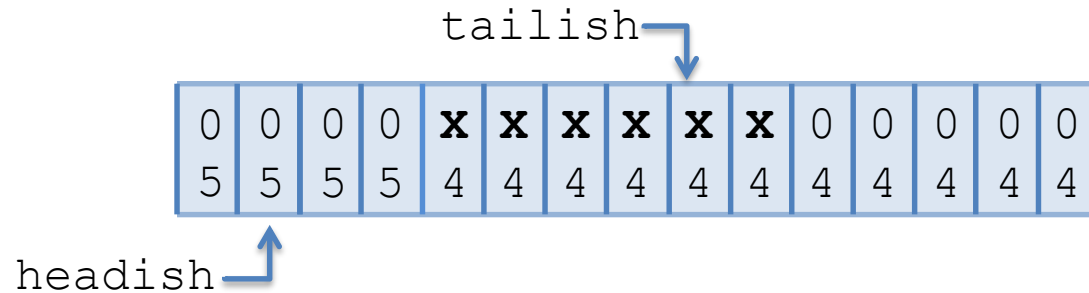
Atomicization!

```
class Queue {
    entry buffer[SIZE];
    geni headish;
    geni tailish;
};
```



Atomicization

```
class Queue {
    atomic<entry> buffer[SIZE];
    geni headish;
    geni tailish;
};
```

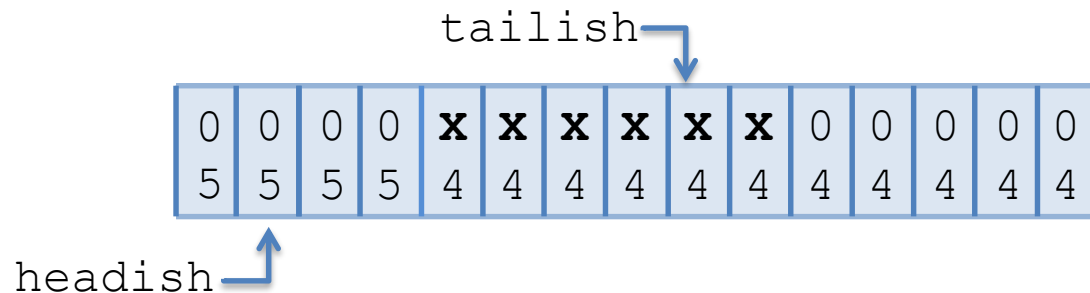


Atomicization

small enough

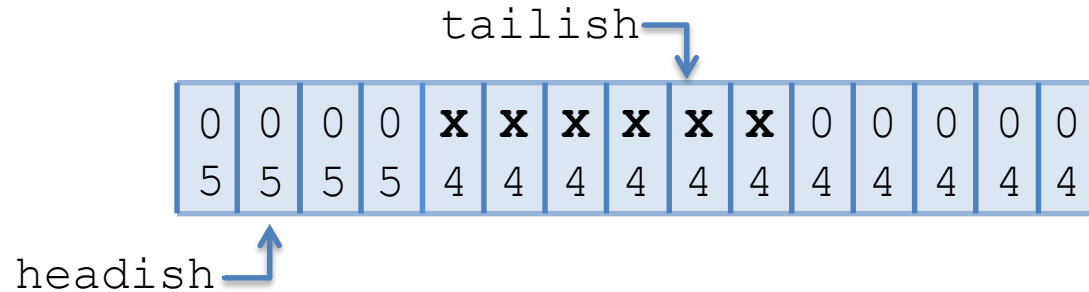
```
struct entry {
    int data;
    int generation;
};
```

```
class Queue {
    atomic<entry> buffer[SIZE];
    geni headish;
    geni tailish;
};
```



Atomicization

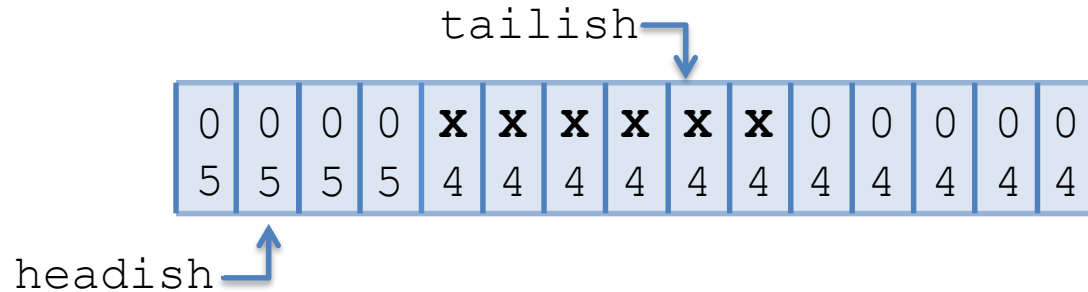
```
class Queue {  
    atomic<entry> buffer[SIZE];  
    laxtomic<geni> headish;  
    laxtomic<geni> tailish;  
};
```



Atomicization



```
class Queue {  
    atomic<entry> buffer[SIZE];  
    laxtomic<geni> headish;  
    laxtomic<geni> tailish;  
};
```



Atomicization



```
class Queue {
    atomic<entry> buffer[SIZE];
    laxtomic<geni> headish;
    laxtomic<geni> tailish;
};
```

Sorry Herb...

[illegible]

headish



Atomicization

```
class Queue {  
    atomic<entry> buffer[SIZE];  
    laxtomic<geni> headish;  
    laxtomic<geni> tailish;  
};
```

```
template <typename T> struct laxtomic : atomic<T> {  
    ...function name...( ..., memory_order = memory_order_relaxed) {...}  
};
```

0	0	0	0	X	X	X	X	X	X	0	0	0	0	0
5	5	5	5	4	4	4	4	4	4	4	4	4	4	4

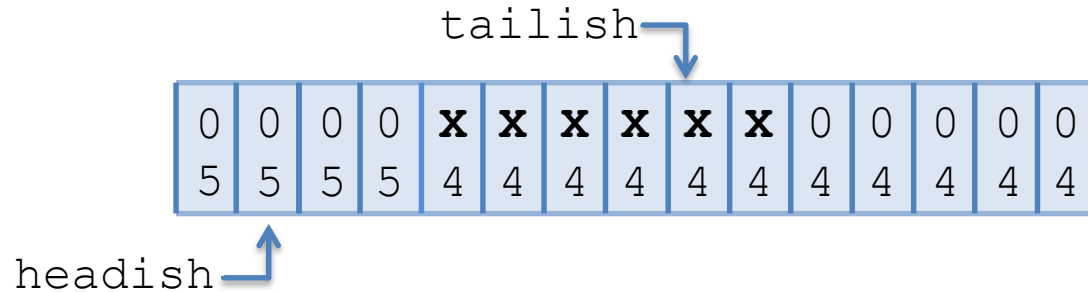
headish



Atomicization

OK?

```
class Queue {  
    atomic<entry> buffer[SIZE];  
    laxatomic<geni> headish;  
    laxatomic<geni> tailish;  
};
```



Atomicization

OK?

“normal” lock free code

```
data.x = 10;  
data.y = 20;  
data.ready.store(true, release);
```

X

4

```
class Queue {  
    atomic<entry> buffer[SIZE];  
    latomic<geni> headish;  
    geni> tailish;
```

```
if (data.ready.load(acquire))  
{  
    x = data.x;  
    y = data.y;  
    ...  
}
```

Atomicization

OK?

“normal” lock free code

```
class Queue {  
    atomic<entry> buffer[SIZE];  
    latomic<geni> headish;  
    geni> tailish;
```

```
data.x = 10;  
data.y = 20;  
data.ready.store(true, release);
```

X
4

```
if (data.ready.load(acquire))  
{  
    x = data.x;  
    y = data.y;  
    ...  
}
```

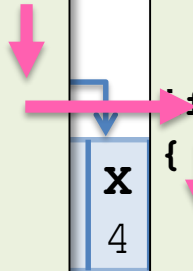
Atomicization

OK?

“normal” lock free code

```
class Queue {
    atomic<entry> buffer[SIZE];
    latomic<geni> headish;
    geni> tailish;
```

```
data.x = 10;
data.y = 20;
data.ready.store(true, release);
data.x = 10; // can't happen!
```



```
x = data.x; // can't happen!
if (data.ready.load(acquire))
{
    x = data.x;
    y = data.y;
    ...
}
```


Atomicization

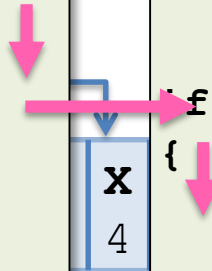
OK?

“normal” lock free code

```
class Queue {
    atomic<entry> buffer[SIZE];
    latomic<geni> headish;
    geni> tailish;
```

```
data.x = 10;
data.y = 20;
data.ready.store(true, release);
data.x = 10; // can't happen!
```

release:
before means before



```
x = data.x; // can't happen!
if (data.ready.load(acquire))
{
    x = data.x;
    y = data.y;
    ...
}
```

acquire:
after means after

Atomicization

OK?

relaxed lock free code

```
class Queue {  
    atomic<entry> buffer[SIZE];  
    laxtomic<geni> headish;  
    geni> tailish;
```

```
data.x = 10;  
data.y = 20;  
data.ready.store(true, mo_relax);  
data.x = 10; // CAN happen!
```



x
4

```
    x = data.x; // CAN happen!  
    if (data.ready.load(mo_relax))  
    {  
        x = data.x;  
        y = data.y;  
        ...  
    }
```



Atomicization

OK?

“normal” lock free code

```
class Queue {  
    atomic<entry> buffer[SIZE];  
    latomic<geni> headish;  
    geni> tailish;
```

```
data.x = 10;  
data.y = 20;  
data.ready.store(true, release);  
data.x = 10; // can't happen!
```

x
4

```
    x = data.x; // can't happen!  
    if (data.ready.load(acquire))  
    {  
        x = data.x;  
        y = data.y;  
        ...
```

there is a *relationship* between
ready and the other data

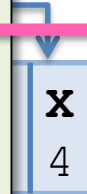
Atomicization

OK?

relaxed lock free code

```
class Queue {  
    atomic<entry> buffer[SIZE];  
    laxtomic<geni> headish;  
    geni> tailish;
```

```
data.x = 10;  
data.y = 20;  
data.ready.store(true, mo_relax);  
data.x = 10; // CAN happen!
```



```
    x = data.x; // CAN happen!  
    if (data.ready.load(mo_relax))  
    {  
        x = data.x;  
        y = data.y;  
        ...
```

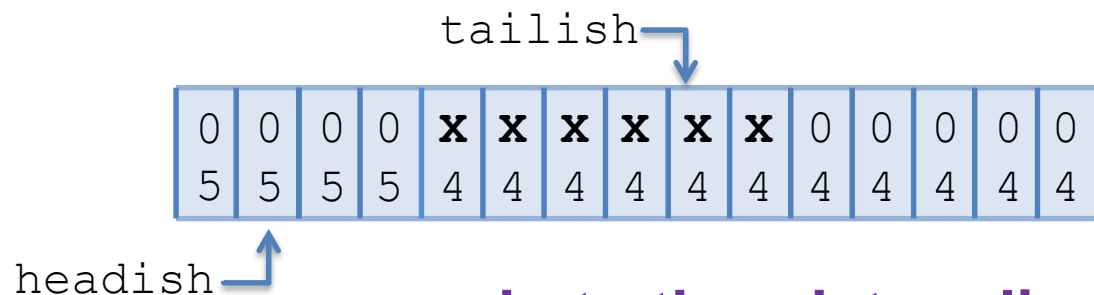
there is **no relationship** between
ready and the other data



Atomicization

OK?

```
class Queue {  
    atomic<entry> buffer[SIZE];  
    laxatomic<geni> headish;  
    laxatomic<geni> tailish;  
};
```

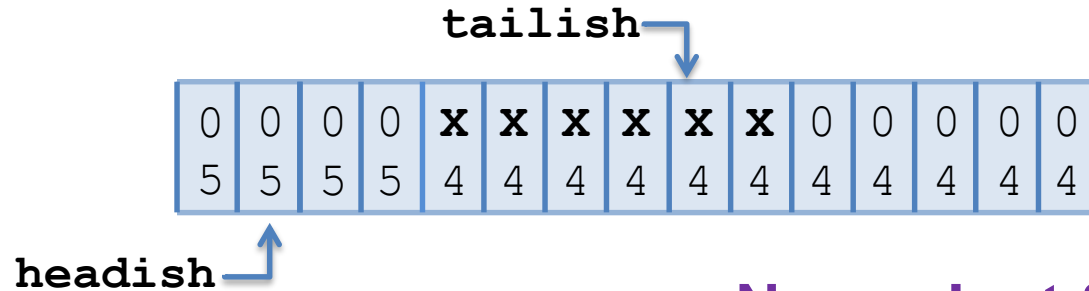


what other data relies on head/tail?

Atomicization

OK

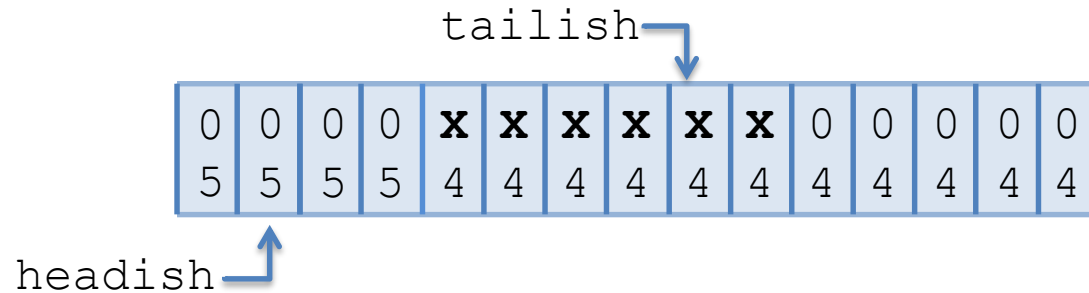
```
class Queue {  
    atomic<entry> buffer[SIZE];  
    laxatomic<geni> headish;  
    laxatomic<geni> tailish;  
};
```



None. Just 'hints'.

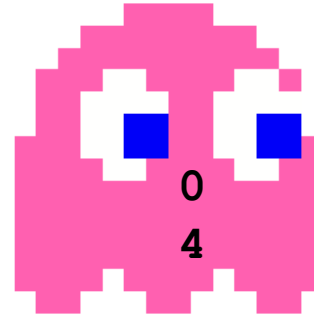
Atomicization

```
class Queue {
    atomic<entry> buffer[SIZE];
    laxtomic<geni> headish;
    laxtomic<geni> tailish;
};
```



Codization

```
void push(int val) {  
    ...  
}
```



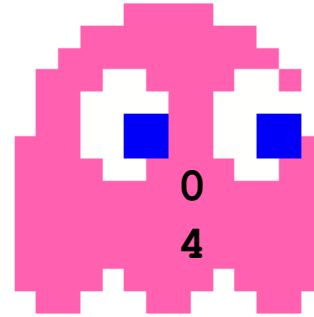
tailish

0	0	0	0	x	x	x	x	x	x	0	0	0	0	0
5	5	5	5	4	4	4	4	4	4	4	4	4	4	4

headish

Codization

```
void push(int val) {  
  ...  
}
```



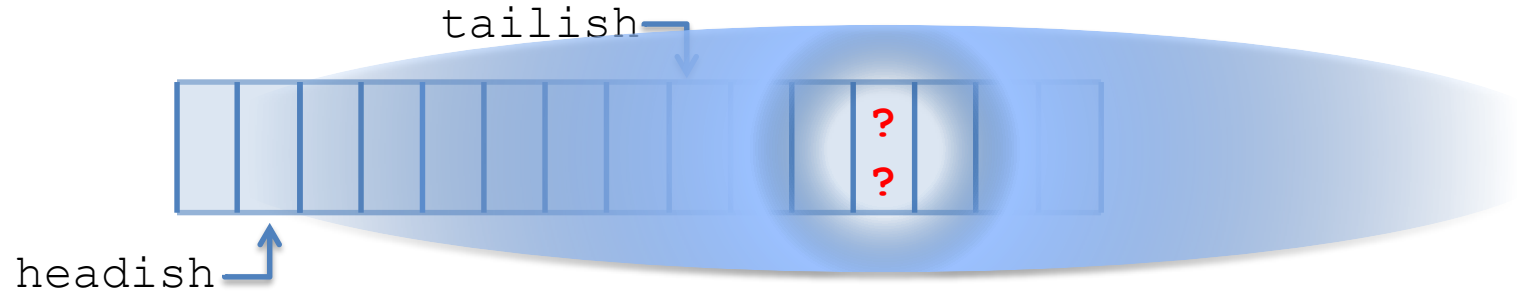
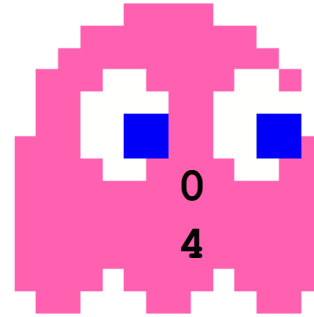
tailish



headish

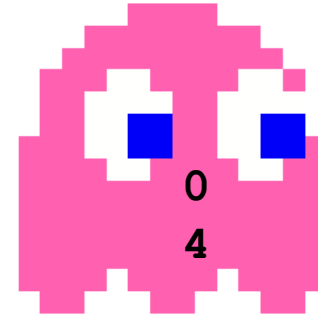
Codization

```
void push(int val) {  
  ...  
}
```



Codization

```
void push(int val) {  
    ...  
}
```

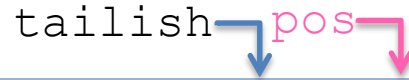


tailish

0	0	0	0	x	x	x	x	x	x	0	0	0	0	0
5	5	5	5	4	4	4	4	4	4	4	4	4	4	4

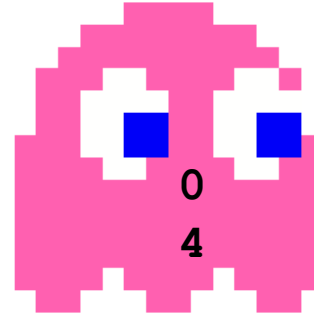
headish



```
void push(int val) {
    geni pos = find_tail(tailish);
    write_value(pos, val);
}
```

[illegible]

Codization

```
void push(int val) {  
    geni pos = find_tail(tailish);  
    write_value(pos, val);  
}
```

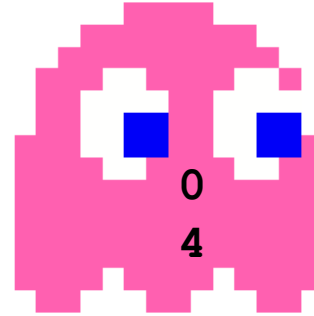




tailish   pos



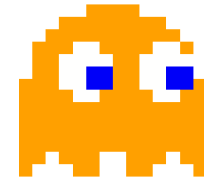
Codization

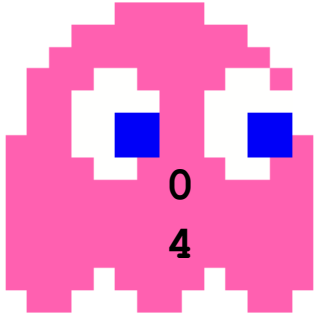
```
void push(int val) {  
    geni pos = find_tail(tailish);  
    write_value(pos, val);  
}
```



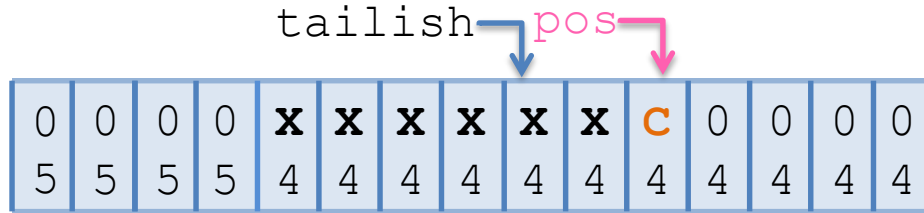
tailish   pos

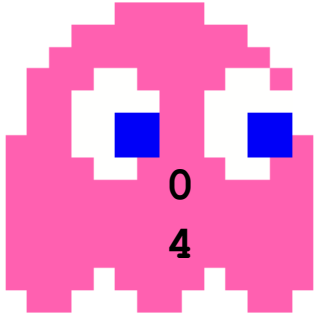
0	0	0	0	x	x	x	x	x	x	c	0	0	0	0
5	5	5	5	4	4	4	4	4	4	4	4	4	4	4



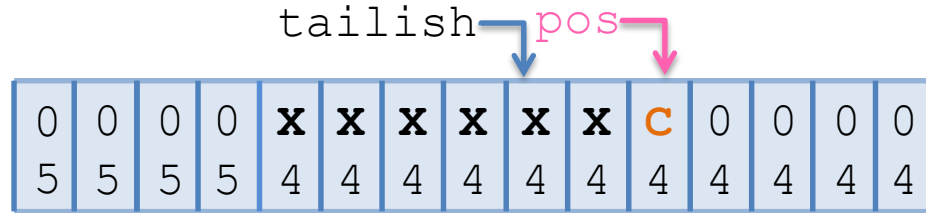


```
void push(int val) {  
  do {  
    pos = find_tail(tailish);  
  } while (!try_write_value(pos, val));  
}
```

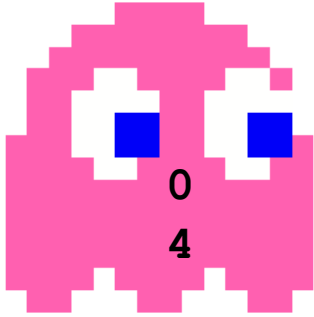




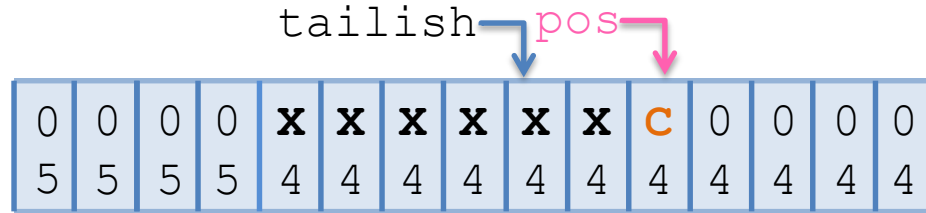
```
void push(int val) {  
    do {  
        pos = find_tail(tailish);  
    } while (!try_write_value(pos, val));  
}
```



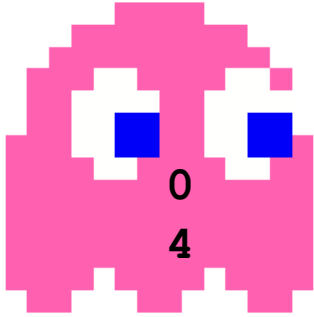
The nature of “lock-free” – you fail **only** when someone else makes progress



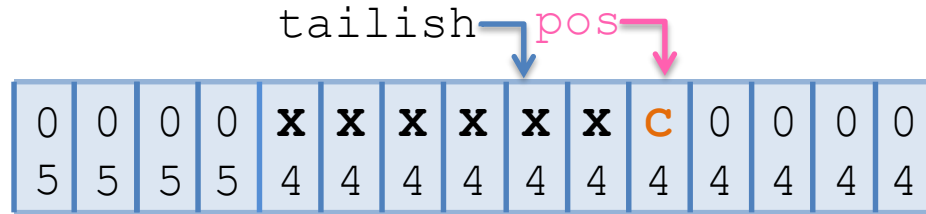
```
void push(int val) {  
    do {  
        pos = find_tail(tailish);  
    } while (!try_write_value(pos, val));  
}
```

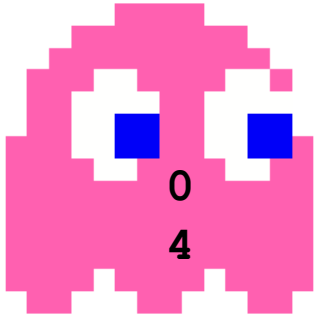


CAS loop: loops on progress
Spin-lock: loops on progress and NON progress

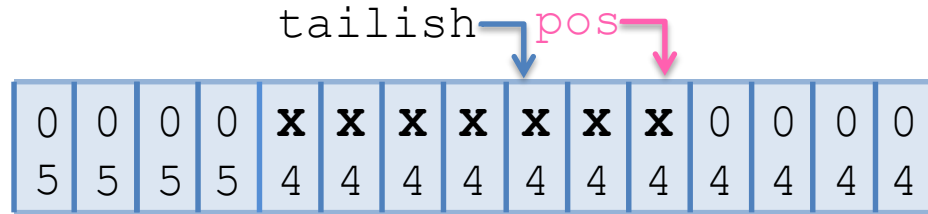


```
void push(int val) {  
    geni pos = tailish; // relaxed load  
    do {  
        pos = find_tail(pos);  
    } while (!try_write_value(pos, val));  
}
```



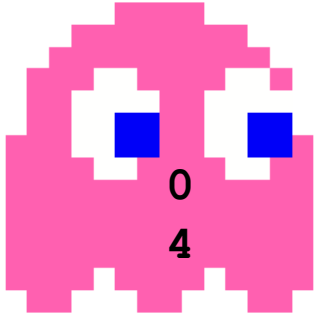


```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```

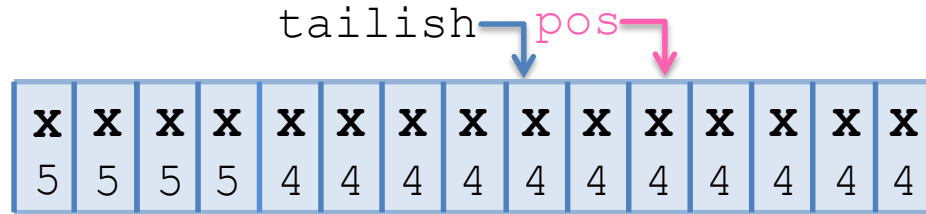


```
bool is_zero(entry e, int gen) {
    return e.data == 0 && e.gen == gen;
}
```

```
geni find_tail(geni pos) {
    while( ! is_zero(buffer[pos.val].load(relaxed), pos.gen) )
        pos++;
    return pos;
}
```



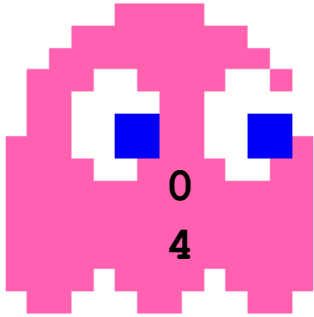
```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



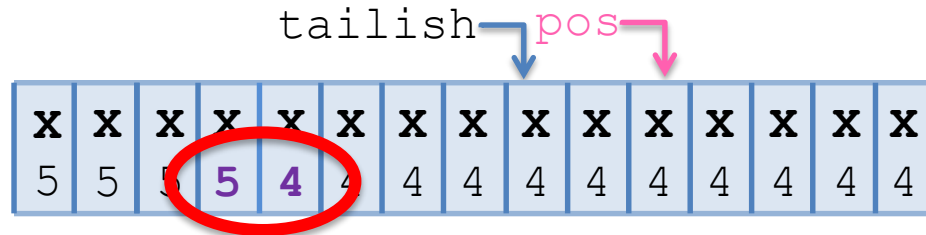
Queue is FULL

```
bool is_zero(entry e, int gen) {
    return e.data == 0 && e.gen == gen;
}
```

```
geni find_tail(geni pos) {
    while( ! is_zero(buffer[pos.val].load(relaxed), pos.gen) )
        pos++;
    return pos;
}
```



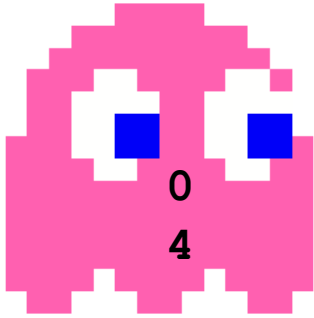
```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



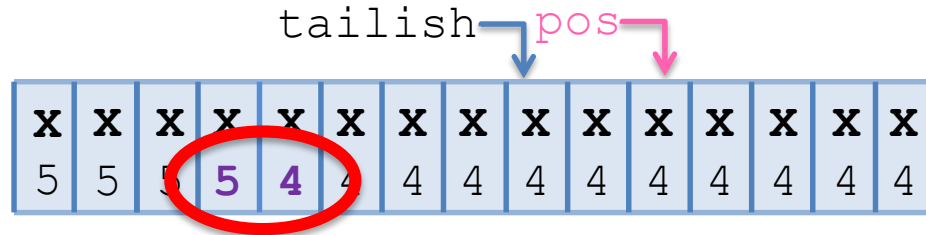
Queue is FULL

```
bool is_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
    || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) {
    while( !is_tail(buffer[pos.val].load(relaxed), pos.gen) )
        pos++;
    return pos;
}
```



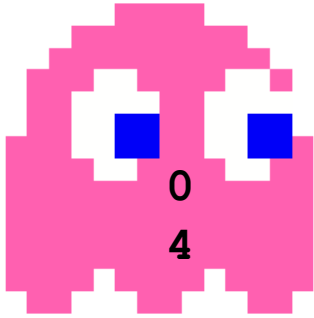
```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



Queue is FULL

```
bool is_tail(entry e, int gen) {
    return e.data == 0 ?
        e.gen == gen : e.gen < gen;
}
```

```
geni find_tail(geni pos) {
    while( !is_tail(buffer[pos.val].load(relaxed), pos.gen) )
        pos++;
    return pos;
}
```



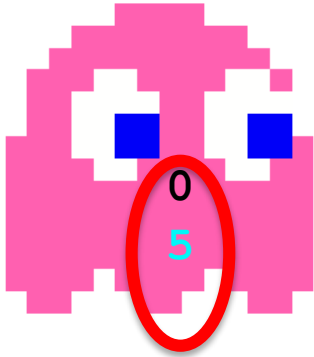
```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



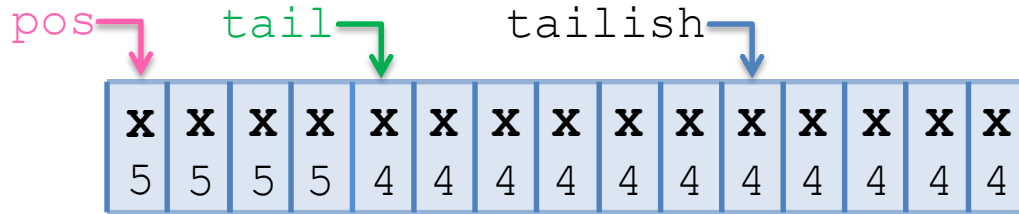
Queue is FULL

```
bool is_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) {
    while( !is_tail(buffer[pos.val].load(relaxed), pos.gen) )
        pos++;
    return pos;
}
```

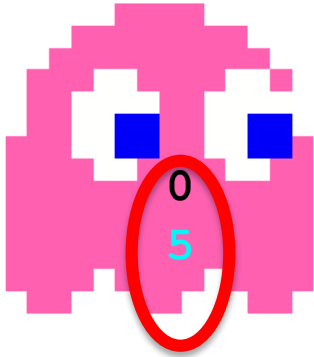
```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



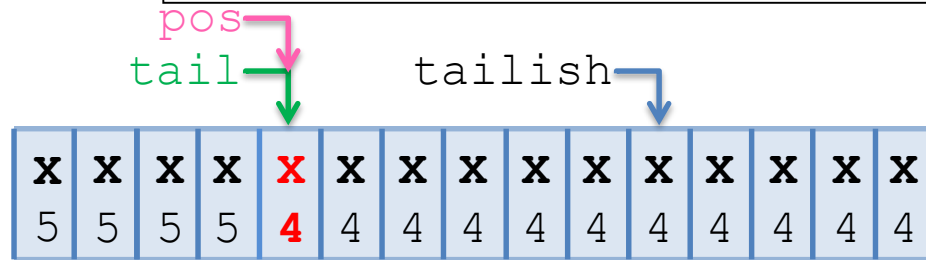
Queue is FULL

```
bool is_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
    || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) {
    while( !is_tail(buffer[pos.val].load(relaxed), pos.gen) )
        pos++;
    return pos;
}
```



```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



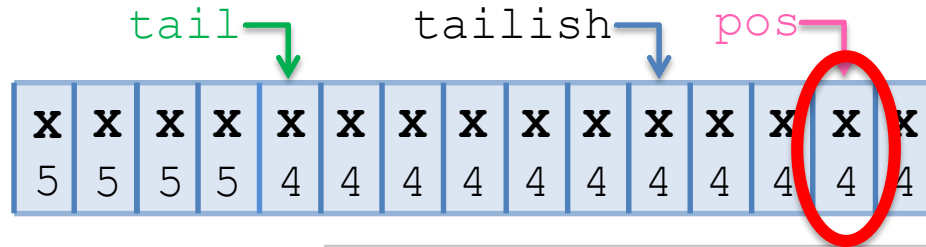
Queue is FULL

```
bool is_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
    || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) {
    while( !is_tail(buffer[pos.val].load(relaxed), pos.gen) )
        pos++;
    return pos;
}
```



```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



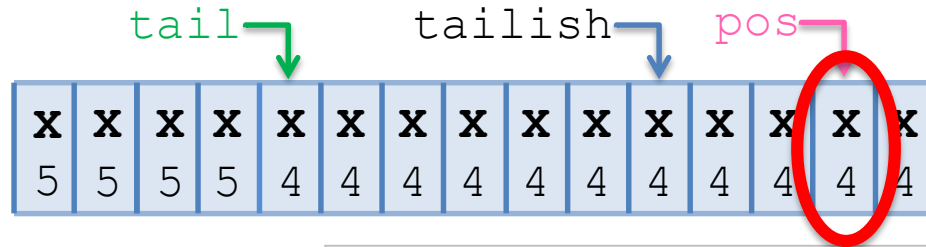
Queue is FULL

```
bool is_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) {
    while( ! is_tail(buffer[pos.val].load(relaxed), pos.gen) )
        pos++;
    return pos;
}
```



```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



Queue is FULL

```
bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) {
    while(!maybe_tail(buffer[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}
```



INVARIANT:
tailish <= real tail

```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



Queue is FULL

```
bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
    || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) { // precondition: pos <= tail
    while(!maybe_tail(buffer[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}
```



```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```

Exercise for the reader:

Don't loop on FULL. Either:

- have push() return false
- *wait* for non full.



Queue is FULL

```
bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) { // precondition: pos <= tail
    while(!maybe_tail(buffer[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}
```



```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos, full);
        if (full) ...
    } while (!try_write_value(pos, val));
}
```

Exercise for the reader:

Don't loop on FULL. Either:

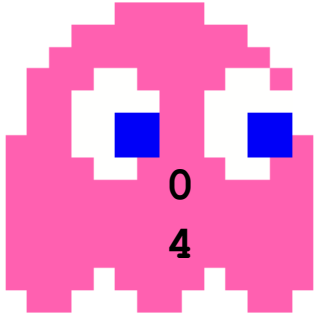
- have push() return false
- *wait* for non full.



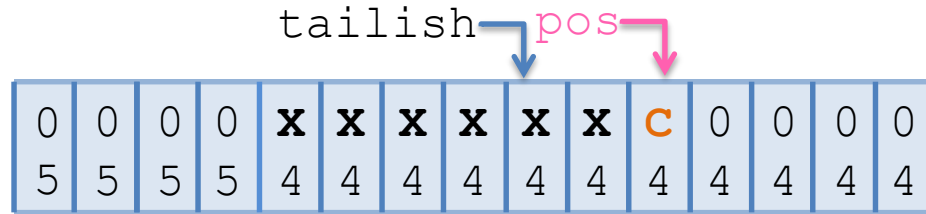
Queue is FULL

```
bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}
```

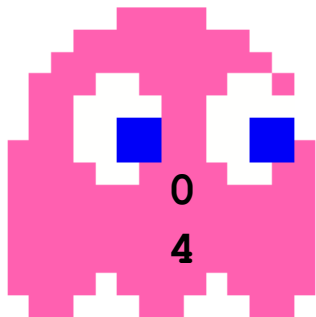
```
geni find_tail(geni pos) { // precondition: pos <= tail
    while(!maybe_tail(buffer[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}
```

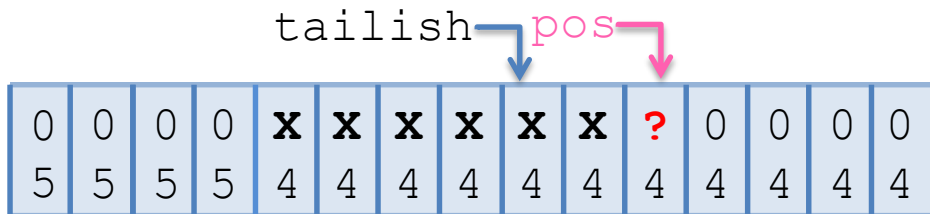
```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



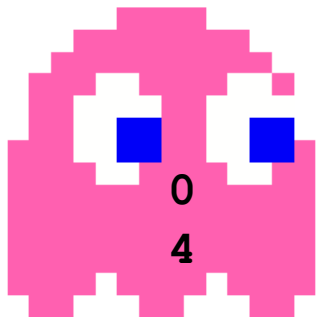
```
bool try_write_value(geni pos, int val) {
    // ...
}
```



```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```



```
bool try_write_value(geni pos, int val) {
    entry old{0, pos.gen};
    entry nu{val, pos.gen};
    return buffer[pos].compare_exchange_weak(old, nu, release, relaxed);
}
```

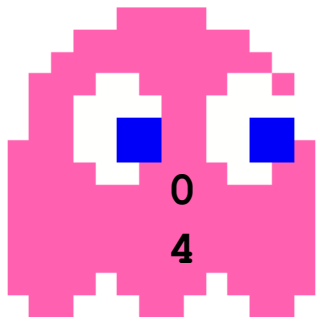


```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
}
```

```
bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) { // precondition: pos <= tail
    while(!maybe_tail(buffer[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}
```

```
bool try_write_value(geni pos, int val) {
    entry old{0, pos.gen};
    entry nu{val, pos.gen};
    return buffer[pos].compare_exchange_weak(old, nu, release, relaxed);
}
```

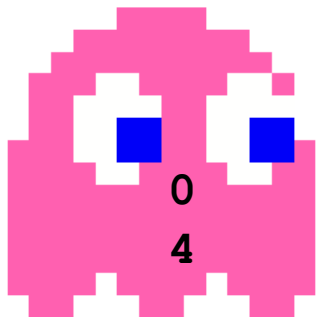


```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
    // ???
}
```

```
bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
    || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) { // precondition: pos <= tail
    while(!maybe_tail(buffer[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}
```

```
bool try_write_value(geni pos, int val) {
    entry old{0, pos.gen};
    entry nu{val, pos.gen};
    return buffer[pos].compare_exchange_weak(old, nu, release, relaxed);
}
```

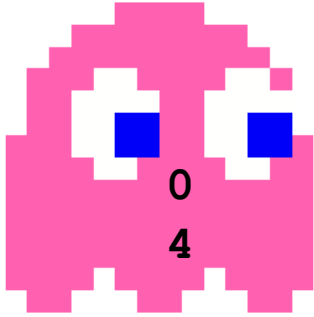


```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
    update_tailish(tailish, pos);
}
```

```
bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}
```

```
geni find_tail(geni pos) { // precondition: pos <= tail
    while(!maybe_tail(buffer[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}
```

```
bool try_write_value(geni pos, int val) {
    entry old{0, pos.gen};
    entry nu{val, pos.gen};
    return buffer[pos].compare_exchange_weak(old, nu, release, relaxed);
}
```

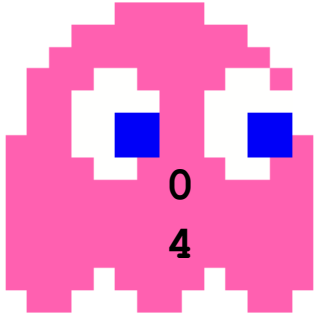



tailish

```
void push(int val) {  
    geni pos = tailish; // relaxed load  
    do {  
        pos = find_tail(pos);  
    } while (!try_write_value(pos, val));  
    update_tailish(tailish, ++pos);  
}
```

pos





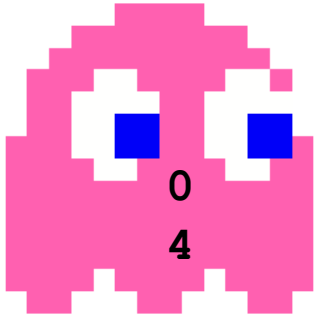
```
void push(int val) {  
    geni pos = tailish; // relaxed load  
    do {  
        pos = find_tail(pos);  
    } while (!try_write_value(pos, val));  
    update_tailish(tailish, ++pos);  
}
```

tailish

pos



```
tailish.compare_exchange(oldtailish, pos); // relaxed (in theory)
```

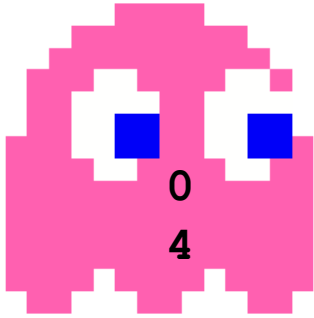
```
void push(int val) {  
    geni pos = tailish; // relaxed load  
    do {  
        pos = find_tail(pos);  
    } while (!try_write_value(pos, val));  
    update_tailish(tailish, ++pos);  
}
```

tailish

pos



```
tailish.compare_exchange(oldtailish, pos); // relaxed (in theory)  
---  
geni recent = tailish;  
while (recent < pos && !tailish.compare_exchange(recent, pos)) //(recent is updated each loop)  
    ;
```



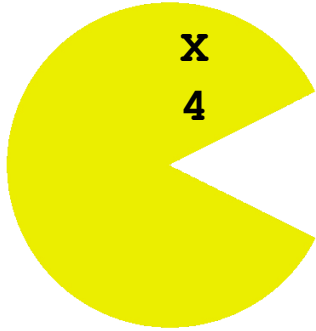
```
void push(int val) {
    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
    } while (!try_write_value(pos, val));
    update_tailish(tailish, ++pos);
}
```

tailish

pos



```
tailish.compare_exchange(oldtailish, pos); // relaxed (in theory)
---
geni recent = tailish;
while (recent < pos && !tailish.compare_exchange(recent, pos)) //(recent is updated each loop)
    ;
---
tailish = pos; // relaxed (from Sebastian Redl)
```

```
int pop() {
```

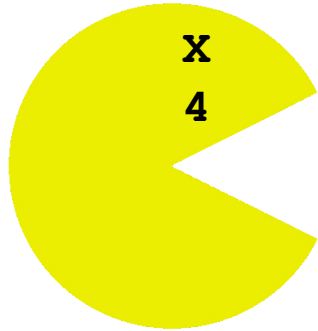
```
}
```

x	x	0	0	0	0	0	0	0	0	x	x	x	x	x
5	5	5	5	5	5	5	5	5	5	4	4	4	4	4

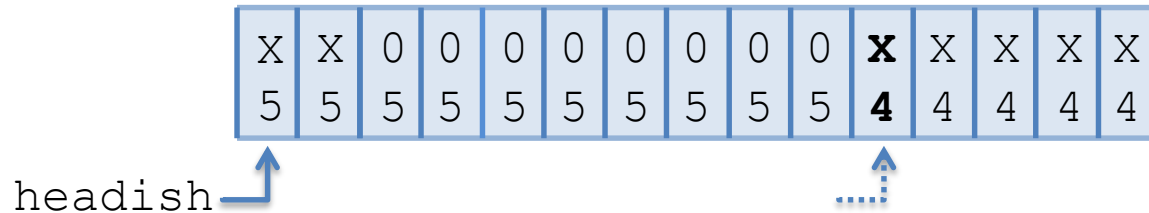
headish 



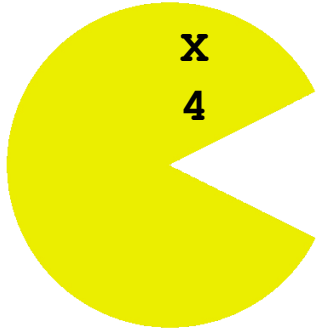
First non-zero data **of correct generation**



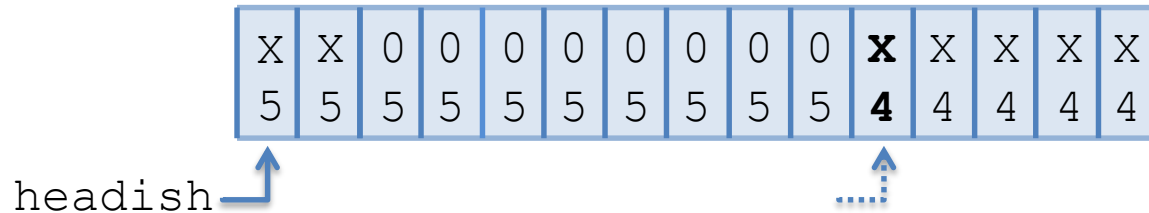
```
void push(int val) {  
    geni pos = tailish; // relaxed load  
    do {  
        pos = find_tail(pos);  
    } while (!try_write_value(pos, val));  
    tailish = pos+1; // thanks Sebastian  
}
```



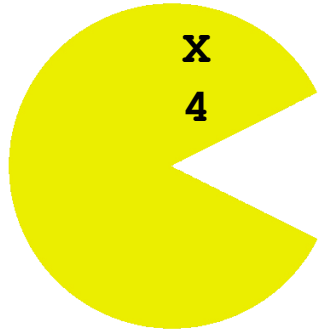
First non-zero data **of correct generation**



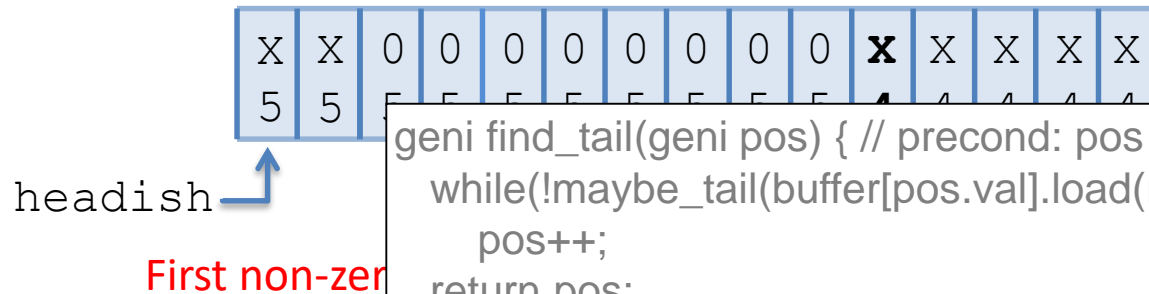
```
int pop() {  
    entry ent;  
    geni pos = headish; // relaxed load  
    do {  
        pos = find_head(pos, ent);  
    } while (!try_remove_value(pos, ent));  
    headish = pos+1; // thanks Sebastian  
    return ent.val;  
}
```



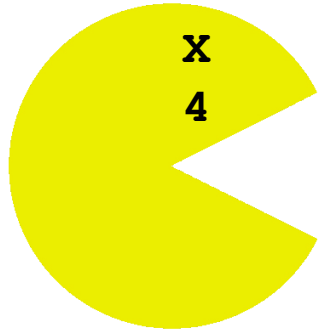
First non-zero data of correct generation



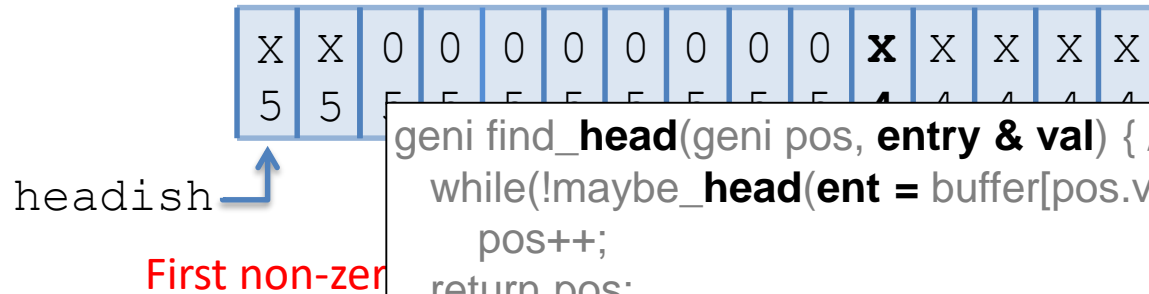
```
int pop() {
  entry ent;
  geni pos = headish; // relaxed load
  do {
    pos = find_head(pos, ent);
  } while (!try_remove_value(pos, ent));
  headish = pos+1; // thanks Sebastian
  return ent.val;
}
```



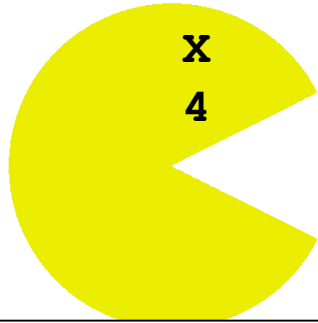
```
geni find_tail(geni pos) { // precondition: pos <= tail
  while(!maybe_tail(buffer[pos.val].load(relaxed), pos.gen))
    pos++;
  return pos;
}
```



```
int pop() {
  entry ent;
  geni pos = headish; // relaxed load
  do {
    pos = find_head(pos, ent);
  } while (!try_remove_value(pos, ent));
  headish = pos+1; // thanks Sebastian
  return ent.val;
}
```



```
geni find_head(geni pos, entry & val) { // precondition: pos <= head
  while(!maybe_head(ent = buffer[pos.val].load(relaxed), pos.gen))
    pos++;
  return pos;
}
```

```
bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}
```

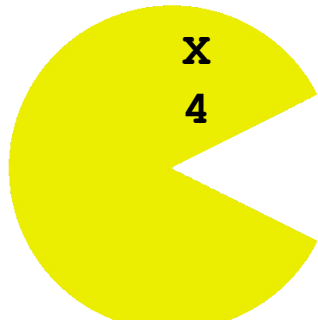
```
int pop() {
    entry ent;
    geni pos = headish; // relaxed load
    do {
        pos = find_head(pos, ent);
    } while (!try_remove_value(pos, ent));
    headish = pos+1; // thanks Sebastian
    return ent.val;
}
```

X	X	0	0	0	0	0	0	0	0	X	X	X	X	X
5	5	5	5	5	5	5	5	5	5	4	4	4	4	4

headish →

First non-zero

```
geni find_head(geni pos, entry & val) { // precondition: pos <= head
    while(!maybe_head(ent = buffer[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}
```



```
bool maybe_head(entry e, int gen) {
    return e.data != 0 && e.gen == gen
        || e.data == 0 && e.gen < gen;
}
```

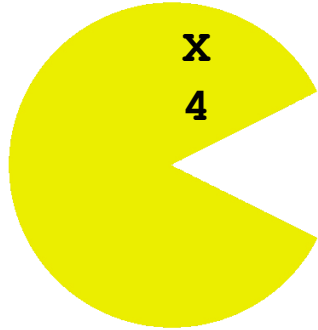
```
int pop() {
    entry ent;
    geni pos = headish; // relaxed load
    do {
        pos = find_head(pos, ent);
    } while (!try_remove_value(pos, ent));
    headish = pos+1; // thanks Sebastian
    return ent.val;
}
```

X	X	0	0	0	0	0	0	0	0	X	X	X	X	X
5	5	5	5	5	5	5	5	5	5	4	4	4	4	4

headish →

First non-zero

```
geni find_head(geni pos, entry & val) { // precondition: pos <= head
    while(!maybe_head(ent = buffer[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}
```



```
int pop() {
  entry ent;
  geni pos = headish; // relaxed load
  do {
    pos = find_head(pos, ent);
  } while (!try_remove_value(pos, ent));
  headish = pos+1; // thanks Sebastian
  return ent.val;
}
```

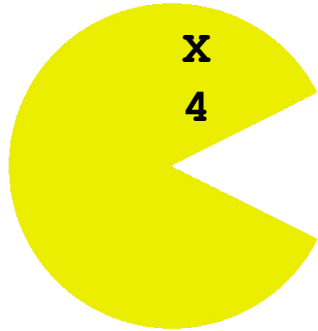
X	X	0	0	0	0	0	0	0	0	x	X	X	X	X
5	5	5	5	5	5	5	5	5	5	4	4	4	4	4

headish

First

```
bool try_remove_value(geni pos, entry old) {

  entry nu{0, pos.gen};
  return buffer[pos].compare_exchange_weak(old, nu, acquire, relaxed);
}
```



```
int pop() {
    entry ent;
    geni pos = headish; // relaxed load
    do {
        pos = find_head(pos, ent);
    } while (!try_remove_value(pos, ent));
    headish = pos+1; // thanks Sebastian
    return ent.val;
}
```

X	X	0	0	0	0	0	0	0	0	X	X	X	X	X
5	5	5	5	5	5	5	5	5	5	4	4	4	4	4

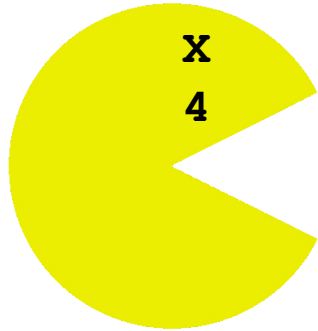
headish =

First

```
bool try_remove_value(geni pos, entry old) {
```

```
    entry nu{0, pos.gen};
    return buffer[pos].compare_exchange_weak(old, nu, acquire, relaxed);
}
```





```
int pop() {
    entry ent;
    geni pos = headish; // relaxed load
    do {
        pos = find_head(pos, ent);
    } while (!try_remove_value(pos, ent));
    headish = pos+1; // thanks Sebastian
    return ent.val;
}
```

X	X	0	0	0	0	0	0	0	0	X	X	X	X	X
5	5	5	5	5	5	5	5	5	5	4	4	4	4	4

headish

First

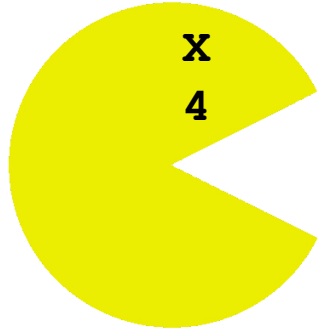
```
bool try_remove_value(geni pos, entry old) {
```

```
    entry nu{0, pos.gen};
```

```
    return buffer[pos].compare_exchange_weak(old, nu, acq_rel, relaxed);
```

```
}
```

?



```
int pop() {
    entry ent;
    geni pos = headish; // relaxed load
    do {
        pos = find_head(pos, ent);
    } while (!try_remove_value(pos, ent));
    headish = pos+1; // thanks Sebastian
    return ent.val;
}
```

X	X	0	0	0	0	0	0	0	0	X	X	X	X	X
5	5	5	5	5	5	5	5	5	5	4	4	4	4	4

Exercise for the reader

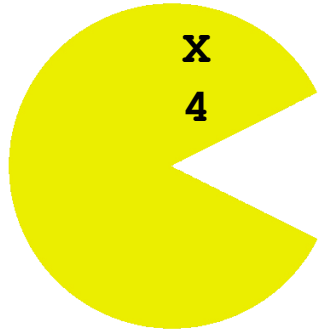
headish =

First

```
bool try_remove_value(geni pos, entry old) {

    entry nu{0, pos.gen};
    return buffer[pos].compare_exchange_weak(old, nu, acq_rel, relaxed);
}
```

?



```
int pop() {
    entry ent;
    geni pos = headish; // relaxed lo
    do {
        pos = find_head(pos, ent);
    } while (!try_remove_value(pos,
    headish = pos+1; // thanks Sebastian
    return ent.val;
}
```

Exercise for the reader:

Don't loop on **EMPTY**. Either:

- have **pop()** return **0**
- *wait* for non **empty**.

X	X	0	0	0	0	0	0	0	0	X	X	X	X	X
5	5	5	5	5	5	5	5	5	5	4	4	4	4	4

Exercise for the reader

headish

First

```
bool try_remove_value(geni pos, entry old) {

    entry nu{0, pos.gen};
    return buffer[pos].compare_exchange_weak(old, nu, acq_rel, relaxed);
}
```

?

```

void push(int val) {

    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
        if (FULL) { ... }
    } while (!try_write_value(pos, val));
    tailish = pos+1; // thanks Sebastian

}

geni find_tail(geni pos) { // precondition: pos <= tail
    while(!maybe_tail(buff[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}

bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}

bool try_write_value(geni pos, int val) {
    entry old{0, pos.gen};
    entry nu{val, pos.gen};
    return buffer[pos].c_e_weak(old, nu, release, relaxed);
}

```

```

int pop() {
    entry ent;
    geni pos = headish; // relaxed load
    do {
        pos = find_head(pos, ent);
        if (FULL) { ... }
    } while (!try_remove_value(pos, ent));
    headish = pos+1; // thanks Sebastian
    return ent.val;
}

geni find_head(geni pos, entry & val) { //precondition: pos <= head
    while(!maybe_head(ent = buff[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}

bool maybe_head(entry e, int gen) {
    return e.data != 0 && e.gen == gen
        || e.data == 0 && e.gen < gen;
}

bool try_remove_value(geni pos, entry old) {

    entry nu{0, pos.gen};
    return buffer[pos].c_e_weak(old, nu, acq_rel, relaxed);
}

```



```
void push(int val) {
```

```

    geni pos = tailish; // relaxed load
    do {
        pos = find_tail(pos);
        if (FULL) { ... }
    } while (!try_write_value(pos, val));
    tailish = pos+1; // thanks Sebastian

}

geni find_tail(geni pos) { // precondition: pos <= tail
    while (!maybe_tail(buff[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}

bool maybe_tail(entry e, int gen) {
    return e.data == 0 && e.gen == gen
        || e.data != 0 && e.gen < gen;
}

bool try_write_value(geni pos, int val) {
    entry old{0, pos.gen};
    entry nu{val, pos.gen};
    return buffer[pos].c_e_weak(old, nu, release, relaxed);
}

```

```
int pop() {
```

```

    entry ent;
    geni pos = headish; // relaxed load
    do {
        pos = find_head(pos, ent);
        if (FULL) { ... }
    } while (!try_remove_value(pos, ent));
    headish = pos+1; // thanks Sebastian
    return ent.val;

}

geni find_head(geni pos, entry & val) { // precondition: pos <= head
    while (!maybe_head(ent = buff[pos.val].load(relaxed), pos.gen))
        pos++;
    return pos;
}

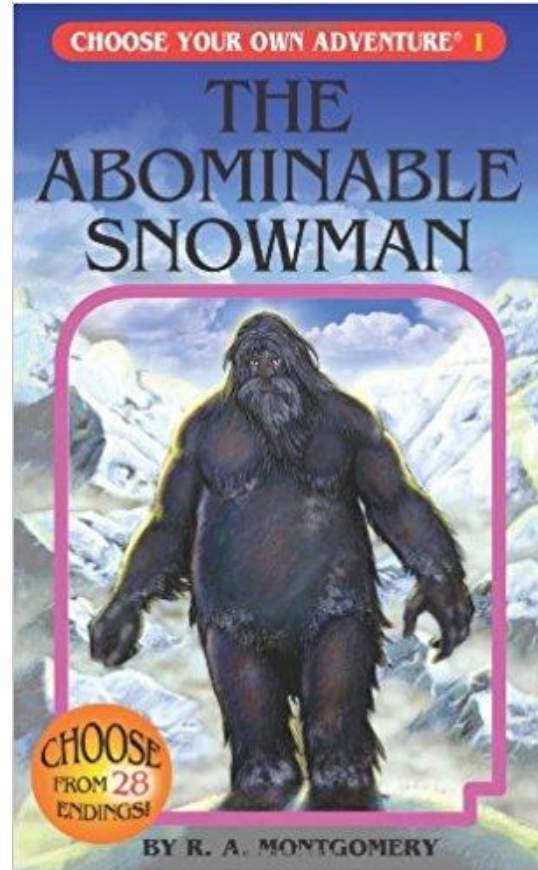
bool maybe_head(entry e, int gen) {
    return e.data != 0 && e.gen == gen
        || e.data == 0 && e.gen < gen;
}

bool try_remove_value(geni pos, entry old) {

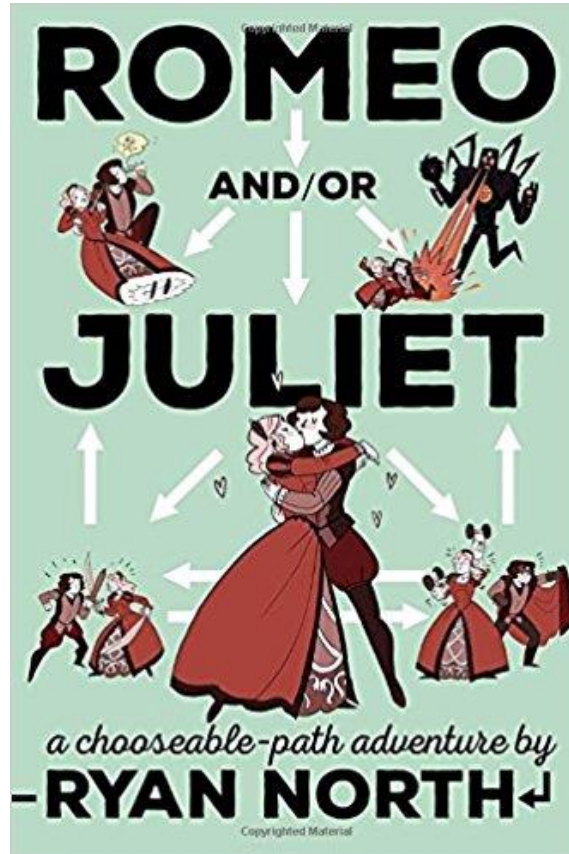
    entry nu{0, pos.gen};
    return buffer[pos].c_e_weak(old, nu, acq_rel, relaxed);
}

```

Choose Your Own Adventure



Chooseable-path Adventure



Choose Your Own Adventure

- Make head/tail *disappear*
- Store non-ints
- Wait on full/empty (with condvar)
- Grow the queue
- Proofs
- *Something else!*

A word about auto

```
auto pos = tailish;
```

```
---
```

```
geni pos = tailish;
```

A word about auto

```
auto pos = tailish; // pos is laxtomic<geni>
---
geni pos = tailish;
```

A word about auto

```
auto pos = tailish; // pos is laxtomic<geni>
---
geni pos = tailish;
---
auto pos = tailish.load(); // pos is geni
```


A word about auto

```
auto pos = tailish; // pos is laxtomic<geni>
---
geni pos = tailish;
---
auto pos = tailish.load(); // pos is geni
---
Iterator it = v.begin();
```

FCD

FCD

Fear

Certainty

Doubt

Atomicization

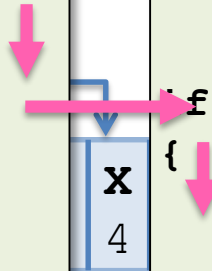
OK?

“normal” lock free code

```
class Queue {  
    atomic<entry> buffer[SIZE];  
    latomic<geni> headish;  
    geni> tailish;
```

```
data.x = 10;  
data.y = 20;  
data.ready.store(true, release);
```

release:
before means before



```
if (data.ready.load(acquire))  
{  
    x = data.x;  
    y = data.y;  
    ...  
}
```

acquire:
after means after

```
int non_atomic_data;  
atomic<bool> initialized;
```

```
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    initialized.store(true, release);  
};
```

```
// User thread  
if (initialized.load(acquire)) {  
    read(non_atomic_data);  
}
```

```
int non_atomic_data;  
atomic<bool> initialized;
```

```
// Thread 1
```

```
void init() {  
    non_atomic_data = 17;  
    initialized.store(true, release);  
};
```



```
// User thread
```

```
if (initialized.load(acquire)) {  
    read(non_atomic_data);  
}
```

```
int non_atomic_data;  
atomic<bool> initialized;  
atomic<bool> counted; // for accounting  
  
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    initialized.store(true, release);  
};  
  
// Thread 2  
atomic<int> keep_track;  
void accounting() {  
    if (initialized && !counted) { // or acquire/relaxed  
        keep_track++;  
        counted.store(true, relaxed);  
    }  
}
```

```
// User thread  
if (initialized.load(acquire)) {  
    read(non_atomic_data);  
}
```

```
int non_atomic_data;  
atomic<bool> initialized;  
atomic<bool> counted; // for accounting  
  
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    initialized.store(true, release);  
};  
  
// Thread 2  
atomic<int> keep_track;  
void accounting() {  
    if (initialized && !counted) { // or acquire/relaxed  
        keep_track++;  
        counted.store(true, relaxed);  
    }  
}
```

```
int non_atomic_data;  
enum State { un, inittd, counted };  
atomic<State> state; // combined!  
  
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    state.store(inittd, release);  
};  
  
// Thread 2  
atomic<int> keep_track;  
void accounting() {  
    if (state == inittd) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```



```
int non_atomic_data;  
enum State { un, initted, counted };  
atomic<State> state; // combined!
```

```
// Thread 1
```

```
void init() {  
    non_atomic_data = 17;  
    state.store(initted, release);  
};
```

```
// Thread 2
```

```
atomic<int> keep_track;  
void accounting() {  
    if (state == initted) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```

```
int non_atomic_data;  
enum State { un, inittd, counted };  
atomic<State> state; // combined!
```

```
// Thread 1
```

```
void init() {  
    non_atomic_data = 17;  
    state.store(inittd, release);  
};
```

```
// Thread 2
```

```
atomic<int> keep_track;  
void accounting() {  
    if (state == inittd) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```

```
// User thread...?
```

```
if (initialized.load(acquire)) {  
    read(non_atomic_data);  
}
```

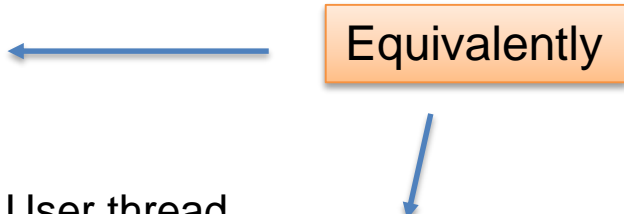
```
int non_atomic_data;  
enum State { un, inittd, counted };  
atomic<State> state; // combined!
```

```
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    state.store(inittd, release);  
};
```

```
// Thread 2  
atomic<int> keep_track;  
void accounting() {  
    if (state == inittd) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```

```
// User thread  
if (state.load(acquire) >= inittd) {  
    read(non_atomic_data);  
}
```

```
int non_atomic_data;  
enum State { un, initted =0x1, counted =0x3 };  
atomic<State> state; // combined!  
  
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    state.store(initted, release);  
};  
  
// Thread 2  
atomic<int> keep_track;  
void accounting() {  
    if (state == initted) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```



← Equivalently ↓

```
// User thread  
if (state.load(acquire) & initted) {  
    read(non_atomic_data);  
}
```

```
int non_atomic_data;  
enum State { un, inittd, counted };  
atomic<State> state; // combined!
```

```
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    state.store(inittd, release);  
};
```

```
// Thread 2  
atomic<int> keep_track;  
void accounting() {  
    if (state == inittd) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```

```
// User thread  
if (state.load(acquire) >= inittd) {  
    read(non_atomic_data);  
}
```

```
int non_atomic_data;  
enum State { un, inittd, counted };  
atomic<State> state; // combined!
```

// Thread 1

```
void init() {  
    non_atomic_data = 17;  
    state.store(inittd, release);  
};
```

// Thread 2

```
atomic<int> keep_track;  
void accounting() {  
    if (state == inittd) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```

// User thread

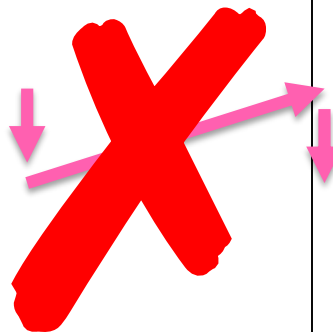
```
if (state.load(acquire) >= inittd) {  
    read(non_atomic_data);  
}
```



```
int non_atomic_data;  
enum State { un, inittd, counted };  
atomic<State> state; // combined!
```

```
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    state.store(inittd, release);  
};
```

```
// Thread 2  
atomic<int> keep_track;  
void accounting() {  
    if (state == inittd) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```



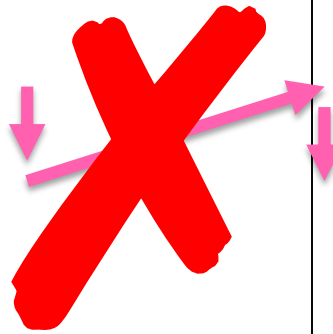
```
// User thread  
if (state.load(acquire) >= inittd) {  
    read(non_atomic_data);  
}
```

UNDEFINED BEHAVIOUR

```
int non_atomic_data;  
enum State { un, inittd, counted };  
atomic<State> state; // combined!
```

```
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    state.store(inittd, release);  
};
```

```
// Thread 2  
atomic<int> keep_track;  
void accounting() {  
    if (state == inittd) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```



```
// User thread  
if (state.load(acquire) >= inittd) {  
    read(non_atomic_data);  
}
```

UNDEFINED BEHAVIOUR

(probably)


```
int non_atomic_data;  
enum State { un, inittd, counted };  
atomic<State> state; // combined!
```

```
// Thread 1  
void init() {  
    non_atomic_data = 17;  
    state.store(inittd, release);  
};
```

```
// Thread 2  
atomic<int> keep_track;  
void accounting() {  
    if (state == inittd) {  
        keep_track++;  
        state.store(counted, relaxed);  
    }  
}
```

```
// User thread  
if (state.load(acquire) >= inittd) {  
    read(non_atomic_data);  
}
```

UNDEFINED BEHAVIOUR

(probably)

```
int r1 = *p;
```

*This is 2 separate loads, the second
depending on the first:*

```
int* ptr = read(p);  
int r1 = read_at(ptr);
```

“The Problem with Threads”

<http://ptolemy.eecs.berkeley.edu/>

<http://ptolemy.eecs.berkeley.edu/publications/papers/06/problemwithThreads/>

“A part of the Ptolemy Project experiment was to see whether **effective software engineering practices** could be developed for an academic research setting. We developed a process that included a code maturity rating system (with four levels, red, yellow, green, and blue), **design reviews, code reviews, nightly builds, regression tests**, and **automated code coverage metrics**. The portion of the kernel that ensured a consistent view of the program structure was written in early 2000, design reviewed to yellow, and code reviewed to green. The **reviewers included concurrency experts**, not just inexperienced graduate students (Christopher Hylands (now Brooks), Bart Kienhuis, John Reekie, and myself were all reviewers). We wrote **regression tests that achieved 100 percent code coverage**. The nightly build and regression tests ran on a two processor SMP machine, which exhibited different thread behavior than the development machines, which all had a single processor. The Ptolemy II **system** itself began to be **widely used**, and every use of the system exercised this code. **No problems were observed until the code **deadlocked** on April 26, 2004, four years later.**”

The Continuing Saga of the Lock-free Queue

Part 3 of N

Tony Van Eerd C++Now May 2018

