

The Julia language and C++

The perfect marriage?

C++Now 2018

Keno Fischer (<https://github.com/Keno>), Bart Janssens (<https://github.com/barche>)

Short bio: Bart

Work: Royal Military Academy (<http://www.rma.ac.be>) of Belgium

C++:

- K-3D (<http://www.k-3d.org>): BoostCon 2007 presentation on Boost.Python (http://www.k-3d.org/k3d/wiki/images/8/8b/Extreme_Object_Models_Using_Boost.Python.pdf) by previous maintanter Tim Shead
- Coolfluid3 (<https://coolfluid.github.io>): C++Now 2013 presentation on using Boost.Proto with Eigen (https://github.com/boostcon/cppnow_presentations_2013/blob/master/fri/proto-eigen-fem.pdf?raw=true)

C++ and Julia:

- CxxWrap.jl (<https://github.com/JuliaInterop/CxxWrap.jl>)
- QML.jl (<https://github.com/barche/QML.jl>)
- Trilinos.jl (<https://github.com/barche/Trilinos.jl>)

Outline

- Why Julia?
- Julia intro: a C++ programmer's perspective
- CxxWrap.jl
- Cxx.jl (<https://github.com/Keno/Cxx.jl>)
- Code repository: <https://github.com/barche/cppnow2018-julia>
(<https://github.com/barche/cppnow2018-julia>)

About Julia

Why another language?

- Solve the "two language problem":
 - Prototype in a simple language
 - Write production code in a fast language

So what is it?

- High-level programming language for scientific computing
- Dynamic language
- Strongly typed, with user-defined types and generics
- JIT-compiled using LLVM
- Native interface to C
- Central concept: **(dynamic) multiple dispatch**

A simple function

```
In [1]: function add(a,b)  
        return a + b  
        end
```

Out[1]: add (generic function with 1 method)

Shorter version:

```
In [2]: add(a,b) = a+b
```

Out[2]: add (generic function with 1 method)

```
In [3]: add(1,2)
```

Out[3]: 3

```
In [4]: add(1.0, 2.0)
```

Out[4]: 3.0

Where are the types?

In [5]: `typeof(1)`

Out[5]: `Int64`

In [6]: `typeof(add(1,2))`

Out[6]: `Int64`

In [7]: `typeof(add(1.0,2))`

Out[7]: `Float64`

A look at the assembly

Each combination of types for the arguments compiles a different version of the code.


```
In [8]: @code_native add(1,2)
```

```
        .section __TEXT,__text,regular,pure_instructions
Filename: In[2]
        pushq    %rbp
        movq     %rsp, %rbp
Source line: 1
        leaq     (%rdi,%rsi), %rax
        popq     %rbp
        retq
        nopw     (%rax,%rax)
```

```
In [9]: @code_native add(1.0,2.0)
```

```
        .section __TEXT,__text,regular,pure_instructions
Filename: In[2]
        pushq    %rbp
        movq     %rsp, %rbp
Source line: 1
        addsd    %xmm1, %xmm0
        popq     %rbp
        retq
        nopw     (%rax,%rax)
```

Equivalence with C++

The Julia function `add(a,b) = a+b` is equivalent to the following C++ function:

```
template<typename A, typename B>
auto add(A a, B b) -> decltype(a + b)
{
    return a + b;
}
```

- Template function valid for any type A and B
- C++ compiles a new version for every combination of types
- Automatic (static) computation of the return type (`decltype` annotation)

Types

```
In [10]: # Define a type  
struct MyNumber  
    n::Int  
end
```

```
In [11]: # Create an instance  
const mynum = MyNumber(2)
```

```
Out[11]: MyNumber(2)
```

What about our add function?

```
In [12]: add(mynum, 2)
```

```
MethodError: no method matching +(::MyNumber, ::Int64)
```

```
Closest candidates are:
```

```
  +(::Any, ::Any, ::Any, ::Any...) at operators.jl:424
```

```
  +(::Complex{Bool}, ::Real) at complex.jl:247
```

```
  +(::Char, ::Integer) at char.jl:40
```

```
...
```

```
Stacktrace:
```

```
[1] add(::MyNumber, ::Int64) at ./In[2]:1
```

```
[2] include_string(::String, ::String) at ./loading.jl:522
```

```
In [13]: add(a::MyNumber, b) = MyNumber(a.n + b)
```

```
Out[13]: add (generic function with 2 methods)
```

```
In [14]: methods(add)
```

```
Out[14]:
```

2 methods for generic function **add**:

- `add(a::MyNumber, b)` at `In[13]:1`
- `add(a, b)` at `In[2]:1`

```
In [15]: add(mynum, 2)
```

```
Out[15]: MyNumber(4)
```

Inheritance

Julia supports single inheritance from abstract base types:

```
In [16]: abstract type MyBase end  
        struct MyConcrete <: MyBase  
            a::Int  
        end  
        # This kind of introspection is built-in:  
        supertype(MyConcrete)
```

Out[16]: MyBase

```
In [17]: geta(x::MyBase) = x.a  
        some_a = MyConcrete(3)  
        geta(some_a)
```

Out[17]: 3

Constness

```
In [18]: some_a.a = 2
```

```
type MyConcrete is immutable
```

```
Stacktrace:
```

```
[1] include_string(::String, ::String) at ./loading.jl:522
```

```
In [19]: # Make a mutable subtype:  
mutable struct MutableConcrete <: MyBase  
    a::Int  
end
```

```
mutable_a = MutableConcrete(4)
```

```
# Still works:
```

```
@show geta(mutable_a)
```

```
# We can change it now:
```

```
mutable_a.a = 42
```

```
geta(mutable_a)
```

```
geta(mutable_a) = 4
```

```
Out[19]: 42
```

Generic types

```
In [20]: struct Point{T,N}  
         coords::NTuple{N,T}  
         end  
         p1 = Point((1,2,3))
```

```
Out[20]: Point{Int64,3}((1, 2, 3))
```

```
In [21]: p2 = Point((2.0, 3.0))
```

```
Out[21]: Point{Float64,2}((2.0, 3.0))
```



```
In [22]: # While this is allowed in C++, it's not in Julia:  
struct Combined{T,N1,N2}  
    coords::NTuple{N1+N2,T}  
end
```

MethodError: no method matching +(::TypeVar, ::TypeVar)

Closest candidates are:

+(::Any, ::Any, ::Any, ::Any...) at operators.jl:424

Stacktrace:

[1] include_string(::String, ::String) at ./loading.jl:522

Multiple dispatch

- In Julia, not writing types means any type can be used
- It still tracks the types and computes the correct return type
- Each combination of argument types yields a newly compiled version of the method
- Type computation can happen dynamically (at runtime) or statically (during JIT compilation)
- The static type computation is equivalent to the C++ version
- Deciding which function to call is **multiple dispatch** (static or dynamic)

Dynamic dispatch

Consider a heterogeneous array:

```
In [23]: A1 = [1, 2.0, [3,4]]
```

```
Out[23]: 3-element Array{Any,1}:  
 1  
 2.0  
 [3, 4]
```

```
In [24]: # Take the norm of every element  
norm.(A1)
```

```
Out[24]: 3-element Array{Real,1}:  
 1  
 2.0  
 5.0
```

Dynamic dispatch

- Because every element has a different type, the `norm` function is chosen at *runtime*, depending on the actual content of the array
- In C++, this can be realized using a common base class and a virtual `norm` function.

Let's add a second array:

```
In [25]: A2 = [[5, 6.0], 7im, 8]
```

```
Out[25]: 3-element Array{Any,1}:  
          [5.0, 6.0]  
          0+7im  
          8
```

Now adding this to the previous array:

```
In [26]: A1 .+ A2
```

```
Out[26]: 3-element Array{Any,1}:  
          [6.0, 7.0]  
          2.0+7.0im  
          [11, 12]
```

This performs a dynamic dispatch based on *both* arguments to +

Julia for C++ programmers: summary

	Julia	C++
Types	✓	✓
Inheritance	Single, abstract	Multiple, concrete
Constness	Type level	Variable level
Methods part of class	✗	✓
Generic types	✓	✓
Computed field types	not yet	✓
Introspection	✓	✗
Multiple dispatch	dynamic	static
Single dispatch	N/A	dynamic
Dot operator overload	✓	✗

CxxWrap

- Github page: <https://github.com/JuliaInterop/CxxWrap.jl> (<https://github.com/JuliaInterop/CxxWrap.jl>)
- Makes use of the native C calling ability of Julia
- Most code for wrapping C++ is written in C++, compiled with your favourite compiler
- Inspired by Boost.Python and pybind11

Calling C functions from Julia

Using `ccall`:

```
In [27]: ccall(:fabs, Float64, (Float64,), -1)
```

```
Out[27]: 1.0
```


The overhead is low:

```
In [28]: using BenchmarkTools
```

```
In [29]: @btime ccall(:fabs, Float64, (Float64,), -1.0)
```

```
3.969 ns (0 allocations: 0 bytes)
```

```
Out[29]: 1.0
```

```
In [30]: @btime abs(-1.0)
```

```
2.538 ns (0 allocations: 0 bytes)
```

```
Out[30]: 1.0
```

CxxWrap approach for functions

- First argument to `ccall` can be a function pointer
- For C-like functions: pass pointer directly
- Otherwise: pass a function pointer that has the actual arguments *and* a pointer to an `std::function` closure

Example

Code from tutorials/cxxwrap/hello_world (tutorials/cxxwrap/hello_world):

```
#include <iostream>
#include <jlcxx/jlcxx.hpp>

void hello() { std::cout << "hello world!" << std::endl; }

JULIA_CPP_MODULE_BEGIN(registry)
    jlcxx::Module& mod = registry.create_module("Hello");

    mod.method("hello", hello);
    mod.method("hello_lambda", [] () { std::cout << "hello lambda!" << std::endl; }
);

JULIA_CPP_MODULE_END
```

CMake:

```
project(Hello)
cmake_minimum_required(VERSION 2.8.12)

find_package(JlCxx REQUIRED)

set(CMAKE_INSTALL_RPATH "${JlCxx_DIR}/../")
set(CMAKE_INSTALL_RPATH_USE_LINK_PATH TRUE)
set(CMAKE_MACOSX_RPATH 1)
set(CMAKE_CXX_FLAGS "${CMAKE_CXX_FLAGS} -std=c++11")

add_library(hello SHARED hello.cpp)
target_link_libraries(hello JlCxx::cxxwrap_julia)
```

Running in Julia

```
In [31]: # Customize this to your actual build directory  
buildroot = joinpath(ENV["HOME"], "src/build/cppnow2018");
```

```
In [32]: using CxxWrap
```

```
In [33]: wrap_modules(joinpath(buildroot, "hello/libhello"))
```

```
In [34]: Hello.hello()
```

```
hello world!
```

```
In [35]: Hello.hello_lambda()
```

```
hello lambda!
```

Behind the scenes

Without the macro:

```
extern "C" void register_julia_modules(void* void_reg)
{
    jlcxx::ModuleRegistry& registry = *reinterpret_cast<jlcxx::ModuleRegistry*>(void_reg);
    try
    {
        jlcxx::Module& mod = registry.create_module("Hello");
        mod.method("hello", hello);
    }
    catch (const std::runtime_error& e) { jl_error(e.what()); }
}
```

- This can be called from Julia using `ccall`
- Other C entry points in `libcxxwrap_julia`

Building a method

The method function is declared as:

```
template<typename R, typename... Args>  
FunctionWrapperBase& method(const std::string& name, std::function<R(Args...)> f  
) ;
```

Here, the `FunctionWrapperBase` object stores all the information needed to build a method:

- The name
- Return type
- Argument types
- Pointers to regular function and `std::function`

In Julia, it looks like this:

```
In [36]: struct CppFunctionInfo
          name::Any
          argument_types::Array{Type,1}
          reference_argument_types::Array{Type,1}
          return_type::Type
          function_pointer::Ptr{Void}
          thunk_pointer::Ptr{Void}
        end
```


The function pointer

The function pointer must be directly callable by Julia. It is a pointer to apply:

```
template<typename R, typename... Args>
struct CallFunctor
{
    using return_type = decltype(ReturnTypeAdapter<R, Args...>()(...));

    static return_type apply(const void* functor, mapped_julia_type<Args>... args)
    {
        try
        {
            return ReturnTypeAdapter<R, Args...>()(functor, args...);
        }
        catch(const std::exception& err)
        {
            jl_error(err.what());
        }
        return return_type();
    }
};
```

Function definition wrap-up

Steps to define the Julia functions:

1. Execute the `register_julia_modules` entry point
2. Julia requests all `CppFunctionInfo` by calling the appropriate C function
3. Methods themselves are defined in Julia, wrapping `ccall` to the appropriate pointers

Difference with Python: No C interface to define methods

So how efficient is this?

Test code:

```
function half_loop(n::Array{Float64,1}, out_arr::Array{Float64,1})  
    test_length = length(n)  
    for i in 1:test_length  
        out_arr[i] = half_x(n[i])  
    end  
end  
  
mod.method("half_loop_cpp",  
    [](jlcxx::ArrayRef<double> in, jlcxx::ArrayRef<double> out)  
    {  
        std::transform(in.begin(), in.end(), out.begin(), [](const double d) { return  
0.5*d; });  
    });
```

Results

Loop over 50 M elements:

	Time (s)
Julia	0.079545
C++	0.081665
ccall	0.145599
cxxwrap, c function	0.143931
cxxwrap, lambda	0.277089
cfunction from C++	0.521564

Exposing types

```
class A
{
public:
    A(int value) : m_value(value) {}

    virtual int get_a() { return m_value; }
private:
    int m_value;
};

JULIA_CPP_MODULE_BEGIN(registry)
    jlccxx::Module& mod = registry.create_module("Types");

    mod.add_type<A>("A")
        .constructor<int>()
        .method("get_a",&A::get_a);

JULIA_CPP_MODULE_END
```

```
In [37]: # Load the module  
wrap_modules(joinpath(buildroot, "types/libtypes"))
```

```
In [38]: a = Types.A(42)
```

```
Out[38]: AAllocated{Ptr{Void} @0x00007f90a1795900}
```

```
In [39]: Types.get_a(a)
```

```
Out[39]: 42
```

Memory management

3 types are created for A:

```
abstract type A <: CxxWrap.CppAny end
```

```
struct ARef <: A  
  cpp_object::Ptr{Void}  
end
```

```
mutable struct AAllocated <: A  
  cpp_object::Ptr{Void}  
end
```

- The AAllocated type runs the destructor upon garbage collect (typically returned by constructors).
- A itself is abstract...

Inheritance

```
class TwiceA : public A
{
public:
    TwiceA(int value) : A(value) {}
    virtual int get_a() { return 2*A::get_a(); }
};

mod.add_type<TwiceA>("TwiceA", jlcxx::julia_type<A>())
    .constructor<int>();
```



```
In [40]: a2 = Types.TwiceA(42)
```

```
Out[40]: TwiceAAllocated{Ptr{Void} @0x00007f90a176a450}
```

```
In [41]: Types.get_a(a2)
```

```
Out[41]: 84
```

```
In [42]: # Only one Julia method exists:  
methods(Types.get_a)
```

```
Out[42]: 1 method for generic function get_a:
```

- `get_a(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:A)`

```
In [43]: supertype(Types.TwiceA)
```

```
Out[43]: A
```

Smart pointers

Some standard library smart pointers are automatically supported:

```
mod.method("smartA", [] (int x) { return std::make_shared<A>(x); } );
```

```
In [44]: smart_a = Types.smartA(1)
```

```
Out[44]: CxxWrap.SmartPointerWithDeref{A,0x0000000121cc0870,Ptr{Void} @0x0000000121caf700,Ptr{Void} @0x0000000121caf730,Ptr{Void} @0x0000000121caf750}(Ptr{Void} @0x0007f90a17b45b0)
```

```
In [45]: Types.get_a(smart_a)
```

```
Out[45]: 1
```

Generic types

The class to wrap:

```
template<typename T1, typename T2>
class SimplePair
{
public:
    typedef T1 x_type;
    typedef T2 y_type;

    SimplePair(const T1 x, const T2 y) : m_x(x), m_y(y) {}
    T1 get_x() const { return m_x; }
    T2 get_y() const { return m_y; }
private:
    T1 m_x;
    T2 m_y;
};
```

Wrapping code for bool:

```
auto pair_type = mod.add_type<jlcxx::Parametric<jlcxx::TypeVar<1>, jlcxx::TypeVar<2>>>("SimplePair");
    // Apply just for bool
    pair_type.apply<SimplePair<bool,bool>>([](auto wrapped)
    {
        typedef typename decltype(wrapped)::type WrappedT;
        typedef typename WrappedT::x_type x_type;
        typedef typename WrappedT::y_type y_type;
        wrapped.template constructor<x_type,y_type>();
        wrapped.method("get_x", &WrappedT::get_x);
        wrapped.method("get_y", &WrappedT::get_y);
    });
```

Automatic combination of types:

```
pair_type.apply_combination<SimplePair,  
                                jlcxx::ParameterList<int, float, double>,  
                                jlcxx::ParameterList<int, float, double>>  
([](auto wrapped)  
{  
    typedef typename decltype(wrapped)::type WrappedT;  
    typedef typename WrappedT::x_type x_type;  
    typedef typename WrappedT::y_type y_type;  
    wrapped.template constructor<x_type,y_type>();  
    wrapped.method("get_x", &WrappedT::get_x);  
    wrapped.method("get_y", &WrappedT::get_y);  
});
```

Testing in Julia

```
In [46]: Types.SimplePair(x::T1,y::T2) where {T1,T2} = Types.SimplePair{T1,T2}(x,y)
```

```
In [47]: p1 = Types.SimplePair(false,false)
```

```
Out[47]: SimplePairAllocated{Bool,Bool}(Ptr{Void} @0x00007f90a1791d40)
```

```
In [48]: Types.get_x(p1)
```

```
Out[48]: false
```

```
In [49]: methods(Types.get_x)
```

Out[49]: 10 methods for generic function **get_x**:

- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Bool,Bool})**
- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Int32,Int32})**
- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Int32,Float32})**
- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Int32,Float64})**
- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Float32,Int32})**
- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Float32,Float32})**
- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Float32,Float64})**
- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Float64,Int32})**
- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Float64,Float32})**
- **get_x(arg1::Union{CxxWrap.SmartPointer{T2}, T2} where T2<:SimplePair{Float64,Float64})**

```
In [50]: p2 = Types.SimplePair{Int32(3),1.0}
```

```
Out[50]: SimplePairAllocated{Int32,Float64}(Ptr{Void} @0x00007f90a17517c0)
```

```
In [51]: p3 = Types.SimplePair("2", 3)
```

```
MethodError: no constructors have been defined for SimplePair{String,Int64}
```

```
Stacktrace:
```

```
[1] SimplePair{::String, ::Int64} at ./In[46]:1
```

```
[2] include_string{::String, ::String} at ./loading.jl:522
```


Template types with non-type parameters

```
template<typename T, int I>
class NonType
{
public:
    NonType(const T x) : m_x(x) {}
    T compute() const { return I*m_x; }
private:
    T m_x;
};
```

Wrapping code

```
mod.add_type<jlcxx::Parametric<jlcxx::TypeVar<1>, jlcxx::TypeVar<2>>>("NonType")
    .apply<NonType<double, 2>>([](auto wrapped)
    {
        typedef typename decltype(wrapped)::type WrappedT;
        wrapped.method("compute", &WrappedT::compute);
    });
```

Fails...

```
error: static_assert failed "No parameters found when applying type. Specialize j
lcxx::BuildParameterList
    for your combination of type and non-type parameters."
```

Building parameter lists

```
template<typename T>
struct BuildParameterList
{
    typedef ParameterList<> type;
};

// Match any combination of types only
template<template<typename...> class T, typename... ParametersT>
struct BuildParameterList<T<ParametersT...>>
{
    typedef ParameterList<ParametersT...> type;
};
```

Need to add:

```
template<typename T, int I>
struct BuildParameterList<NonType<T, I>>
{
typedef ParameterList<T, std::integral_constant<int, I>> type;
};
```

Finally, in Julia:

```
In [53]: nt = Types.NonType{Float64,Int32(2)}(2)
```

```
Out[53]: NonTypeAllocated{Float64,2}(Ptr{Void} @0x00007f90a1c66490)
```

```
In [54]: Types.compute(nt)
```


```
Out[54]: 4.0
```

Trilinos.jl (<https://github.com/barche/Trilinos.jl>)

- Interface to some parts of the Trilinos (<https://trilinos.org/>) C++ library
- Focus on the Tpetra matrix library and iterative solvers
- Uses CxxWrap.jl
- Example: 2D Laplace in C++ and Julia

Laplace example:

Solve the equation $\nabla^2 \varphi + f = 0$ with $f = 2h^2((1 - x^2) + (1 - y^2))$ and 0 on the boundary:



C++ code

```
template<typename MatrixT>
void fill_laplace2d(MatrixT& A, const CartesianGrid& g)
{
    Teuchos::TimeMonitor local_timer(*fill_time);
    const auto& rowmap = *A.getRowMap();
    const local_t n_my_elms = rowmap.getNodeNumElements();

    // storage for the per-row values
    global_t row_indices[5] = {0,0,0,0,0};
    scalar_t row_values[5] = {4.0,-1.0,-1.0,-1.0,-1.0};

    for(local_t i = 0; i != n_my_elms; ++i)
    {
        const global_t global_row = rowmap.getGlobalElement(i);
        const local_t row_n_elms = laplace2d_indices(row_indices, global_row, g);
        row_values[0] = 4.0 - (5-row_n_elms);
        A.replaceGlobalValues(global_row, Teuchos::ArrayView<global_t>(row_indices,row_n_elms), Teuchos::ArrayView<scalar_t>(row_values,row_n_elms));
    }
}
```


Julia code

```
function fill_laplace2d!(A, g::CartesianGrid)
    rowmap = Tpetra.getRowMap(A)
    n_my_elms = Tpetra.getNodeNumElements(rowmap)

    # storage for the per-row values
    row_indices = [0,0,0,0,0]
    row_values = [4.0,-1.0,-1.0,-1.0,-1.0]

    for i in 0:n_my_elms-1
        global_row = Tpetra.getGlobalElement(rowmap,i)
        row_n_elms = laplace2d_indices!(row_indices, global_row, g)
        row_values[1] = 4.0 - (5-row_n_elms)
        Tpetra.replaceGlobalValues(A, global_row, Teuchos.ArrayView(row_indices,row_n
        _elms), Teuchos.ArrayView(row_values,row_n_elms))
    end
end
```

Timing comparison for 1000 x 1000 matrix (in ms)

	C++	Julia
Graph construction	190.7	115.2
Source term	41.17	32.01
Matrix filling	137.1	102.5
Dirichlet setup	0.899	0.761
Check time	41.06	29.42

Conclusions

- C++ integration with minimal overhead
- Expanding on C++ code in pure Julia is possible
- Some CxxWrap todo's:
 - Better Array wrappers
 - std containers
 - exploit Julia dot operator overload for member access