

-Os Matters

Mark Zeren

C++Now, May 10, 2018




Credit to Ole Agesen

Fast Growth Company

- We add features, ergo
- We add code.
- We under-allocate tech debt cleanup, ergo
- We never remove code.
- (It's C++. It's big!)

Growth



We don't see the graph

- Code is brittle
- Code is slow
- Code is hard to understand

Contributors

- Add a feature
- Duck

Everyone has seen this

Mature Product

- Features being added elsewhere in the company
- The most important feature we can provide is to do what we are doing, just better.

Complexity

Mature Product

- Motivated but,
 - Many lines of attack
 - Many opinions
- Requires
 - Staffing
 - Leadership
 - Workflow

The Herding Problem

A classic problem in
software engineering

Metrics

How do we collectively
make forward progress?


Of course we measure:

- Ops / second
- Resources / op
- Max capacity
- Uptime
- Micro-benchmarks
- Coverage...

But...

- Do not drive decreases in *complexity*
- (Except maybe coverage)
- Some drive increases in complexity
- Latency of integration tests
- How to incentivize complexity reduction
- At every commit?

Enter: Binary size



mylib: Remove temporaries

Remove unnecessary string copies.

Testing Done:

- * Unit tests.

- * mylib size:

text	data	bss	dec	filename
194527	21900	74873	291300	mylib-after
194647	21900	74873	291420	mylib-before
-120	0	0	-120	bytes

```
bash-4.1$ ls -ld ../src-*
```

```
drwxr-xr-x 19 mzeren mts 4096 Apr 24 13:17 ../src-base
```

```
drwxr-xr-x 18 mzeren mts 4096 Apr  3 18:21 ../src-this
```

```
drwxr-xr-x 18 mzeren mts 4096 Apr  5 18:14 ../src-that
```

```
bash-4.1$ ls -ld ../bld-*
```

```
drwxr-xr-x 19 mzeren mts 4096 Apr 25 20:28 ../bld-base
```

```
drwxr-xr-x 18 mzeren mts 4096 Apr 25 21:15 ../bld-this
```

```
drwxr-xr-x 18 mzeren mts 4096 Apr 25 21:34 ../bld-that
```

```
bash-4.1$
```



```
bash-4.1$ size mylib.so
```

text	data	bss	dec	hex	filename
194527	21900	74873	74873	12479	mylib.so

```
bash-4.1$ src-base/scripts/mylib-size bld-this bld-base  
mylib size:
```

text	data	bss	dec	filename
194527	21900	74873	291300	mylib-after
194647	21900	74873	291420	mylib-before
-120	0	0	-120	bytes

```
bash-4.1$
```


- Part of your edit compile loop
- Copy-n-paste into review description

“You’ve got to be kidding”

No, not really...

- People keep multiple projects anyway.
- You want to run locally
- It's like a unit test
- Did I pass my test?
- Did I bloat my app?
- It is part of self-review
- It is part of design

Early Days



Manual de-virtualization, base class simplification

OK, but context please?

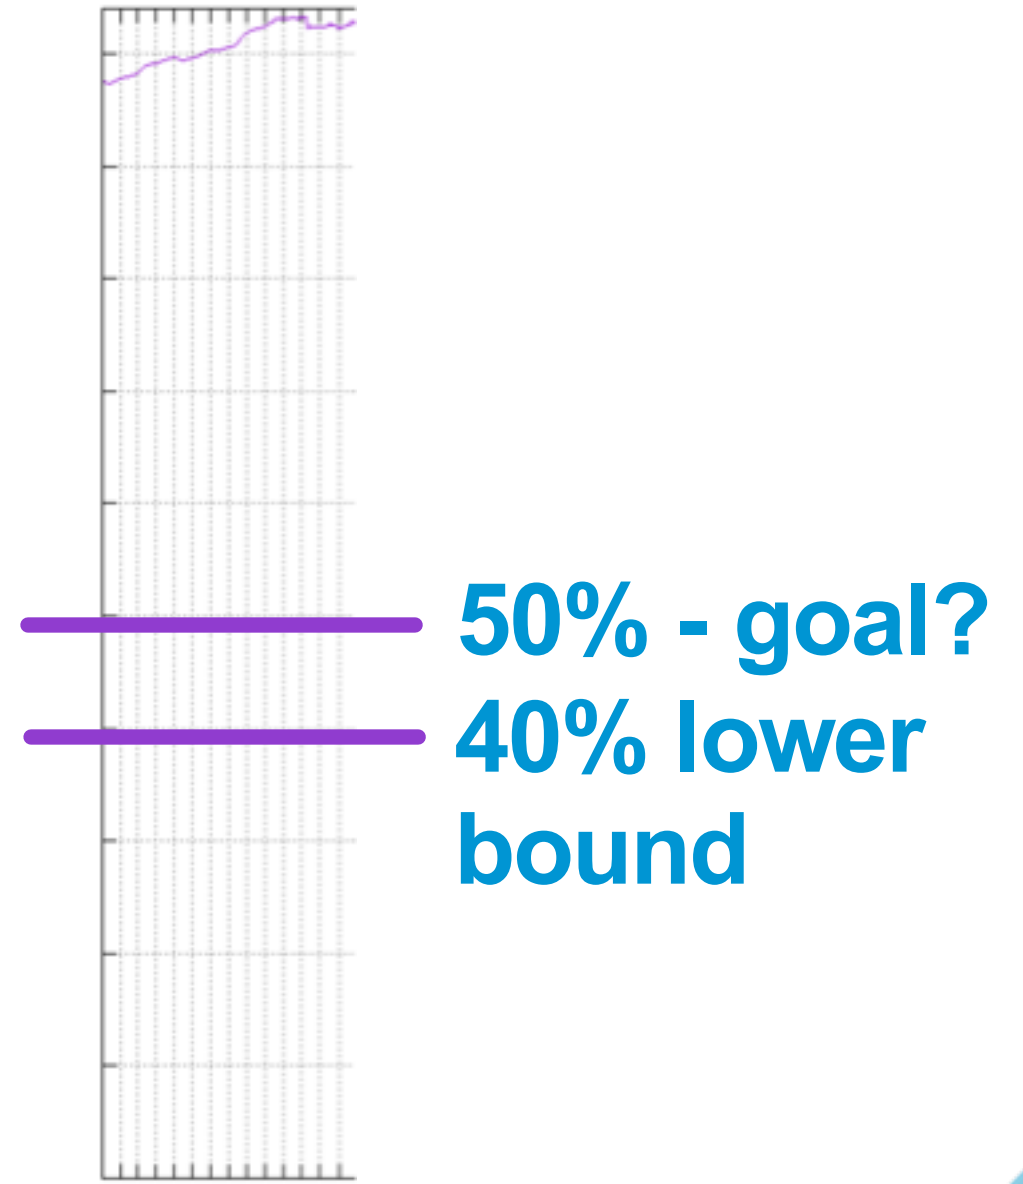


OK, but context please?




1%
gcc 4.3 -O2
Release
Stripped
Nightly
Years


OK, how far can we go?



OK, how far can we go?




What else do you see here?




Ole

Tools




- **Libraries, emitted code, compilers, etc.**
- **Percentages: 0.5%, 1%, 5%**
- **Discontinuities**

Humans

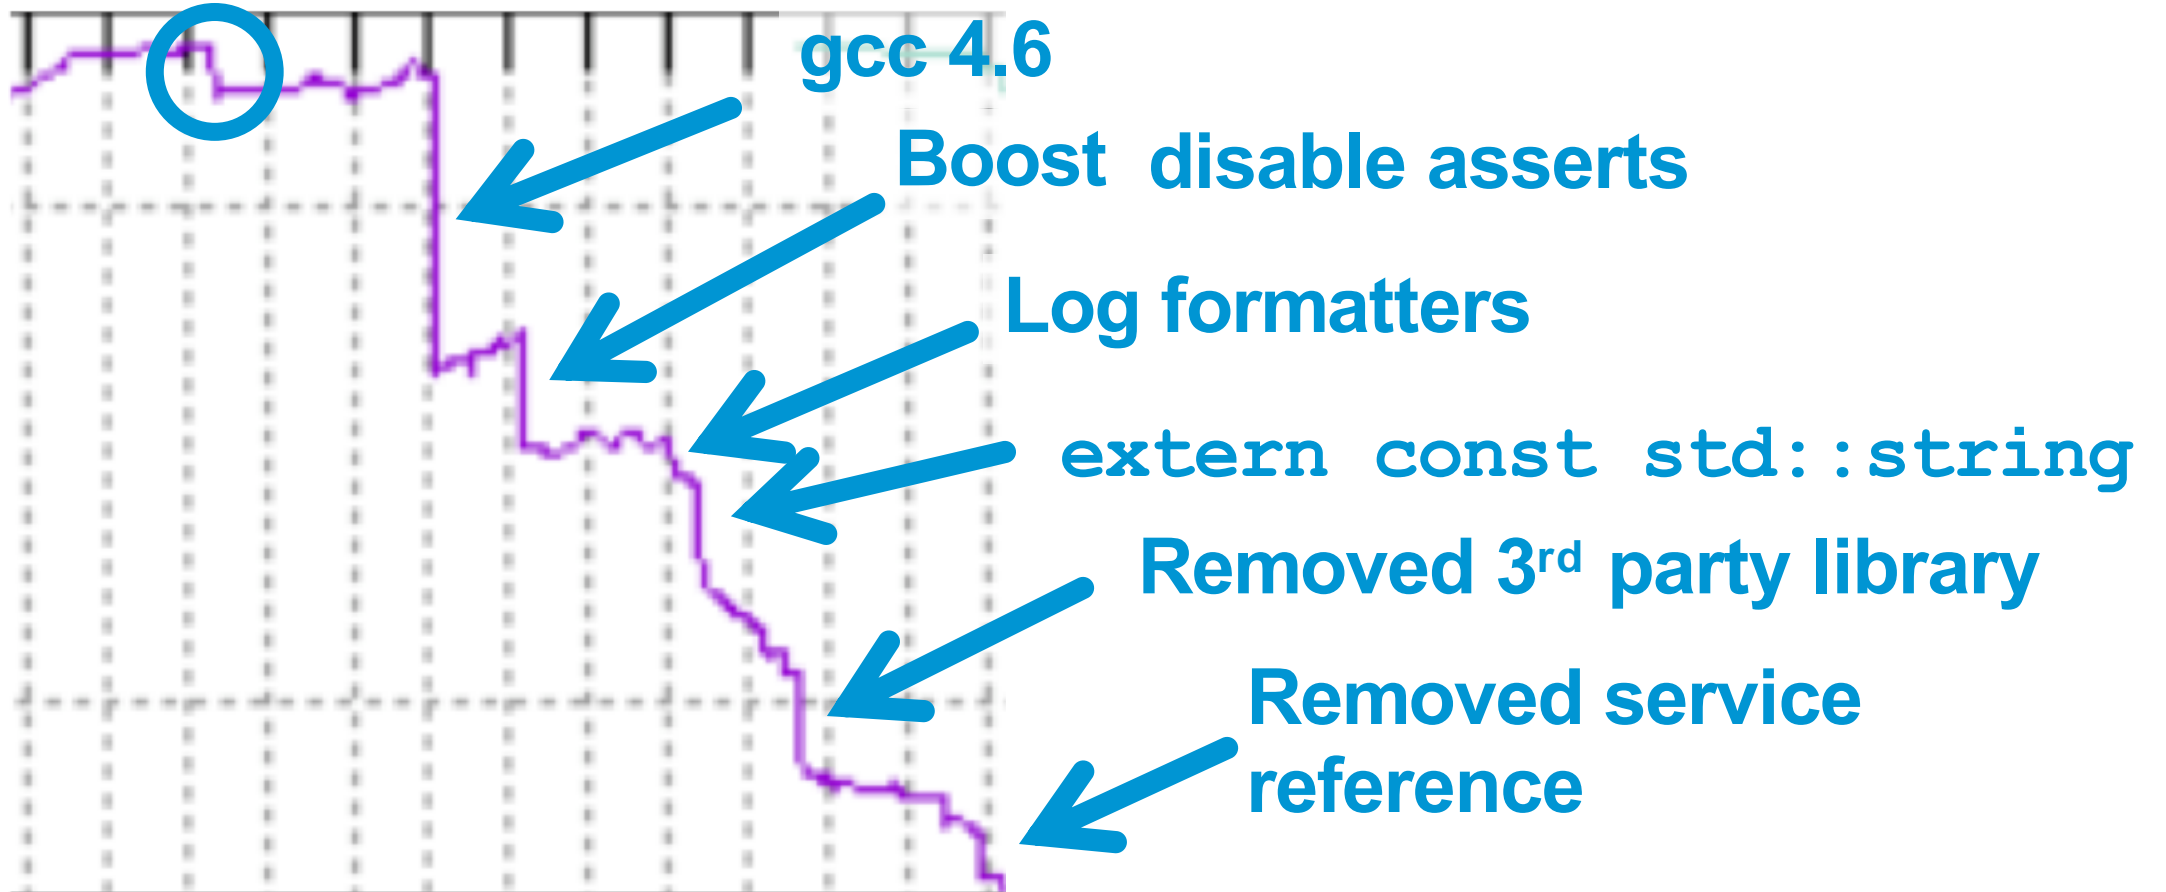


- Step by step, easy, small changes
- Absolute numbers: 150 bytes, 1K, 20K, 0K

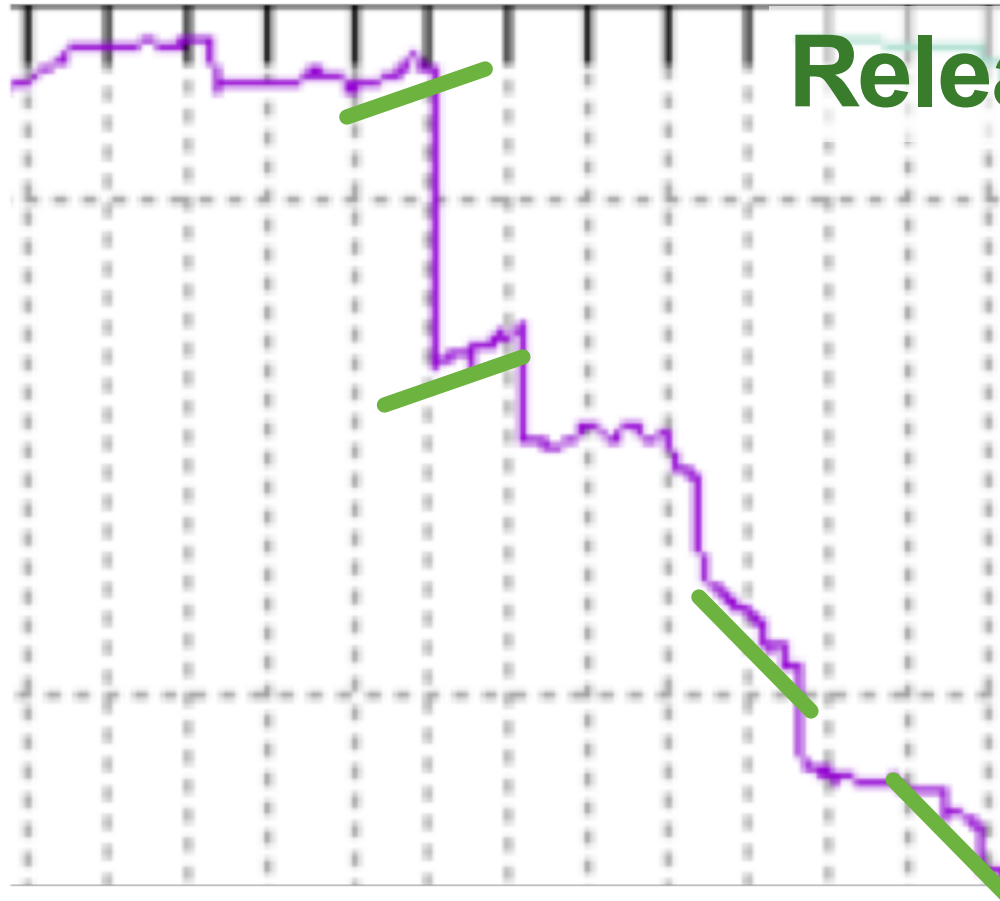
Gaining Momentum ...



Gaining Momentum



Gaining Momentum



Release crunch

Motivated team

mylib: Remove string copies

- * Switch to const& locals to avoid string copies.
- * Other no-op cleanup.

Testing Done:

- * Unit tests
- * mylib size:

text	data	bss	dec	filename
194527	21900	74873	291300	mylib-after
194647	21900	74873	291420	mylib-before
-120	0	0	-120	bytes

mylib: Remove more string copies, etc.

- * More const& locals.
- * More no-op cleanup.

Testing Done:

- * Unit tests
- * mylib size:

text	data	bss	dec	filename
194757	21900	74873	291530	mylib-after
194647	21900	74873	291420	mylib-before
110	0	0	110 bytes	


the big hammer

-Os
the big hammer

-Os

- ~~the big hammer~~
+ the small hammer

The small hammer...



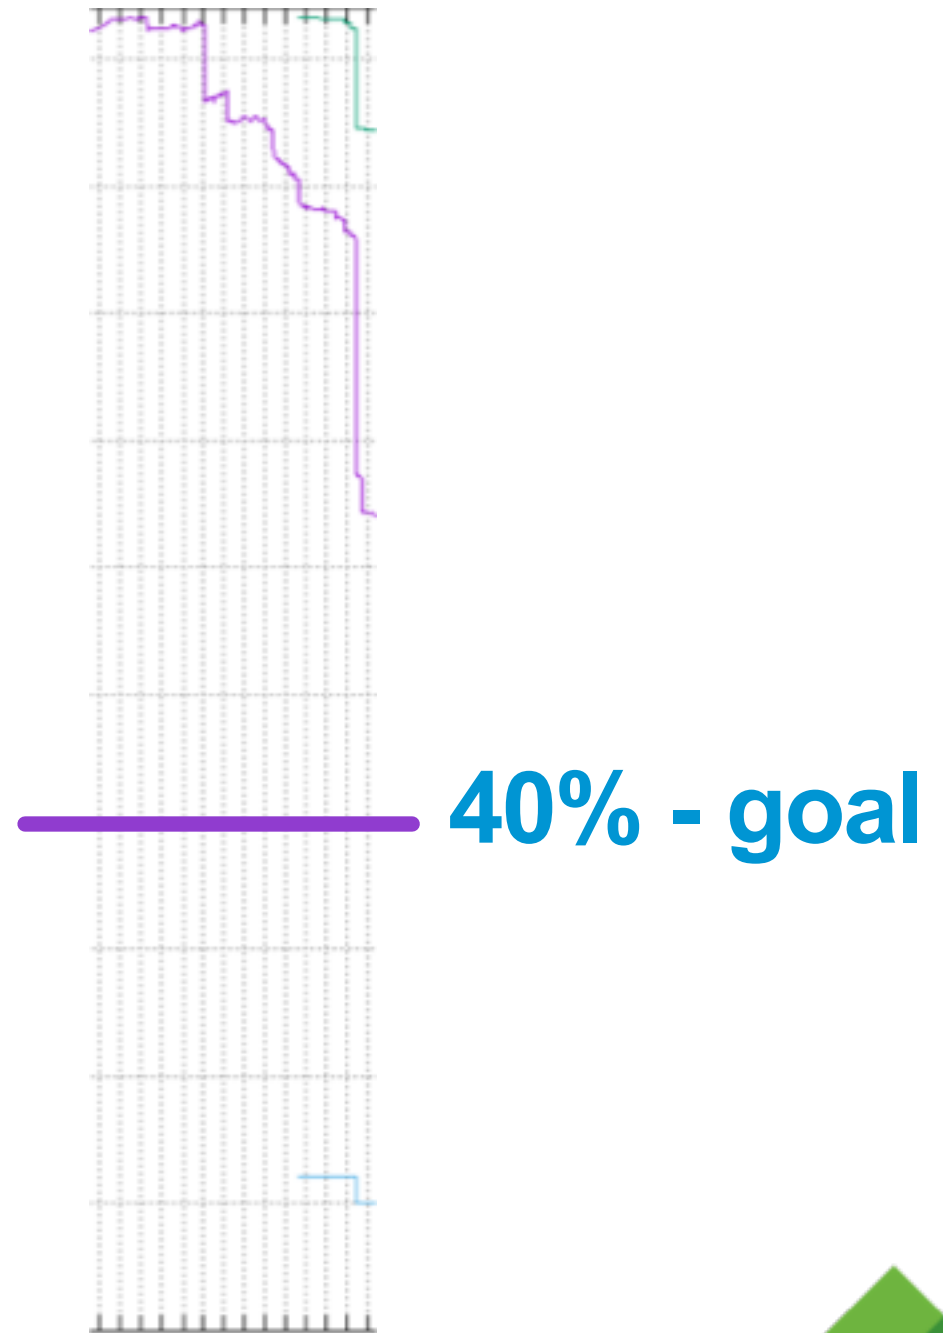
Requirement

- Measure our most important problem: complexity


Rationale

- Binary is huge anyway
 - Larger than I-cache, TLB, predictors, etc.
- We measured
 - Throughput delta was in the noise
 - Micro benchmarks of course show changes


Kobayashi Maru!





How does this movie end?





40% - goal











Real Work Too

- Algorithmic optimization
- Removing whole subsystems
- Adding entirely new features

C++11 Things

- out parameters -> return values
- Reference types to value types
- `bind` -> lambda
- Range based for
- `final`

General Things

- Simplify class structure
- Remove dead features
- Remove layers of indirection
- Remove un-used instrumentation
- Remove simulators and test code
- Rationalize logging
- Exception factories
- Etc.

app: Add frob::ping

* Add the ping method to the frob service.

Testing Done:

* Unit tests

* mylib size:

text	data	bss	dec	filename
196527	21900	74873	293300	mylib-after
194647	21900	74873	291420	mylib-before
2120	0	0	2120 bytes	

app: Add frob::pong

* Add the pong method to the frob service.

Testing Done:

* Unit tests

* mylib size:

text	data	bss	dec	file
206527	21900	74873	303300	mylib-after
196647	21900	74873	293420	mylib-before
10120	0	0	10120	bytes



Review

- Small object copies
- Class design
- Use a library
- Avoid library
- Etc.


Self Review

- If there are two ways, choose the smaller one.

The Rule is to Measure

The Measure is Not the Rule

Sometimes an increase in size is a
decrease in complexity



The Measure is not the Rule

- `bind -> lambda`
- **RAII**
- Range based for
- `unique_ptr`
- `<algorithm>`
- First instantiation of a template
- Etc.

Outside of VMware





JASON TURNER

Rich Code For Tiny Computers: A Simple Commodore 64 Game in C++ 17

```
Compiler Explorer - C++ - Fullscreen
Source on GitHub  Donate  Share  About

78
79     VIC_II &vic;
80 };
81
82     return Frame(*this);
83 }
84 };
85
86 }
87
88 int main()
89 {
90     VIC_II vic;
91
92     while (true) {
93         const auto frame = vic.frame();
94         if (const auto joy = JoyStick(1); joy.fire)
95         {
96             ++vic.background();
97         }
98     }
99 }
100
101
102
103
104

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15

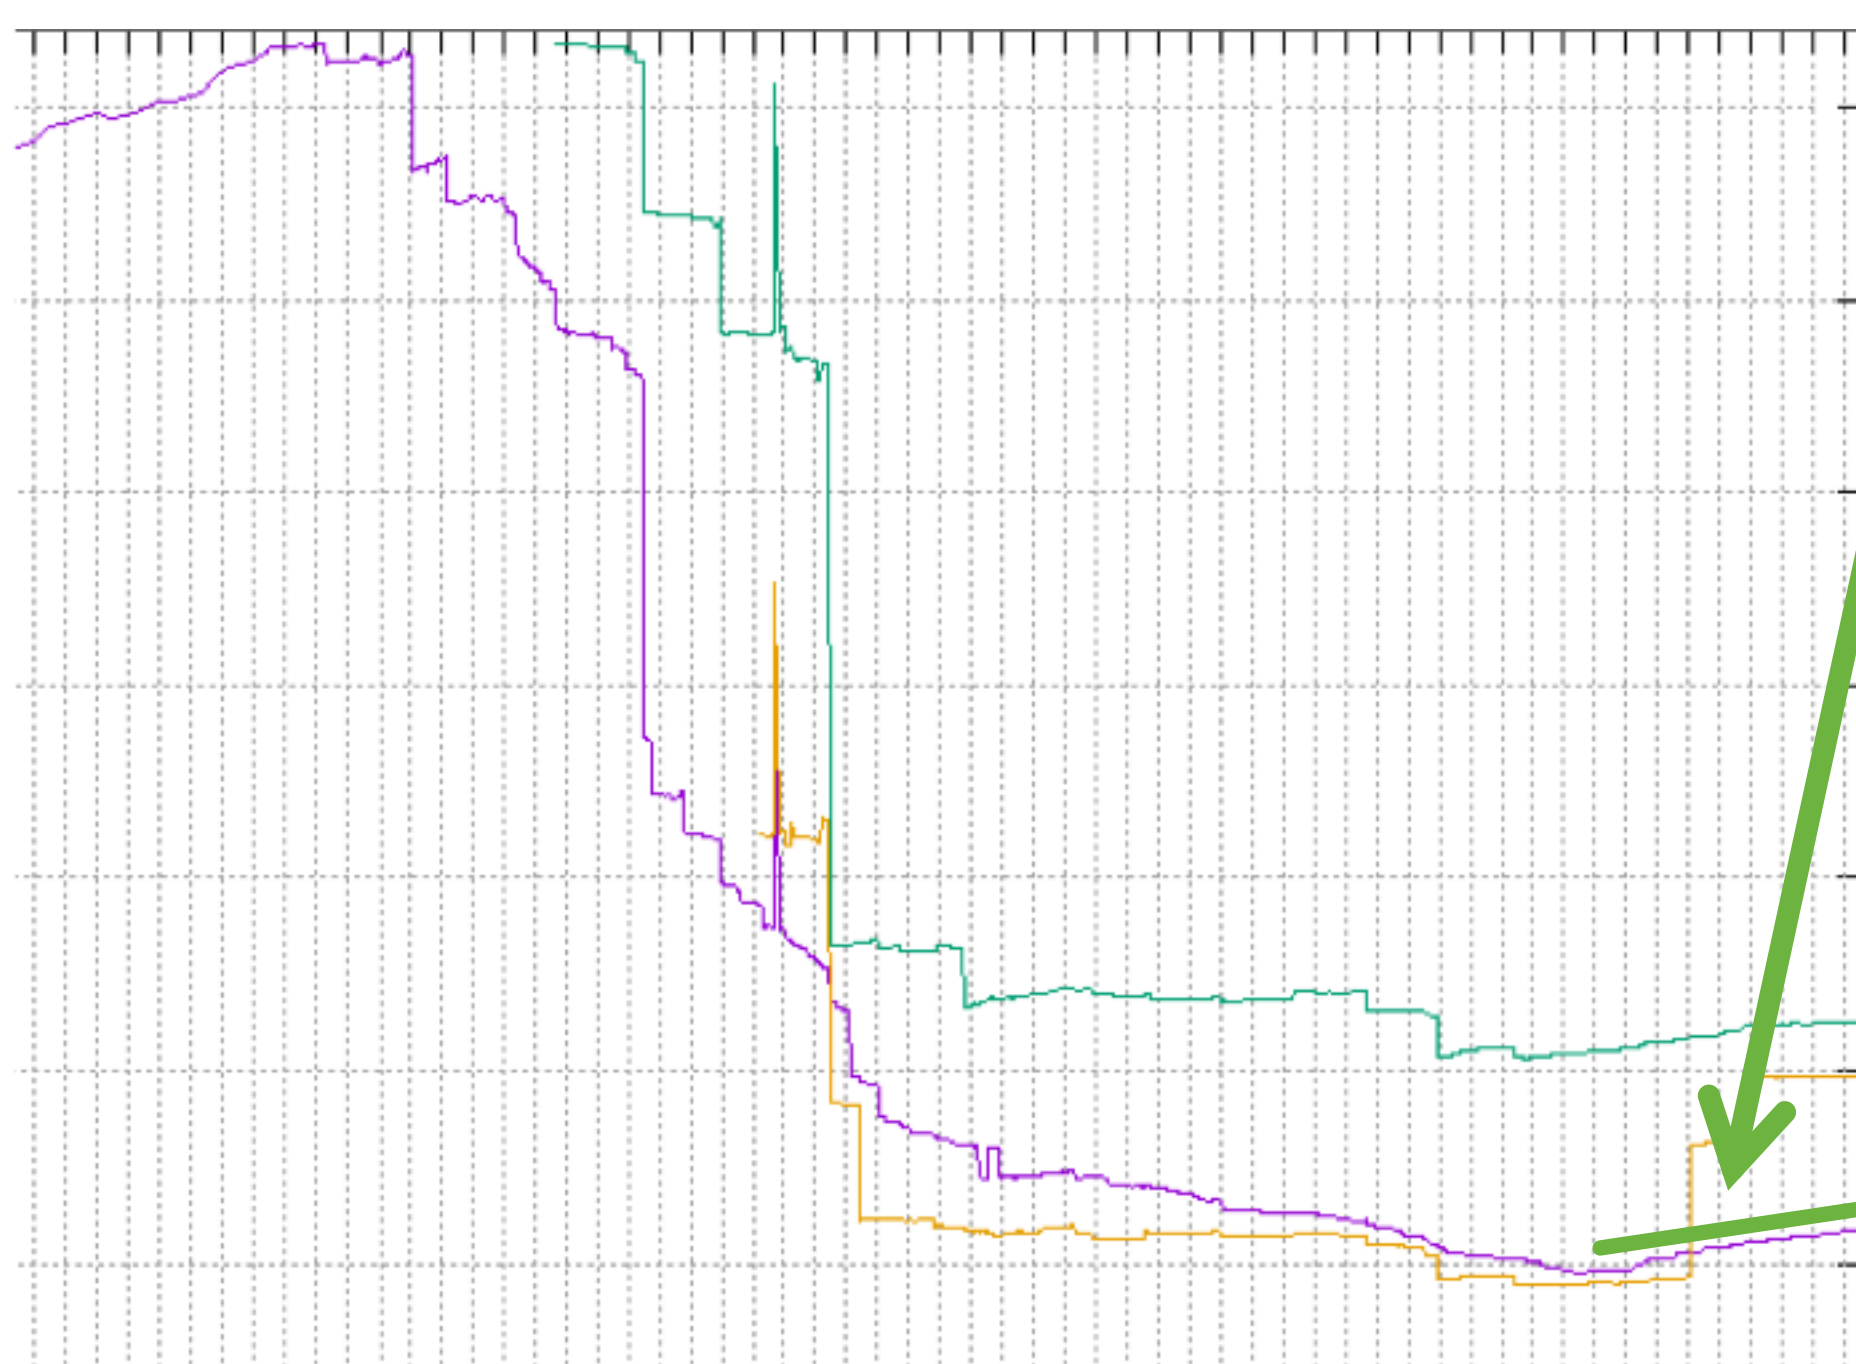
main:
# @main
2 jmp .LBB0_1
3 .LBB0_4:
# %if.end
4 incb 53280
5 .LBB0_1:
# %while.cond.i.i
6 movzbl 53281, %eax
7 cmpb $-16, %al
8 jne .LBB0_1
9 testb $16, 56321
10 jne .LBB0_4
11 incb 53281
12 jmp .LBB0_4
13
14
15

clang version 4.0.0 (http://llvm.org/gn/gn/clang.git/refs/tags/llvm-4.0.0)
Target: x86_64-pc-linux-gnu
```

Outside of VMware

[**https://github.com/google/bloaty**](https://github.com/google/bloaty)

Wrapping Up



- **Business cycle**
- **Adding value**

Future Directions

- Bloaty!
 - We have scripts of a certain age
 - `size` has an arbitrary notion of `.text`
 - Select and disassemble
 - Bloaty should be in your toolbox.

Future Directions

- Back to `-O2`?
 - Occasional Gotchas
 - Some builds really should be `-Os`
 - Test what you deploy
 - Combine with clang plugins
 - Build system integration

Summing Up

- Managing complexity is part of the product lifecycle
- –Os size is a convenient proxy for complexity
- Make it part of the edit, test loop
- Make it part of the review cycle
- Make it part of design

Questions?

Mark Zeren

mzeren@vmware.com

cpplang.slack.com